🍰 Interview Cake

You've built an inflight entertainment system with on-demand movie streaming.

Users on longer flights like to start a second movie right when their first one ends, but they complain that the plane usually lands before they can see the ending. **So you're building a feature for choosing two movies whose total runtimes will equal the exact flight length.**

Write a function that takes an integer `flightLength` (in minutes) and a vector of integers `movieLengths` (in minutes) and returns a boolean indicating whether there are two numbers in `movieLengths` whose sum equals `flightLength`.

When building your function:

- Assume your users will watch *exactly* two movies
- Don't make your users watch the same movie twice
- Optimize for runtime over memory

Gotchas

We can do this in $O(n)$ time, where n is the length of `movieLengths`.

Remember: your users shouldn't watch the same movie twice. **Are you sure your function won't give a false positive if the vector has one element that is half `flightLength`?**

Breakdown

How would we solve this by hand? We know our two movie lengths need to sum to `flightLength`. So for a given `firstMovieLength`, we need a `secondMovieLength` that equals `flightLength - firstMovieLength`.

To do this by hand we might go through `movieLengths` from beginning to end, treating each item as `firstMovieLength`, and for each of those check if there's a `secondMovieLength` equal to `flightLength - firstMovieLength`.

How would we implement this in code? We could nest two loops (the outer choosing `firstMovieLength`, the inner choosing `secondMovieLength`). That'd give us a runtime of $O(n^2)$. We can do better.

To bring our runtime down we'll probably need to replace that inner loop (the one that looks for a matching `secondMovieLength`) with something faster.

We could sort the `movieLengths` first—then we could use binary search↧ to find `secondMovieLength` in $O(\lg n)$ time instead of $O(n)$ time. But sorting would cost $O(n lg(n))$, and we can do even better than that.

Could we check for the existence of our `secondMovieLength` in constant time?

What data structure gives us convenient constant-time lookups?

An unordered set!

So we could throw all of our `movieLengths` into an unordered set first, in $O(n)$ time. *Then* we could loop through our possible `firstMovieLengths` and replace our inner loop with a simple check in our unordered set. This'll give us $O(n)$ runtime overall!

Of course, we need to add some logic to make sure we're not showing users the same movie twice...

But first, we can tighten this up a bit. Instead of two sequential loops, can we do it all in one loop? (Done carefully, this will give us protection from showing the same movie twice as well.)

Solution

We make one pass through `movieLengths`, treating each item as the `firstMovieLength`. At each iteration, we:

1. See if there's a `matchingSecondMovieLength` we've seen already (stored in our `movieLengthsSeen` unordered set) that is equal to `flightLength - firstMovieLength`. If there

> is, we short-circuit and return `true`.

2. Keep our `movieLengthsSeen` unordered set up to date by throwing in the current `firstMovieLength`.

C++ ▾

```cpp
bool canTwoMoviesFillFlight(const vector<int>& movieLengths, int flightLength)
{
    // movie lengths we've seen so far
    unordered_set<int> movieLengthsSeen;

    for (int firstMovieLength : movieLengths) {

        int matchingSecondMovieLength = flightLength - firstMovieLength;
        if (movieLengthsSeen.find(matchingSecondMovieLength) != movieLengthsSeen.end()) {
            return true;
        }

        movieLengthsSeen.insert(firstMovieLength);
    }

    // we never found a match, so return false
    return false;
}
```

We know users won't watch the same movie twice because we check `movieLengthsSeen` for `matchingSecondMovieLength` *before* we've put `firstMovieLength` in it!

Complexity

$O(n)$ time, and $O(n)$ space. Note while optimizing runtime we added a bit of space cost.

Bonus

1. What if we wanted the movie lengths to sum to something *close* to the flight length (say, within 20 minutes)?
2. What if we wanted to fill the flight length as nicely as possible with *any* number of movies (not just 2)?
3. What if we knew that `movieLengths` was *sorted*? Could we save some space and/or time?

What We Learned

The trick was to use an unordered set to access our movies *by length*, in $O(1)$ time.

Using hash-based data structures, like unordered maps or unordered sets, is *so common* in coding challenge solutions, it should always be your *first* thought. Always ask yourself, right from the start: "Can I save time by using an unordered map?"

Want more coding interview help?

Check out **interviewcake.com** for more advice, guides, and practice questions.