

Implement a queue with 2 stacks. Your queue should have an enqueue and a dequeue method and it should be "first in first out" (FIFO).

Optimize for the time cost of m calls on your queue. These can be any mix of enqueue and dequeue calls.

Assume you already have a stack implementation and it gives $O(1)$ time push and pop.

Gotchas

We can get $O(m)$ runtime for m calls. Crazy, right?

Breakdown

Let's call our stacks `stack1` and `stack2`.

To start, we could just push items onto `stack1` as they are enqueued. So if our first 3 calls are enqueues of `a`, `b`, and `c` (in that order) we push them onto `stack1` as they come in.

But recall that stacks are last in, first out. If our next call was a `dequeue()` we would need to return `a`, but it would be on the bottom of the stack.

Look at what happens when we pop `c`, `b`, and `a` one-by-one from `stack1` to `stack2`.

Notice how their order is reversed.

We can pop each item 1-by-1 from `stack1` to `stack2` until we get to `a`.

We could return `a` immediately, but what if our next operation was to enqueue a new item `d`? Where would we put `d`? `d` should get dequeued after `c`, so it makes sense to put them next to each-other . . . but `c` is at the bottom of `stack2`.

Let's try moving the other items back onto `stack1` before returning. This will restore the ordering from before the dequeue, with `a` now gone. So if we enqueue `d` next, it ends up on top of `c`, which seems right.

So we're basically storing everything in `stack1`, using `stack2` only for temporarily "flipping" all of our items during a dequeue to get the bottom (oldest) element.

This is a complete solution. But we can do better.

What's our time complexity for m operations? At any given point we have $O(m)$ items inside our data structure, and if we dequeue we have to move all of them from `stack1` to `stack2` and back again. One dequeue operation thus costs $O(m)$. The number of dequeues is $O(m)$, so our worst-case runtime for these m operations is $O(m^2)$.

Not convinced we can have $O(m)$ dequeues and also have each one deal with $O(m)$ items in the data structure? What if our first $.5m$ operations are enqueues, and the second $.5m$ are alternating enqueues and dequeues. For each of our $.25m$ dequeues, we have $.5m$ items in the data structure.

We can do better than this $O(m^2)$ runtime.

What if we didn't move things back to `stack1` after putting them on `stack2`?

Solution

Let's call our stacks `inStack` and `outStack`.

For enqueue, we simply push the enqueued item onto `inStack`.

For dequeue on an empty outStack, the oldest item is at the bottom of inStack. So we dig to the bottom of inStack by pushing each item one-by-one onto outStack until we reach the bottom item, which we return.

After moving everything from inStack to outStack, the item that was enqueued the *2nd* longest ago (after the item we just returned) is at the top of outStack, the item enqueued *3rd* longest ago is just below it, etc. **So to dequeue on a non-empty outStack**, we simply return the top item from outStack.

With that description in mind, let's write some code!

```
class QueueTwoStacks
{
private:
    stack<int> inStack_;
    stack<int> outStack_;

public:
    void enqueue(int item)
    {
        inStack_.push(item);
    }

    int dequeue()
    {
        if (outStack_.empty()) {
            // Move items from inStack to outStack, reversing order
            while (!inStack_.empty()) {
                int newestInStackItem = inStack_.top();
                inStack_.pop();
                outStack_.push(newestInStackItem);
            }

            // If outStack is still empty, raise an error
            if (outStack_.empty()) {
                throw runtime_error("Can't dequeue from empty queue!");
            }
        }

        int result = outStack_.top();
        outStack_.pop();
        return result;
    }
};
```

Complexity

Each enqueue is clearly $O(1)$ time, and so is each dequeue when `outStack_` has items. Dequeue on an empty `outStack_` is order of the number of items in `inStack_` at that moment, which can vary significantly.

Notice that the more expensive a dequeue on an empty `outStack_` is (that is, the more items we have to move from `inStack_` to `outStack_`), **the more $O(1)$ -time dequeues off of a non-empty `outStack_` it wins us in the future.** Once items are moved from `inStack_` to `outStack_` they just sit there, ready to be dequeued in $O(1)$ time. An item never moves "backwards" in our data structure.

We might guess that this "averages out" so that in a set of m enqueues and dequeues the total cost of all dequeues is actually just $O(m)$. To check this rigorously, we can use the accounting method,¹ **counting the time cost *per item* instead of per enqueue or dequeue.**

So let's look at the worst case for a single item, which is the case where it is enqueued and then later dequeued. In this case, the item enters `inStack_` (costing 1 push), then later moves to `outStack_` (costing 1 pop and 1 push), then later comes off `outStack_` to get returned (costing 1 pop).

Each of these 4 pushes and pops is $O(1)$ time. **So our total cost *per item* is $O(1)$.** Our m enqueue and dequeue operations put m or fewer items into the system, giving a total runtime of $O(m)$.

What We Learned

People often struggle with the runtime analysis for this one. The trick is to think of the cost *per item passing through our queue*, rather than the cost per `enqueue()` and `dequeue()`.

This trick generally comes in handy when you're looking at the time cost of not just one call, but " m " calls.

Want more coding interview help?

Check out **interviewcake.com** for more advice, guides, and practice questions.