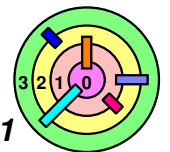



# Housekeeping (Lecture 4 - 9/9/2013)

- ➡ Warmup #1 due at 11:45pm this Friday, 9/13/2013
  - if you have code from a previous semester, be very careful and ***not copy any code from it***
    - it's best if you just get rid of it
  - get started soon
    - if you are stuck, make sure you come to see me during office hour next Monday
- ➡ **Grading guidelines** is the **ONLY** way we will grade and we will grade on `nunki.usc.edu` in our grading account (which you don't have access to)
  - due to our **fairness** policy
  - it's a good idea to run your code against the grading guidelines
- ➡ You need to keep up with the lecture materials
  - anything you don't understand fully, come to see me soon
    - or post a message to the class Google Group



# Housekeeping (Lecture 4 - 9/9/2013)

- ➡ Please do not set your class Google Group e-mail delivery preference to "No email"
- ➡ When you have a question, please do one of the follow and *not both*
  - ▢ post to the class Google Group
    - if your classmates have responded, I may or may not respond
      - ◆ if I don't respond, it does not necessarily mean that your classmates got the right answer
      - ◆ it may mean that they are pointing you in the right direction
      - ◆ you need to learn how to figure out what the right answer is
  - ▢ send me an e-mail directly
    - it may take 24 hours for me to reply
    - for kernel projects, my response may simply be that you should post your question to the class Google Group because I should not be giving you a straight answer
      - ◆ not different from a HW in a regular class, if you ask for an answer, you usually won't get a straight answer



# Creating a POSIX Thread


➡ `man pthread_create`

## SYNOPSIS

```
#include <pthread.h>
```

```
int pthread_create(  
    pthread_t *thread,  
    const pthread_attr_t *attr,  
    void *(*start_routine)(void *),  
    void *arg);
```

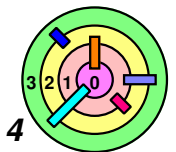
Compile and link with `-pthread`.



# Creating a POSIX Thread

```
start_servers( ) {  
    pthread_t thread;  
    int i;  
    for (i=0; i<nr_of_server_threads; i++)  
        pthread_create(&thread,          // thread ID  
                        0,                // default attributes  
                        server,           // start routine  
                        argument);        // argument  
}  
  
void *server(void *arg) {  
    // perform service  
    return(0);  
}
```

- `pthread_create()` returns 0 if successful
- POSIX 1003.1c standard



# Creating a POSIX Thread

➡ These are the same:

- keep thread handle in the stack


```
pthread_t thread;  
pthread_create(&thread, ...);
```

- keep thread handle in the heap

```
pthread_t *thread_ptr =  
    (pthread_t*)malloc(sizeof(pthread_t));  
pthread_create(thread_ptr, ...);
```

- need to make sure that eventually you will call the following to not leak memory

```
free(thread_ptr);
```




# Creating a Win32 Thread

```
start_servers( ) {
    HANDLE thread;
    DWORD id;
    int i;
    for (i=0; i<nr_of_server_threads; i++)
        thread = CreateThread(
            0,           // security attributes
            0,           // default # of stack pages allocated
            server,       // start routine
            arg,          // argument
            0,           // default attributes
            0,           // creation flags
            &id);         // thread ID
}
```


```
DWORD WINAPI server(void *arg) {
    // perform service
    return(0);
}
```

— We won't talk about Win32 much




# Complications

```
rlogind(int r_in, int r_out, int l_in, int l_out) {  
    pthread_t in_thread, out_thread;  
  
    pthread_create(&in_thread,  
                  0,  
                  incoming,  
                  r_in, l_out);  
    pthread_create(&out_thread,  
                  0,  
                  outgoing,  
                  l_in, r_out);  
  
}
```



# Complications

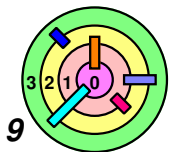
```
rlogind(int r_in, int r_out, int l_in, int l_out) {  
    pthread_t in_thread, out_thread;  
  
    pthread_create(&in_thread,  
                  0,  
                  incoming,  
                  r_in, l_out); // Cannot do this ...  
    pthread_create(&out_thread,  
                  0,  
                  outgoing,  
                  l_in, r_out); // Cannot do this ...  
    /* How do we wait till they are done? */  
}
```






# Multiple Arguments

```
typedef struct {  
    int first, second;  
} two_ints_t;  
  
rlogind(int r_in, int r_out, int l_in, int l_out) {  
    pthread_t in_thread, out_thread;  
  
    two_ints_t in={r_in, l_out}, out={l_in, r_out};  
    pthread_create(&in_thread,  
                  0,  
                  incoming,  
                  &in);  
  
    ...  
    /* How do we wait till they are done? */  
}
```




# Multiple Arguments

- ➡ Need to be careful how to pass multiple arguments to `pthread_create()`
- there is no way to pass multiple arguments in either POSIX or Win32
  - passing address of a *local* variable (like the previous example) only works if we are certain if this storage doesn't go out of scope until the thread is done with it
  - passing address of a *static* or a *global* variable only works if we are certain that only one thread at a time is using the storage
  - passing address of a *dynamically* allocated storage only works if we can free the storage when, and only when, the thread is finished with it
    - this would not be a problem if the language supports garbage collection
- ➡ Ask yourself, "How can you be sure?"
- if the answer is, "I hope it works", then you need a different solution



## When Is The Child Thread Done?

```
rlogind(int r_in, int r_out, int l_in, int l_out) {  
    pthread_t in_thread, out_thread;  
    two_ints_t in={r_in, l_out}, out={l_in, r_out};  
  
    pthread_create(&in_thread, 0, incoming, &in);  
    pthread_create(&out_thread, 0, outgoing, &out);  
  
    pthread_join(in_thread, 0);  
    pthread_join(out_thread, 0);  
}
```



# Thread Termination




## Thread return values

- which threads receive these values
- how do they do it?
  - clearly, receiving thread must wait until the producer thread produced it, i.e., producer thread has terminated
  - so we must have a way for one thread to wait for another thread to terminate
- must have a way to say which thread you are waiting for
  - need a unique identifier
  - tricky if it can be reused



## To wait for another thread to terminate

```
int pthread_join(thread_t thread,  
                 (void **)ret_value);
```



# Thread Termination

➡ How does a thread terminate?

1) return from its "start routine"


◆ return a value of type (void\*)

2) call `pthread_exit (ret_value)`

◆ `ret_value` is of type (void\*)

```
parent() {
    pthread_t thread;
    void *result;
    pthread_create(&thread,
                  0, child, 0);
    pthread_join(thread,
                 (void*)&result);
    switch ((int)result) {
    case 1: ...
    case 2: ...
    }
    ...
}
```

```
void *child(void *arg) {
    ...
    if (terminate_now) {
        pthread_exit((void*)1);
    }
    return((void*)2);
}
```



# Thread Termination




Difference between `pthread_exit()` and `exit()`

- `pthread_exit()` terminates only the calling thread
- `exit()` terminates the process, including all threads running in it
- it will not wait for any thread to terminate
- what will this code do?

```
int main(int argc, char *argv[]) {  
    // create all the threads  
    return(0);  
}
```

- when `main()` returns, `exit()` will be called
  - ◆ as a result, none of the created child thread may get a chance to run



# Thread Termination



Difference between `pthread_exit()` and `exit()`


- `pthread_exit()` terminates only the calling thread
- `exit()` terminates the process, including all threads running in it
- it will not wait for any thread to terminate
- what about this code?

```
int main(int argc, char *argv[]) {  
    // create all the threads  
    pthread_exit(0); // exit the main thread  
    return(0);  
}
```

- here, `pthread_exit()` will terminate the main thread, so `exit()` is never called
  - ◆ as it turns out, this special case is taken care of in the pthread library implementation




You should use `pthread_join()` unless you are absolutely sure



# Thread Termination


- ➡ Any thread can join with any other thread
  - there's **no parent/child relationships among threads**
    - unlike process termination and `wait()`
- ➡ What happens if a thread terminates and no other thread wants to join with this thread?
  - it also goes into a **zombie** state
    - all the thread related information is freed up, **except for the thread ID and return code**
- ➡ What if two threads want to join with the same thread?
  - after the first thread joins, the thread ID and return code are freed up and the thread ID may get reused
  - so don't do this!






# Detached Threads

```
start_servers( ) {  
    pthread_t thread;  
    int i;  
    for (i=0; i<nr_of_server_threads; i++) {  
        pthread_create(&thread, 0, server, 0);  
        pthread_detach(thread);  
    }  
    ...  
}  
  
server( ) {  
    ...  
}
```



# Types


```
pthread_create(&tid,  
              0,  
              (void * (*)(void *)) func,  
              (void *) 1);
```



# Types


```
pthread_create(&tid,  
              0,  
              (void * (*)(void *)) func,  
              (void *) 1);
```

```
int func = 4; // func definition 1
```



# Types

```
pthread_create(&tid,  
              0,  
              (void * (*)(void *)) func,  
              (void *) 1);  
  
int func = 4; // func definition 1  
  
void func(int i) { // func definition 2  
    ...  
}
```




# Types

```
pthread_create(&tid,  
              0,  
              (void (*)(void *))func,  
              (void *)1);
```

```
int func = 4; // func definition 1
```

```
void func(int i) { // func definition 2  
    ...  
}
```

```
void *func(void *arg) { // func definition 3  
    int i = (int)arg;  
    ...  
    return(0);  
}
```



# Thread Attributes

```
pthread_t thread;  
pthread_attr_t thr_attr;
```

```
pthread_attr_init(&thr_attr);  
/* establish some attributes */
```


```
...
```

```
pthread_create(&thread, &thr_attr, startroutine, arg);
```

```
pthread_attr_destroy(&thr_attr);
```

```
...
```


- thread attribute only needs to be valid when a thread is created
  - therefore, it can be destroyed as soon as the thread is created



# Stack Size

```
pthread_t thread;  
pthread_attr_t thr_attr;  
  
pthread_attr_init(&thr_attr);  
pthread_attr_setstacksize(&thr_attr, 20*1024*1024);  
...  
pthread_create(&thread, &thr_attr, startroutine, arg);  
pthread_attr_destroy(&thr_attr);
```

- ▢ the above code set the stack size to 20MB
- ▢ the default stack size is very large
  - if you need to create a lot of threads, you need to control the stack size
  - default stack size is probably around 1MB in Solaris and 8MB in some Linux implementations



# Example

```
#include <stdio.h>
#include <pthread.h>
#include <string.h>

#define M 3
#define N 4
#define P 5


int A[M][N];
int B[N][P];
int C[M][P];

void *matmult(void *arg) {
    int row = (int)arg, col;
    int i, t;

    for (col=0; col < P; col++) {
        t = 0;
        for (i=0; i<N; i++)
            t += A[row][i] * B[i][col];
        C[row][col] = t;
    }
    return(0);
}
```

```
main( ) {
    int i;
    pthread_t thr[M];
    int error;


    /* initialize the matrices ... */
    ...
    // create the worker threads
    for (i=0; i<M; i++) {
        if (error = pthread_create(
            &thr[i],
            0,
            matmult,
            (void *)i)) {
            fprintf(stderr,
                "pthread_create: %s",
                strerror(error));
            exit(1);
        }
    }
    // wait for workers to finish
    for (i=0; i<M; i++)
        pthread_join(thr[i], 0)
    /* print the results ... */
}
```






# Compiling It


```
% gcc -o mat mat.c -lpthread
```



## 2.2.3 Synchronization



# Mutual Exclusion



# Threads and Mutual Exclusion


Thread 1:

$x = x + 1;$

Thread 2:

$x = x + 1;$

- ▢ looks like it doesn't matter how you execute,  $x$  will be incremented by 2 in the end
  - choices are
    - ◆ thread 1 executes  $x = x + 1$  then thread 2 executes  $x = x + 1$
    - ◆ thread 2 executes  $x = x + 1$  then thread 1 executes  $x = x + 1$
  - are there other choices?




# Threads and Mutual Exclusion

Thread 1:

```
x = x+1;
/*
ld    r1,x
add   r1,1
st    r1,x
*/
```

Thread 2:

```
x = x+1;
/*
ld    r1,x
add   r1,1
st    r1,x
*/
```



Unfortunately, machines do not execute high-level language statements

- they execute machine instructions
- now if thread 1 executes the first (or two) machine instructions
- context switch!
  - how can this happen?
- then thread 2 executes all 3 machine instructions
- then later thread 1 executes the remaining machine instructions
- x would have only increased by 1



# Threads and Synchronization

```
// shared by both threads
```

```
pthread_mutex_t m = PTHREAD_MUTEX_INITIALIZER;
```

```
int x;
```

```
...
```

```
pthread_mutex_lock(&m);
```

```
x = x+1;
```


```
pthread_mutex_unlock(&m);
```



Locking a mutex is  
like getting the key to  
a safe-deposit box

*critical section*

- = code between `pthread_mutex_lock()` and `pthread_mutex_unlock()` for a particular *mutex* is called a *critical section* with respect to that mutex
- for a mutex, the OS guarantees that only *one* critical section can be executing at any point in time
- *all the critical sections with respect to a particular mutex* are "*mutually exclusive*"
- how it's really done in the kernel will be covered in Ch 5



# Set Up


```
pthread_mutex_t m = PTHREAD_MUTEX_INITIALIZER;
```

➡ If a mutex cannot be initialized statically, do:

```
int pthread_mutex_init(  
    pthread_mutex_t *mutexp,  
    pthread_mutexattr_t *attrp)
```

```
int pthread_mutex_destroy(  
    pthread_mutex_t *mutexp)
```

➡ Usually, mutex attributes are not used



# Taking Multiple Locks




Mutex is not a cure-all

— when you have more than one locks, you may get into trouble

```
proc1( ) {  
    pthread_mutex_lock(&m1);  
    /* use object 1 */  
    pthread_mutex_lock(&m2);  
    /* use objects 1 and 2 */  
    pthread_mutex_unlock(&m2);  
    pthread_mutex_unlock(&m1);  
}
```

```
proc2( ) {  
    pthread_mutex_lock(&m2);  
    /* use object 2 */  
    pthread_mutex_lock(&m1);  
    /* use objects 1 and 2 */  
    pthread_mutex_unlock(&m1);  
    pthread_mutex_unlock(&m2);  
}
```







# Taking Multiple Locks

```
proc1( ) {  
    pthread_mutex_lock(&m1);  
    /* use object 1 */  
    pthread_mutex_lock(&m2);  
    /* use objects 1 and 2 */  
    pthread_mutex_unlock(&m2);  
    pthread_mutex_unlock(&m1);  
}
```

```
proc2( ) {  
    pthread_mutex_lock(&m2);  
    /* use object 2 */  
    pthread_mutex_lock(&m1);  
    /* use objects 1 and 2 */  
    pthread_mutex_unlock(&m1);  
    pthread_mutex_unlock(&m2);  
}
```

➡ Graph representation



# Necessary Conditions For Deadlocks

➡ All 4 conditions below must be met in order for a deadlock to be possible (no guarantee that a deadlock may occur)

1) **Bounded resources**

⇒ only a finite number of threads can have concurrent access to a resource

2) **Wait for resources**


⇒ threads wait for resources to be freed up, without releasing resources that they hold

3) **No preemption**

⇒ resources cannot be *revoked* from a thread


4) **Circular wait**

⇒ there exists a set of waiting threads, such that each thread is waiting for a resource held by another




# Dealing with Deadlock


- ➡ Deadlock is a programming bug
  - one of the oldest bug
  - it's a tricky one because it only deadlocks *sometimes*
- ➡ Hard
  - is the system deadlocked?
  - will this move lead to deadlock?
  - this is *detection*
    - if you can detect deadlocks, what do you do after you have detected them?
- ➡ Easy
  - restrict use of mutexes so that deadlock cannot happen
  - this is *prevention*




# Deadlock Prevention: Lock Hierarchies




- organize mutexes into levels
- must not try locking a mutex at level  $i$  if already holding a mutex at level  $j$  if  $i < j$ , otherwise it's okay
  - e.g., if hold mutexes at levels 2 & 3, can only wait for a mutex at level 3 or 4



# Deadlock Prevention: Lock Hierarchies




**What if you cannot organize your mutexes in such strict order for deadlock detection?**



# Deadlock Prevention: Conditional Locking

```
proc1( ) {  
    pthread_mutex_lock(&m1);  
    /* use object 1 */  
    pthread_mutex_lock(&m2);  
    /* use objects 1 and 2 */  
    pthread_mutex_unlock(&m2);  
    pthread_mutex_unlock(&m1);  
}  
  
proc2( ) {  
    while (1) {  
        pthread_mutex_lock(&m2);  
  
        if (!pthread_mutex_trylock(&m1))  
            break;  
        pthread_mutex_unlock(&m2);  
    }  
    /* use objects 1 and 2 */  
    pthread_mutex_unlock(&m1);  
    pthread_mutex_unlock(&m2);  
}
```




## Queueing For A Mutex

```
// shared by both threads
pthread_mutex_t m = PTHREAD_MUTEX_INITIALIZER;
int x;
...
pthread_mutex_lock(&m);


x = x+1;


} critical section
pthread_mutex_unlock(&m);
```



- ➡ Can think of `pthread_mutex_lock(&m)` as getting into a *queue* and *wait* there *indefinitely* for mutex `m` to become available
- ➡ multiple threads would join this queue
  - ➡ queue is served one at a time, like a supermarket checkout
  - ➡ when it's your thread's turn, `pthread_mutex_lock()` returns with the mutex locked, your thread can execute critical section code, and then release the mutex

