



# Housekeeping (Lecture 5 - 9/11/2013)

- ➡ Warmup #1 due at 11:45pm this Friday, 9/13/2013
  - if you have code from a previous semester, be very careful and **not copy any code from it**
    - it's best if you just get rid of it
- ➡ **Grading guidelines** is the **ONLY** way we will grade and we will grade on `nunki.usc.edu` in our grading account (which you don't have access to)
  - due to our **fairness** policy
  - it's a good idea to run your code against the grading guidelines
- ➡ Please note that the jobs of the TAs and course producer is **NOT** to do your work for you (and not to teach you C or Unix)
  - that's your job (and learn some stuff on your own)
- ➡ Signing rowsheets will start next Monday
  - if you are registered for one section but would attend lecture in another section, please send me e-mail so I can add your name to the rowsheet




# Beyond Mutexes


- ➡ **Mutex is necessary when shared data is being modified**
  - what if some threads just want to look at (i.e., read) a piece of data?
    - mutex would work, but too restrictive and inefficient (lock threads out when it's not necessary)
- ➡ **Producer-Consumer problem (a.k.a., bounded-buffer problem)**
- ➡ **Readers-Writers problem**
- ➡ **Barrier**




# Producer-Consumer Problem



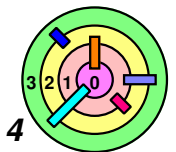
- ➡ A circular buffer is used
- ➡ Most of the time, no interference
  - if you use a *single mutex* to lock the entire array of buffers, it's an overkill (*i.e., too inefficient*)
- ➡ When does it require synchronization?



# Producer-Consumer Problem



- ➡ A circular buffer is used
- ➡ Most of the time, no interference
  - if you use a *single mutex* to lock the entire array of buffers, it's an overkill (*i.e., too inefficient*)
- ➡ When does it require synchronization?
  - producer needs to be blocked when all slots are full
  - consumer needs to be blocked when all slots are empty



# Guarded Commands


```

when (guard) [
  /*
    once the guard is true,
    execute this code atomically
  */
  ...
]


```

} **command sequence**

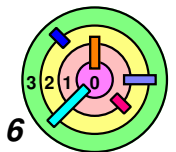
- ⇒ this means that the command sequence *can be* executed (*atomically*) *at any time* the guard is evaluated to be true
  - *atomically* mean that it's executed without interruption
  - since it's executed atomically and without interruption, it is as if it's executed in an instance of time (duration = 0), including the time to evaluate the guard
    - ◇ *evaluting the guard* and *executing the command sequence* altogether is an *atomic operation* if the guard is true



# Semaphores



- ➡ A semaphore,  $S$ , is a nonnegative integer on which there are exactly two operations defined by two guarded commands
- ➡  $P(S)$  operation (implemented as a guarded command):
    - **when**  $(S > 0)$  [
      - $S = S - 1;$
  - ➡  $V(S)$  operation (implemented as a guarded command):
    - $[S = S + 1;]$
  - ➡ there are no other means for manipulating the value of  $S$ 
    - other than initializing it



# Mutexes with Semaphores

```
semaphore S = 1;
```

```
void OneAtATime( ) {
    P(S);
    ...
    /* code executed mutually
       exclusively */
    ...
    V(S);
}
```


— P(S) operation:

○ when (S > 0) [  
     S = S - 1;  
   ]

— V(S) operation:

○ [S = S + 1;]

— this is known as a *binary semaphore*



# Implement A Mutex With A Binary Semaphore

➡ Instead of doing

```
pthread_mutex_lock (&m) ;  
x = x+1;  
pthread_mutex_unlock (&m) ;
```


— do:

```
S = 1;  
P (S) ;  
x = x+1;  
V (S) ;
```

➡ So, you can lock a data structure using a binary semaphore

— this looks just like mutex, what have we really gained?

— if you use it this way, nothing





# Mutexes with Semaphores

```
semaphore S = N;
```

```
void NAtATime( ) {
    P(S);
    ...
    /* code executed mutually
       exclusively */
    ...
    V(S);
}
```

— P(S) operation:

○ when (S > 0) [  
     S = S - 1;  
   ]

— V(S) operation:

○ [S = S + 1;]

— this is known as a *counting semaphore*

— can be used to solve the producer-consumer problem




Main difference between a semaphore and a mutex

— if a thread locks a mutex, it's holding the lock

○ therefore, it must be that thread that unlocks it

— one thread performs a P operation on a semaphore, *another thread* performs a V operation on the same semaphore

○ this is often why you would use a semaphore




# Producer/Consumer with Semaphores

```
Semaphore empty = B;
Semaphore occupied = 0;
int nextin = 0;
int nextout = 0;
```

```
void Produce(char item) {
    P(empty);
    buf[nextin] = item;
    nextin = nextin + 1;
    if (nextin == B)
        nextin = 0;
    V(occupied);
}
```

```
char Consume( ) {
    char item;
    P(occupied);
    item = buf[nextout];
    nextout = nextout + 1;
    if (nextout == B)
        nextout = 0;
    V(empty);
    return(item);
}
```



# Producer/Consumer with Semaphores

```

Semaphore empty = B;
Semaphore occupied = 0;
int nextin = 0;
int nextout = 0;

```

```

void Produce(char item) {
    P(empty);
    buf[nextin] = item;
    nextin = nextin + 1;
    if (nextin == B)
        nextin = 0;
    V(occupied);
}

```



```

char Consume( ) {
    char item;
    P(occupied);
    item = buf[nextout];
    nextout = nextout + 1;
    if (nextout == B)
        nextout = 0;
    V(empty);
    return(item);
}

```

empty	8
occupied	0
nextin	0
nextout	0

Consumer      Producer



# Producer/Consumer with Semaphores



```
Semaphore empty = B;
Semaphore occupied = 0;
int nextin = 0;
int nextout = 0;
```

```
void Produce(char item) {
    → P(empty);
    buf[nextin] = item;
    nextin = nextin + 1;
    if (nextin == B)
        nextin = 0;
    V(occupied);
}
```

```
char Consume( ) {
    char item;
    →✗ P(occupied);
    item = buf[nextout];
    nextout = nextout + 1;
    if (nextout == B)
        nextout = 0;
    V(empty);
    return(item);
}
```

empty	8
occupied	0
nextin	0
nextout	0

Consumer ←      → Producer



# Producer/Consumer with Semaphores

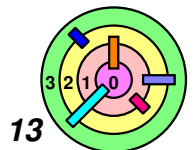

```
Semaphore empty = B;
Semaphore occupied = 0;
int nextin = 0;
int nextout = 0;
```

```
void Produce(char item) {
    P(empty);
    ➔ buf[nextin] = item;
    nextin = nextin + 1;
    if (nextin == B)
        nextin = 0;
    V(occupied);
}
```

```
char Consume( ) {
    char item;
    ➔✗ P(occupied);
    item = buf[nextout];
    nextout = nextout + 1;
    if (nextout == B)
        nextout = 0;
    V(empty);
    return(item);
}
```

empty	7
occupied	0
nextin	0
nextout	0

Consumer      Producer



# Producer/Consumer with Semaphores



```
Semaphore empty = B;
Semaphore occupied = 0;
int nextin = 0;
int nextout = 0;
```

```
void Produce(char item) {
    P(empty);
    buf[nextin] = item;
    → nextin = nextin + 1;
    if (nextin == B)
        nextin = 0;
    V(occupied);
}
```

```
char Consume( ) {
    char item;
    →✗ P(occupied);
    item = buf[nextout];
    nextout = nextout + 1;
    if (nextout == B)
        nextout = 0;
    V(empty);
    return(item);
}
```

empty	7
occupied	0
nextin	0
nextout	0

Consumer      Producer



# Producer/Consumer with Semaphores

```
Semaphore empty = B;
Semaphore occupied = 0;
int nextin = 0;
int nextout = 0;
```



```
void Produce(char item) {
    P(empty);
    buf[nextin] = item;
    nextin = nextin + 1;
    → if (nextin == B)
        nextin = 0;
    V(occupied);
}
```

```
char Consume( ) {
    char item;
    → × P(occupied);
    item = buf[nextout];
    nextout = nextout + 1;
    if (nextout == B)
        nextout = 0;
    V(empty);
    return(item);
}
```

empty	7
occupied	0
nextin	1
nextout	0

Consumer

Producer



# Producer/Consumer with Semaphores

```
Semaphore empty = B;
Semaphore occupied = 0;
int nextin = 0;
int nextout = 0;
```



```
void Produce(char item) {
    P(empty);
    buf[nextin] = item;
    nextin = nextin + 1;
    if (nextin == B)
        nextin = 0;
    → V(occupied);
}
```

```
char Consume( ) {
    char item;
    →✗ P(occupied);
    item = buf[nextout];
    nextout = nextout + 1;
    if (nextout == B)
        nextout = 0;
    V(empty);
    return(item);
}
```

empty	7
occupied	0
nextin	1
nextout	0

Consumer

Producer





# Producer/Consumer with Semaphores

```
Semaphore empty = B;
Semaphore occupied = 0;
int nextin = 0;
int nextout = 0;
```


```
void Produce(char item) {
    P(empty);
    buf[nextin] = item;
    nextin = nextin + 1;
    if (nextin == B)
        nextin = 0;
    V(occupied);
}
```

```
char Consume( ) {
    char item;
    → P(occupied);
    item = buf[nextout];
    nextout = nextout + 1;
    if (nextout == B)
        nextout = 0;
    V(empty);
    return(item);
}
```


empty	7
occupied	1
nextin	1
nextout	0

Consumer

Producer



⇒ note: producer  
continue to produce



# Producer/Consumer with Semaphores

```
Semaphore empty = B;
Semaphore occupied = 0;
int nextin = 0;
int nextout = 0;
```


```
void Produce(char item) {
    P(empty);
    buf[nextin] = item;
    nextin = nextin + 1;
    if (nextin == B)
        nextin = 0;
    V(occupied);
}
```

```
char Consume( ) {
    char item;
    P(occupied);
    → item = buf[nextout];
    nextout = nextout + 1;
    if (nextout == B)
        nextout = 0;
    V(empty);
    return(item);
}
```


empty	7
occupied	0
nextin	1
nextout	0

Consumer

Producer



⇒ note: producer  
continue to produce



# Producer/Consumer with Semaphores

```
Semaphore empty = B;
Semaphore occupied = 0;
int nextin = 0;
int nextout = 0;
```

```
void Produce(char item) {
    P(empty);
    buf[nextin] = item;
    nextin = nextin + 1;
    if (nextin == B)
        nextin = 0;
    V(occupied);
}
```


```
char Consume( ) {
    char item;
    P(occupied);
    item = buf[nextout];
    nextout = nextout + 1;
    if (nextout == B)
        nextout = 0;
    V(empty);
    return(item);
}
```




empty	7
occupied	0
nextin	1
nextout	0

Consumer

Producer



⇒ note: producer  
continue to produce



# Producer/Consumer with Semaphores

```
Semaphore empty = B;
Semaphore occupied = 0;
int nextin = 0;
int nextout = 0;
```


```
void Produce(char item) {
    P(empty);
    buf[nextin] = item;
    nextin = nextin + 1;
    if (nextin == B)
        nextin = 0;
    V(occupied);
}
```

```
char Consume( ) {
    char item;
    P(occupied);
    item = buf[nextout];
    nextout = nextout + 1;
    if (nextout == B)
        nextout = 0;
    V(empty);
    return(item);
}
```




empty	7
occupied	0
nextin	1
nextout	1

Consumer      Producer



⇒ note: producer  
continue to produce



# Producer/Consumer with Semaphores

```
Semaphore empty = B;
Semaphore occupied = 0;
int nextin = 0;
int nextout = 0;
```


```
void Produce(char item) {
    P(empty);
    buf[nextin] = item;
    nextin = nextin + 1;
    if (nextin == B)
        nextin = 0;
    V(occupied);
}
```

```
char Consume( ) {
    char item;
    P(occupied);
    item = buf[nextout];
    nextout = nextout + 1;
    if (nextout == B)
        nextout = 0;
    V(empty);
    return item;
}
```




empty	7
occupied	0
nextin	1
nextout	1

Consumer      Producer



⇒ note: producer  
continue to produce



# Producer/Consumer with Semaphores


```
Semaphore empty = B;
Semaphore occupied = 0;
int nextin = 0;
int nextout = 0;
```

```
void Produce(char item) {
    P(empty);
    buf[nextin] = item;
    nextin = nextin + 1;
    if (nextin == B)
        nextin = 0;
    V(occupied);
}
```


```
char Consume( ) {
    char item;
    P(occupied);
    item = buf[nextout];
    nextout = nextout + 1;
    if (nextout == B)
        nextout = 0;
    V(empty);
    return item;
}
```

empty	8
occupied	0
nextin	1
nextout	1

Consumer      Producer



⇒ note: producer  
continue to produce



# Producer/Consumer with Semaphores


```
Semaphore empty = B;
Semaphore occupied = 0;
int nextin = 0;
int nextout = 0;
```

```
void Produce(char item) {
    P(empty);
    buf[nextin] = item;
    nextin = nextin + 1;
    if (nextin == B)
        nextin = 0;
    V(occupied);
}
```


```
char Consume( ) {
    char item;
    P(occupied);
    item = buf[nextout];
    nextout = nextout + 1;
    if (nextout == B)
        nextout = 0;
    V(empty);
    return(item);
}
```

empty	8
occupied	0
nextin	1
nextout	1

Consumer      Producer



⇒ note: producer  
continue to produce




# Producer/Consumer with Semaphores

```
Semaphore empty = B;
Semaphore occupied = 0;
int nextin = 0;
int nextout = 0;
```

```
void Produce(char item) {
    P(empty);
    buf[nextin] = item;
    nextin = nextin + 1;
    if (nextin == B)
        nextin = 0;
    V(occupied);
}

char Consume( ) {
    char item;
    P(occupied);
    item = buf[nextout];
    nextout = nextout + 1;
    if (nextout == B)
        nextout = 0;
    V(empty);
    return item;
}
```

- if produce and consume at same rate, no one waits
- if producer is fast and consumer slow, producer may wait
- if consumer is fast and producer slow, consumer may wait





# Producer/Consumer with Semaphores

```
Semaphore empty = B;
Semaphore occupied = 0;
int nextin = 0;
int nextout = 0;
```

```
void Produce(char item) {
    P(empty);
    buf[nextin] = item;
    nextin = nextin + 1;
    if (nextin == B)
        nextin = 0;
    V(occupied);
}
```

```
char Consume( ) {
    char item;
    P(occupied);
    item = buf[nextout];
    nextout = nextout + 1;
    if (nextout == B)
        nextout = 0;
    V(empty);
    return(item);
}
```




Mutex is more "coarse grain"

— you may use one mutex to control access to the number of empty and occupied cells, nextin, and nextout




Semaphore is more "fine grain"



# POSIX Semaphores

```
#include <semaphore.h>
sem_t semaphore;
int err;
int pshared = 0;    // not shared among processes
int init_value = B; // initial value

err = sem_init(&semaphore, pshared, init_value);
err = sem_destroy(&semaphore);
err = sem_wait(&semaphore);    /* P operation */
err = sem_trywait(&semaphore); /* conditional P
                                operation
                                */
err = sem_post(&semaphore);    /* V operation */
```



# Producer-Consumer with POSIX Semaphores


```
void produce(char item) {
    sem_wait(&empty);
    buf[nextin++] = item;
    if (nextin >= B)
        nextin = 0;
    sem_post(&occupied);
}
```

```
char consume() {
    char item;
    sem_wait(&occupied);
    item = buf[nextout++];
    if (nextout >= B)
        nextout = 0;
    sem_post(&empty);
    return(item);
}
```

---

```
void Produce(char item) {
    P(empty);
    buf[nextin] = item;
    nextin = nextin + 1;
    if (nextin == B)
        nextin = 0;
    V(occupied);
}
```

```
char Consume( ) {
    char item;
    P(occupied);
    item = buf[nextout];
    nextout = nextout + 1;
    if (nextout == B)
        nextout = 0;
    V(empty);
    return(item);
}
```




# Implementation Of Guarded Commands

```

when (guard) [
    /* command sequence */
    ...
]

```

- ➡ In general, the *guard* can be *complicated* and involving the evaluation of several variables
- can think of the guard as a predicate (i.e., evaluates to either true or false) *that keeps changing its value, continuously and by multiple threads simultaneously*
  - how can we "capture" the instance of time when it evaluates to true so we can execute the command sequence atomically?
    - we have to "sample" it, i.e., take snap shot of all the variables that are involved and then evaluate it
    - once we see that it evaluates to "true", we have to execute the command sequence atomically
    - a mutex is involved, but how?
      - ◆ need to be efficient



# Implementation Of Guarded Commands

```


when (guard) [
    /* command sequence */
    ...
]

```

➡ POSIX provides *condition variables* for programmers to implement guarded commands

⇒ a *condition variable* is a *queue of threads* waiting for some sort of notification (an "event")

- threads waiting for a guard to become true join such a queue
  - ◇ they wait for a *condition* to be *signaled*
- threads that do something to potentially change the value of a guard can then wake up the threads that were waiting in the queue
  - ◇ they can *signal* or *broadcast* the *condition* (sometimes called an "event")
  - ◇ *no guarantee* that the guard will be true when it's time for *another thread* to evaluate the guard



# Implementation Of Guarded Commands

```

when (guard) [
    /* command sequence */
    ...
]

```


➡ POSIX provides *condition variables* for programmers to implement guarded commands

— conceptually, the "event" (signaling/broadcasting of a condition) happens in an instance of time (duration of this "event" is zero)

- if you are not waiting for it, you'll miss it
- how do you make sure you won't miss an event?

1) `pthread_cond_wait(pthread_cond_t *cv, pthread_mutex_t *mutex)`

- should only call `pthread_cond_wait()` if you have the mutex *locked*
- atomically unblocks mutex and wait for the "event"
- when the event is signaled, `pthread_cond_wait()` returns with the mutex *locked*



# Implementation Of Guarded Commands

```
when (guard) [  
    /* command sequence */  
    ...  
]
```




POSIX provides *condition variables* for programmers to implement guarded commands

— conceptually, the "event" (signaling/broadcasting of a condition) happens in an instance of time (duration of this "event" is zero)

- if you are not waiting for it, you'll miss it
- how do you make sure you won't miss an event?

2) `pthread_cond_broadcast(pthread_cond_t *cv)`  
`pthread_cond_signal(pthread_cond_t *cv)`

- should only call `pthread_cond_broadcast()` or `pthread_cond_signal()` if you have the corresponding mutex *locked*



# Implementation Of Guarded Commands

➡ **Synchronization:** mutex, condition variables, guards, critical sections

— with respect to a mutex, a thread can be


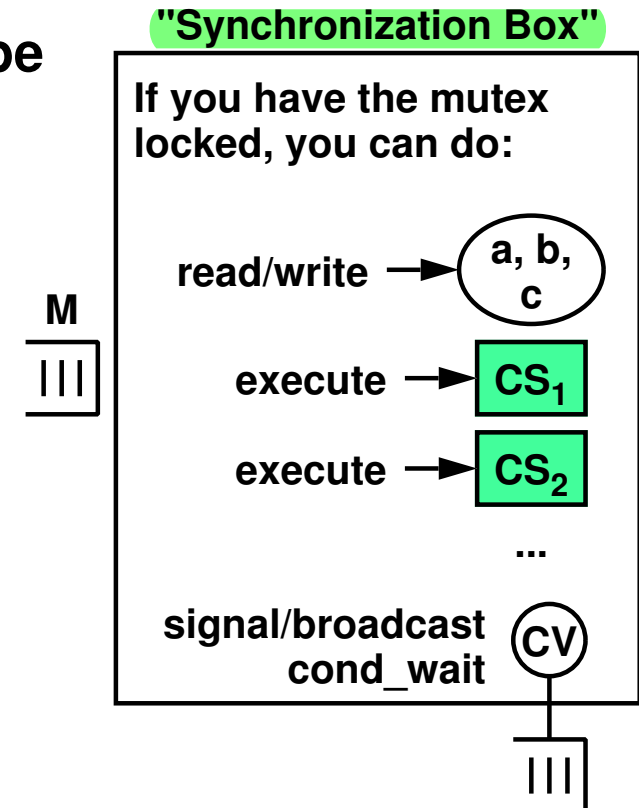
- waiting in the mutex queue
- got the lock and inside the "synchronization box"
  - ◆ only one thread can be inside the "synchronization box"

○ waiting in the CV queue

○ or outside

— with respect to a mutex, a, b, c are variables that can affect the value of the guard predicate

- can only access them if a thread is inside the "synchronization box"


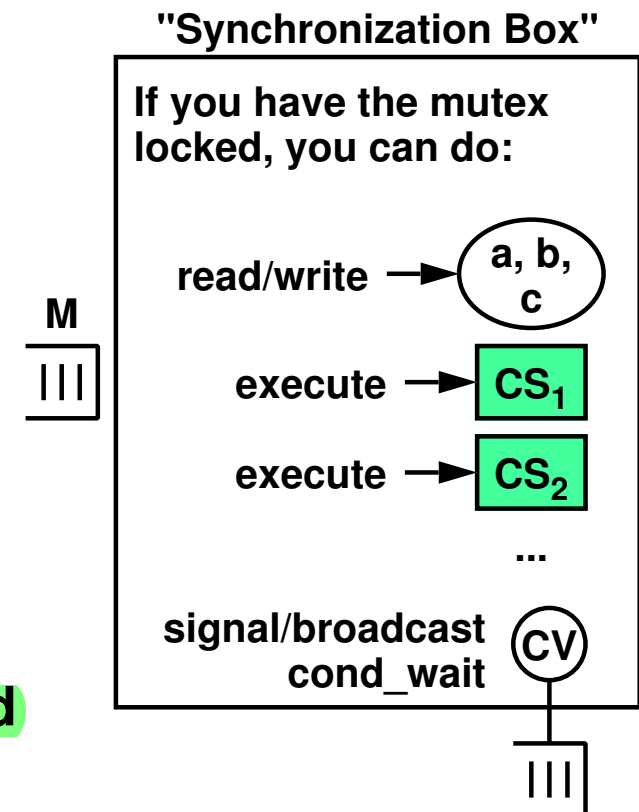




# Implementation Of Guarded Commands


➡ **Synchronization:** mutex, condition variables, guards, critical sections

- when you **signal** CV
  - **one** thread in the CV queue gets moved to the mutex queue
- when you **broadcast** CV
  - **all** threads in the CV queue get moved to the mutex queue
- you can **only** get **added** to the CV queue if you have the mutex locked
- you can only **modify** the variables in the guard if you have the mutex locked
- you can only **read** the variables in the guard (i.e., evaluate the guard) if you have the mutex locked
- you can only **execute** critical section code if you have the mutex locked



# POSIX Condition Variables

<i>Guarded command</i>	<i>POSIX implementation</i>
<pre> when (guard) [     statement 1;     ...     statement n; ]</pre>	<pre> pthread_mutex_lock(&amp;mutex); while (!guard)     pthread_cond_wait(         &amp;cv,         &amp;mutex); statement 1; ... statement n; pthread_mutex_unlock(&amp;mutex);</pre>
<pre> [ /* code    * modifying    * the guard    */ ]</pre>	<pre> pthread_mutex_lock(&amp;mutex); /*code modifying the guard:*/ ... pthread_cond_broadcast(&amp;cv); pthread_mutex_unlock(&amp;mutex);</pre>



# Set Up


```
pthread_cond_t cv = PTHREAD_COND_INITIALIZER;
```

➡ If a condition variable cannot be initialized statically, do:



```
int pthread_cond_init(  
    pthread_cond_t *cvp,  
    pthread_condattr_t *attrp)
```

```
int pthread_cond_destroy(  
    pthread_cond_t *cvp)
```

➡ Usually, condition variable attributes are not used




# Readers-Writers Problem



## Readers-Writers Pseudocode

```
reader( ) {  
    when (writers == 0) [  
        readers++;  
    ]  
    /* read */  
    [readers--;]  
}
```

```
writer( ) {  
    when ((writers == 0) &&  
        (readers == 0))  
    [  
        writers++;  
    ]  
    /* write */  
    [writers--;]  
}
```



# Pseudocode with Assertions

```

reader( ) {
  when (writers == 0) [
    readers++;
  ]
  // sanity check
  assert((writers == 0) &&
        (readers > 0));
  /* read */
  [readers--;]
}


```

```

writer( ) {
  when ((writers == 0) &&
        (readers == 0))
  [
    writers++;
  ]
  // sanity check
  assert((readers == 0) &&
        (writers == 1));
  /* write */
  [writers--;]
}

```

— the sanity checks are really not necessary



# Pseudocode with Assertions

```

reader( ) {
  when (writers == 0) [
    readers++;
  ]
  /* read */
  [readers--;]
}


```

```

writer( ) {
  when ((writers == 0) &&
    (readers == 0))
  [
    writers++;
  ]
  /* write */
  [writers--;]
}

```

- remember, since `readers` is part of the guard in a guarded commend, in the implementation of `[readers--;]`, you must signal/broadcast the corresponding condition used to implement that guard
  - in this case, only have to signal if `readers` becomes 0



# Pseudocode with Assertions

```

reader( ) {
  when (writers == 0) [
    readers++;
  ]
  /* read */
  [readers--;]
}


```

```

writer( ) {
  when ((writers == 0) &&
    (readers == 0))
  [
    writers++;
  ]
  /* write */
  [writers--;]
}

```

- also, since `writers` is part of the guard in a guarded command, in the implementation of `[writers--;]`, you must signal/broadcast the corresponding condition used to implement that guard






# Pseudocode with Assertions

```
reader( ) {  
  when (writers == 0) [  
    readers++;  
  ]  
  /* read */  
  [readers--;]  
}
```

```
writer( ) {  
  when ((writers == 0) &&  
        (readers == 0))  
  [  
    writers++;  
  ]  
  /* write */  
  [writers--;]  
}
```

- don't have to worry about this readers
  - you need to look at your program logic and figure when signal/broad conditions won't be useful



# Solution with POSIX Threads

```

reader( ) {
    pthread_mutex_lock(&m);
    while (!(writers == 0))
        pthread_cond_wait(
            &readersQ, &m);
    readers++;
    pthread_mutex_unlock(&m);
    /* read */
    pthread_mutex_lock(&m);
    if (--readers == 0)
        pthread_cond_signal(
            &writersQ);
    pthread_mutex_unlock(&m);
}


```

```

writer( ) {
    pthread_mutex_lock(&m);
    while (!(readers == 0) &&
        (writers == 0))
        pthread_cond_wait(
            &writersQ, &m);
    writers++;
    pthread_mutex_unlock(&m);
    /* write */
    pthread_mutex_lock(&m);
    writers--;
    pthread_cond_signal(
        &writersQ);
    pthread_cond_broadcast(
        &readersQ);
    pthread_mutex_unlock(&m);
}

```

- one mutex (m) and two condition variables (readersQ and writersQ)



# The Starvation Problem

```


reader( ) {
    when (writers == 0) [
        readers++;
    ]
    /* read */
    [readers--;]
}

writer( ) {
    when ((writers == 0) &&
        (readers == 0))
    [
        writers++;
    ]
    /* write */
    [writers--;]
}

```

- ➡ Can the writer never get a chance to write?
- yes, if there are always readers
  - so, this implementation can be unfair to writers

- ➡ Solution
- once a write arrives, shut the door on new readers
    - writers now means the number of writers *wanting* to write
    - use **active\_writers** to make sure that only one writer can do the actual writing at a time



# Solving The Starvation Problem

```

reader( ) {
    when (writers == 0) [
        readers++;
    ]
    /* read */
    [readers--;]
}


```

```

writer( ) {
    [writers++;]
    when ((readers == 0) &&
        (active_writers == 0))
    [
        active_writers++;
    ]
    /* write */
    [ writers--;
        active_writers--;
    ]
}

```


- now it's unfair to the readers
- isn't writing more important than reading anyway?



# Improved Reader


```
reader( ) {  
    pthread_mutex_lock(&m);  
    while (!(writers == 0))  
        pthread_cond_wait(  
            &readersQ, &m);  
    readers++;  
    pthread_mutex_unlock(&m);  
    /* read */  
    pthread_mutex_lock(&m);  
    if (--readers == 0)  
        pthread_cond_signal(  
            &writersQ);  
    pthread_mutex_unlock(&m);  
}
```

⇒ exactly the same as before!




## Improved Writer

```
writer( ) {  
    pthread_mutex_lock(&m);  
    writers++;  
    while (!((readers == 0) &&  
            (active_writers == 0))) {  
        pthread_cond_wait(&writersQ, &m);  
    }  
    active_writers++;  
    pthread_mutex_unlock(&m);  
    /* write */  
    pthread_mutex_lock(&m);  
    writers--;  
    active_writers--;  
    if (writers > 0)  
        pthread_cond_signal(&writersQ);  
    else  
        pthread_cond_broadcast(&readersQ);  
    pthread_mutex_unlock(&m);  
}
```




# New, From POSIX!

```
int pthread_rwlock_init(  
    pthread_rwlock_t *lock,  
    pthread_rwlockattr_t *att);  
int pthread_rwlock_destroy(  
    pthread_rwlock_t *lock);  
int pthread_rwlock_rdlock(  
    pthread_rwlock_t *lock);  
int pthread_rwlock_wrlock(  
    pthread_rwlock_t *lock);  
int pthread_rwlock_tryrdlock(  
    pthread_rwlock_t *lock);  
int pthread_rwlock_trywrlock(  
    pthread_rwlock_t *lock);  
int pthread_timedrwlock_rdlock(  
    pthread_rwlock_t *lock, struct timespec *ts);  
int pthread_timedrwlock_wrlock(  
    pthread_rwlock_t *lock, struct timespec *ts);  
int pthread_rwlock_unlock(  
    pthread_rwlock_t *lock);
```



# Barriers




➡ When a thread reaches a barrier, it must stop (do nothing) and simply wait for other threads to arrive at the same barrier

— when all the threads that were suppose to arrive at the barrier have all arrived at the barrier, they are all given the signal to proceed forward

○ the barrier is then reset

➡ Ex: fork/join (fork to create parallel execution)






## A Solution?

```
int count = 0;
pthread_mutex_t m;
pthread_cond_t BarrierQueue;
void barrier() {
    pthread_mutex_lock(&m);
    if (++count < n) {
        pthread_cond_wait(&BarrierQueue, &m);
    } else {
        count = 0;
        pthread_cond_broadcast(&BarrierQueue);
    }
    pthread_mutex_unlock(&m);
}
```

- the idea here is to have the last thread broadcast the condition while all the other threads are blocked at waiting for the condition to be signaled
- as it turns out, `pthread_cond_wait()` might return spontaneously, so this won't work


➤ [http://pubs.opengroup.org/onlinepubs/009604599/functions/pthread\\_cond\\_signal.html](http://pubs.opengroup.org/onlinepubs/009604599/functions/pthread_cond_signal.html)



# A Solution?

```
int count = 0;
pthread_mutex_t m;
pthread_cond_t BarrierQueue;
void barrier() {
    pthread_mutex_lock(&m);
    if (++count < n) {
        while (count < n)
            pthread_cond_wait(&BarrierQueue, &m);
    } else {
        pthread_cond_broadcast(&BarrierQueue);
        count = 0;
    }
    pthread_mutex_unlock(&m);
}
```

- if the  $n^{\text{th}}$  thread wakes up all the other blocked threads, not all these threads will see `count == n`
  - moving `count = 0` around won't help




# Barrier in POSIX Threads

```
int count = 0;
pthread_mutex_t m;
pthread_cond_t BarrierQueue;

void barrier() {
    pthread_mutex_lock(&m);
    if (++count < number) {
        int my_generation = generation;
        while(my_generation == generation)
            pthread_cond_wait(&waitQ, &m);
    } else {
        count = 0;
        generation++;
        pthread_cond_broadcast(&waitQ);
    }
    pthread_mutex_unlock(&m);
}
```

- don't use count in the guard since its problematic!
- introduce a new guard



# More From POSIX!

```
int pthread_barrier_init(  
    pthread_barrier_t *barrier,  
    pthread_barrierattr_t *attr,  
    unsigned int count);  
int pthread_barrier_destroy(  
    pthread_barrier_t *barrier);  
int pthread_barrier_wait(  
    pthread_barrier_t *barrier);
```

