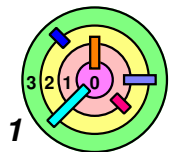


Housekeeping (Lecture 10 - 9/30/2013)

- ➡ Warmup #2 due at 11:45pm on Friday, 10/4/2013
 - if you have code from a previous semester, be very careful and **not copy any code from it**
 - it's best if you just get rid of it
 - get started soon
- ➡ **Grading guidelines** is the **ONLY** way we will grade and we will grade on `nunki.usc.edu` in our grading account (which you don't have access to)
 - it's a good idea to run your code against the grading guidelines
- ➡ After submission, make sure you **Verify Your Submission**
- ➡ Have you installed **Ubuntu 11.10** on your laptop/desktop?
- ➡ Do you have partners for kernel assignments?
 - work with your potential partners for warmup 2
 - again, work at high level and must **not** share code

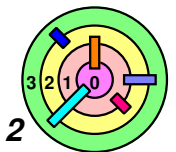


Housekeeping (Lecture 10 - 9/30/2013)



Warmup #2 note

- for the *average number of packets in Q1* statistics, only consider packets that have been served by the server
 - I contradicted myself when I said that you should only consider packets that have made it into Q2
 - that's inconsistent with the spec, I apologize
 - please stick to the spec



Buddy Systems

➡ Data Structure

1) doubly-linked list (not circular) FREE list indexed by k

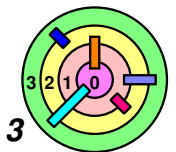
- ◆ links stored in actual blocks
- ◆ $\text{FREE}[k]$ points to first available block of size 2^k

2) each block contains

- ◆ in-use bit
- ◆ size
- ◆ NEXT and PREV links for FREE list

➡ Can get greater variety in block sizes using Fibonacci sequence of block sizes so $b_j = b_{j-1} + b_{j-2}$

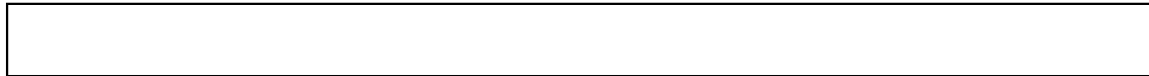
— ratio of successive block sizes is $2/3$ instead of $1/2$



Example of Buddy Algorithm

Ex: 16 locations

0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15



k free[k]

0	Ω
1	Ω
2	Ω
3	Ω
4	0

1) allocate a block of size 2

0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15



k free[k]

0	Ω
1	Ω 2
2	Ω 4
3	Ω 8
4	Ω Ω

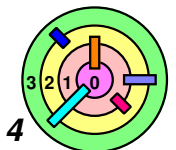
2) allocate a block of size 4

0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15



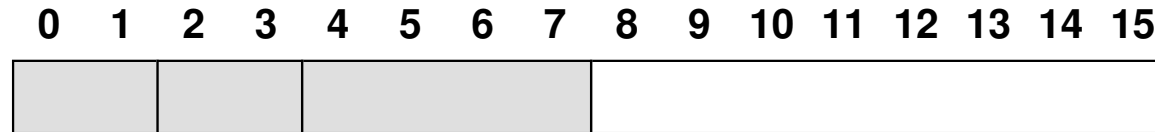
k free[k]

0	Ω
1	Ω 2
2	Ω 4 Ω
3	Ω 8
4	Ω Ω



Example of Buddy Algorithm

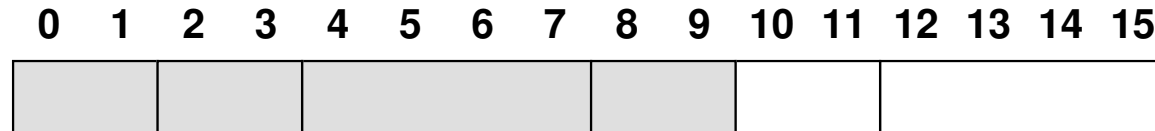
3) allocate a block of size 2



k free[k]

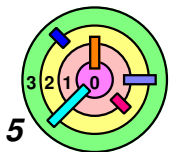
0	Ω
1	0 2 Ω
2	0 4 Ω
3	0 8
4	0 Ω

4) allocate a block of size 2



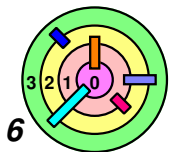
k free[k]

0	Ω
1	0 2 10
2	0 4 12
3	0 8 Ω
4	0 Ω

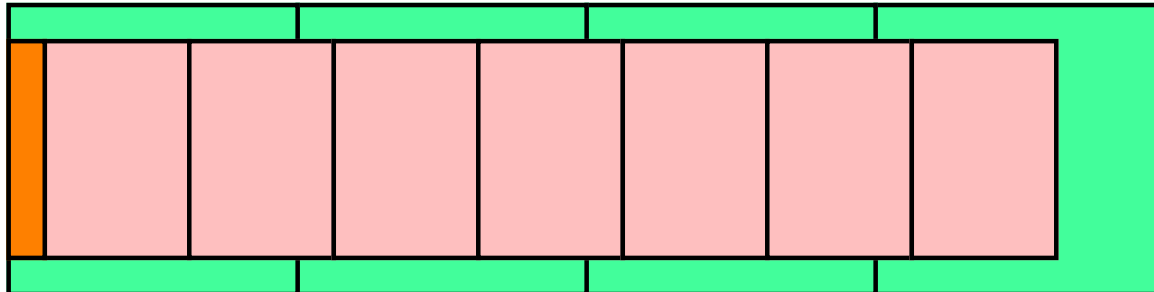
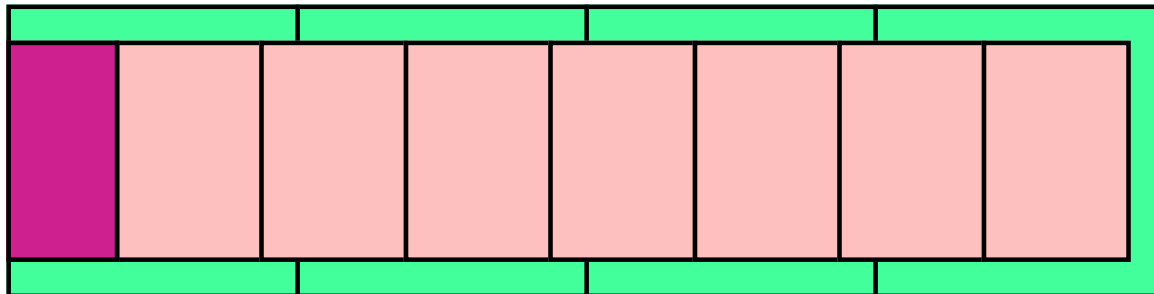
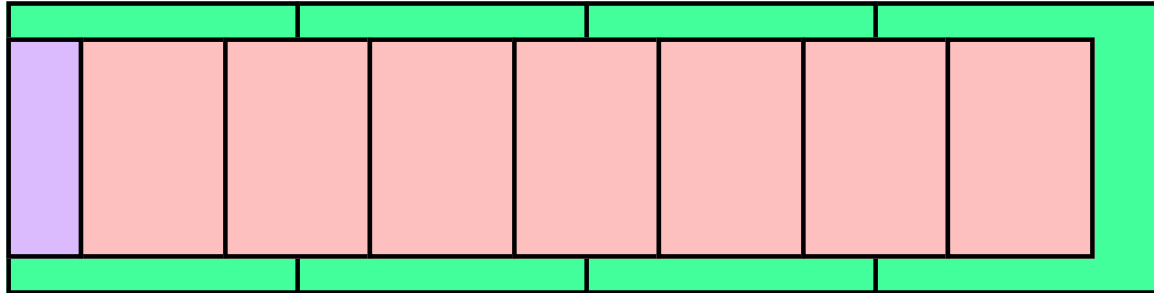


3.3 Dynamic Storage Allocation

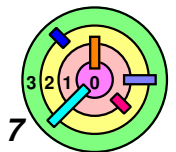
- ➡ Best-fit & First-fit Algorithms
- ➡ Buddy System
- ➡ *Slab Allocation*



Slab Allocation



— see weenix kernel code!



Slab Allocation



Objects are allocated and freed frequently

— **allocation involves**

- **finding an appropriate-sized storage**
- **initialize it**
 - ◆ **pointers need to point at the right places**
 - ◆ **may even need to initialize synchronization data structures**

— **deallocation involves**

- **tearing down the data structures**
- **freeing the storage**

— **lots of "overhead"**

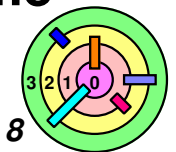


Difficulties with dynamic storage allocation

— **you cannot predict what an application will ask for**

— **but it's not true for the kernel**

- **e.g., can allocate a slab of process control blocks at a time**
 - ◆ **return one of them from a slab**



Slab Allocation

➡ *Slab Allocation*

- sets up a separate cache for each type of object to be managed
- contiguous sets of pages called *slabs*, allocated to hold objects
 - we will cover "pages" later, won't get into too much detail now

➡ Whenever a *slab* is allocated, a constructor is called to initialize all the objects it hold

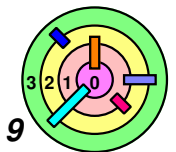
- this is where you pay for initialization, but it's done in a *batch*

➡ As *objects* are being allocated, they are taken from the set of existing slabs in the cache

- objects are considered "*preallocated*" since they have all been initialized already

➡ As *objects* are being freed, they are simply marked as free

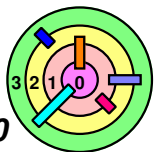
- don't have to free up storage
- when appropriate can free up an entire slab



3.4 Linking & Loading

➡ *Static Linking & Loading*

➡ **Shared Libraries**



Remember This?

```

main:
→ pushl %ebp
  movl %esp, %ebp
  pushl %esi
  pushl %edi
  subl $8, %esp
  ...
  pushl $1
  movl -12(%ebp), %eax
  pushl %eax
  call sub
  addl $8, %esp
  movl %eax, -16(%ebp)
  ...
  addl $8, %esp
  movl $0, %eax
  popl %edi
  popl %esi
  movl %ebp, %esp
  popl %ebp
  ret

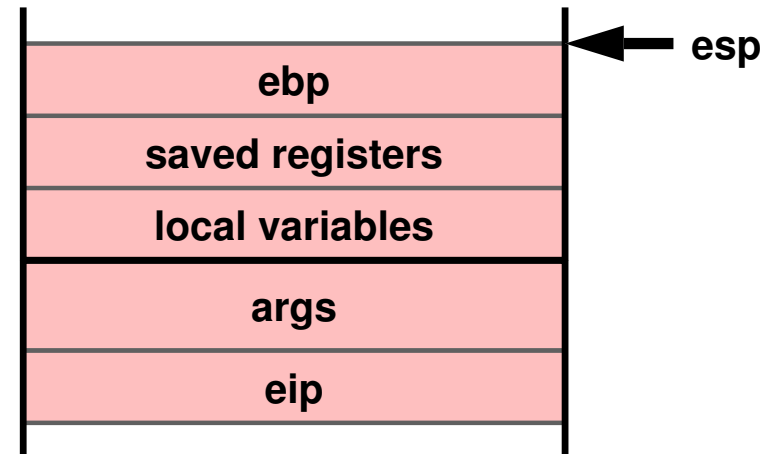
```

set up
stack frame

push args

pop args;
get result

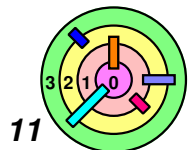
set return
value and
restore frame



```

int main() {
    int i;
    int a;
    ...
    i = sub(a, 1);
    ...
    return(0);
}

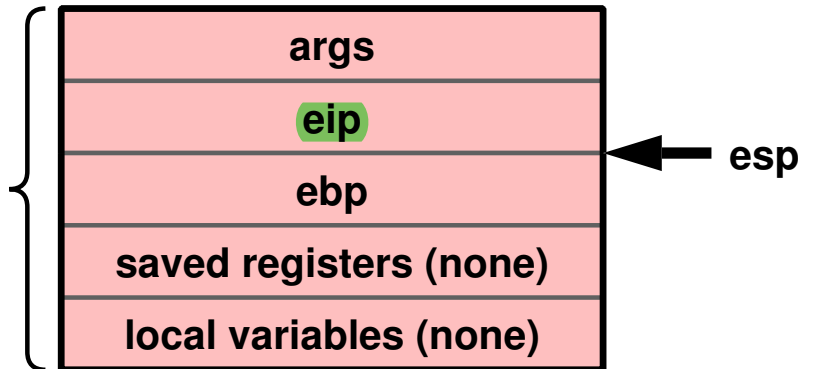
```



Something Simpler

```
int main(int argc,
        char *[]) {
    return(argc);
}
```

stack frame
of main()



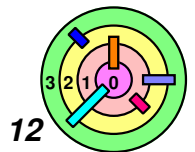
main:

```
→ pushl %ebp
   movl %esp, %ebp
   movl 8(%ebp), %eax
   movl %ebp, %esp
   popl %ebp
   ret
```

set up
stack frame

set return
value and
restore frame

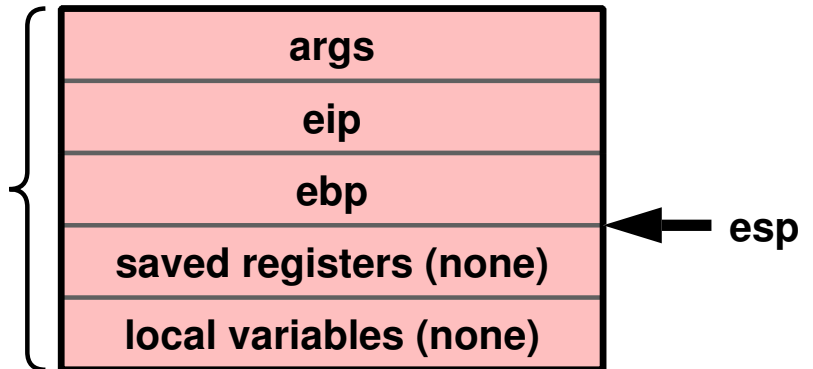
➡ Does location matter?



Something Simpler

```
int main(int argc,
        char *[]) {
    return(argc);
}
```

stack frame
of main()



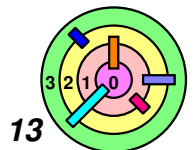
main:

```
    pushl %ebp
    → movl %esp, %ebp
    movl 8(%ebp), %eax
    movl %ebp, %esp
    popl %ebp
    ret
```

set up
stack frame

set return
value and
restore frame

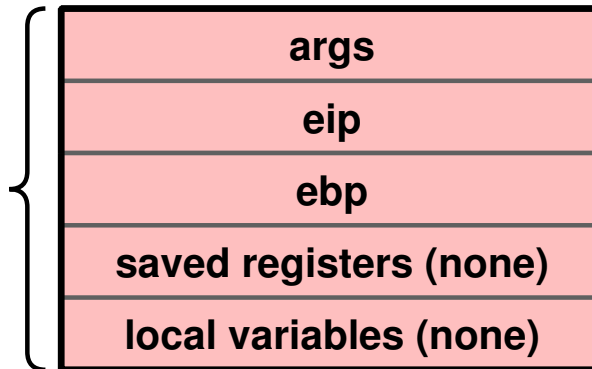
➡ Does location matter?



Something Simpler

```
int main(int argc,
        char *[]) {
    return(argc);
}
```

stack frame
of main()

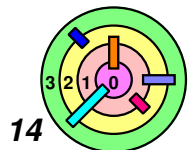


```
main:
    pushl %ebp
    movl %esp, %ebp
    → movl 8(%ebp), %eax
    movl %ebp, %esp
    popl %ebp
    ret
```

set up
stack frame

set return
value and
restore frame

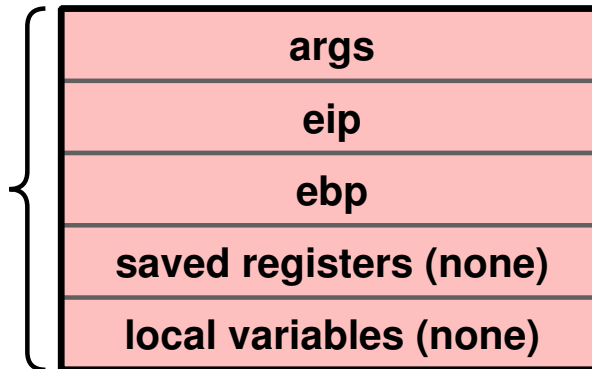
➡ Does location matter?



Something Simpler

```
int main(int argc,
        char *[]) {
    return(argc);
}
```

stack frame
of main()



main:

pushl %ebp

movl %esp, %ebp

movl 8(%ebp), %eax

→ movl %ebp, %esp

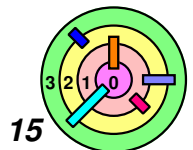
popl %ebp

ret

set up
stack frame

set return
value and
restore frame

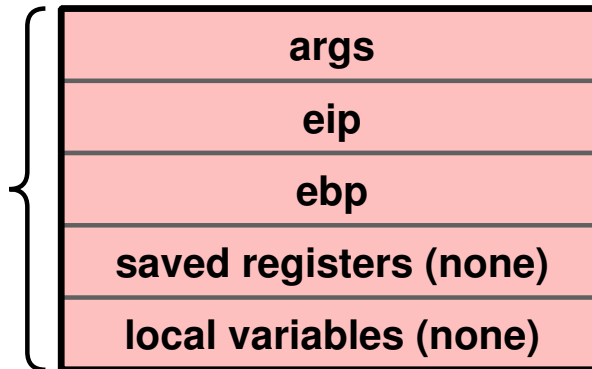
➡ Does location matter?



Something Simpler

```
int main(int argc,
        char *[]) {
    return(argc);
}
```

stack frame
of main()



ebp,
esp

main:

pushl %ebp

movl %esp, %ebp

movl 8(%ebp), %eax

movl %ebp, %esp

→ popl %ebp

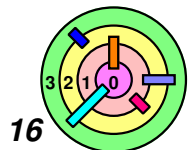
ret

set up
stack frame

set return
value and
restore frame



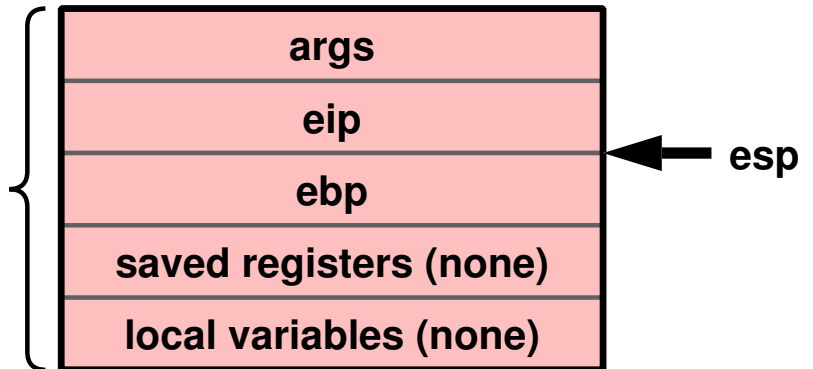
Does location matter?



Something Simpler

```
int main(int argc,
        char *[]) {
    return(argc);
}
```

stack frame
of main()



main:

pushl %ebp

movl %esp, %ebp

movl 8(%ebp), %eax

movl %ebp, %esp

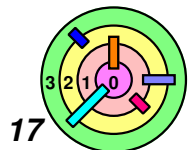
popl %ebp

→ ret

} set up
stack frame

} set return
value and
restore frame

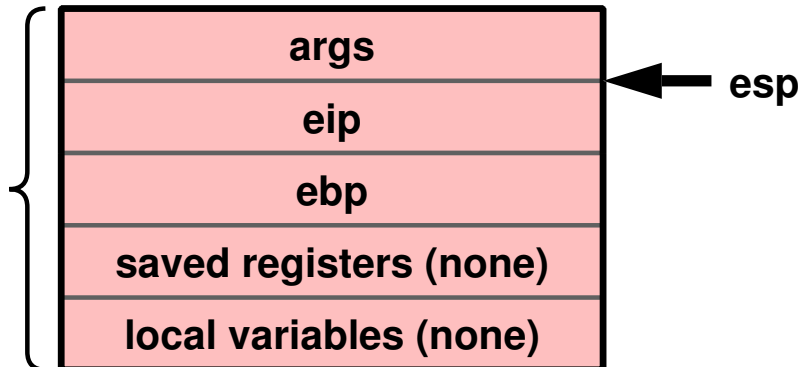
➡ Does location matter?



Something Simpler

```
int main(int argc,
        char *[]) {
    return(argc);
}
```

stack frame
of main()

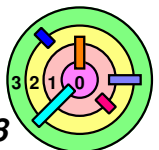


```
main:
    pushl %ebp          } set up
    movl %esp, %ebp     } stack frame
    movl 8(%ebp), %eax
    movl %ebp, %esp
    popl %ebp           } set return
    ret                 } value and
                        } restore frame
```



Does location matter?

- if everything can be accessed relative to the *frame pointer*, then you don't need to know the actual address of an object
 - just use relative-addresses



Location Matters ...

```
int X = 6;
int *aX = &X;

int main( ) {
    void subr(int);
    int y = X;
    subr(y);
    return(0);
}

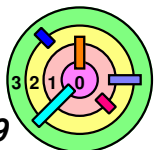
void subr(int i) {
    printf("i = %d\n", i);
}
```



Why does it matter here?

— need to put the address of x into aX

- what's the address of x?
- remember, both x and aX are in the *data segment*
- *who* would put the actual value into aX?



Coping



Relocation

- modify internal references in memory depending on where module is expected to be *loaded*
 - one of the *exec* system calls loads a program into memory
 - ◆ everything is laid out carefully in memory
- modules requiring relocation are said to be *relocatable*
- the act of modifying such a module to *resolve these references* is called *relocation*
- the program that performs relocation is called a *linker*



Two main functions of a *linker*

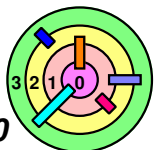
- 1) *relocation*
- 2) *symbol resolution*

textbook
is wrong



A *loader* loads a program into memory

- a "relocating loader" may perform additional relocation



A Slight Revision

```
extern int X;
int *aX = &X;

int main( ) {
    void subr(int);
    int y = *aX;
    subr(y);
    return(0);
}
```

main.c

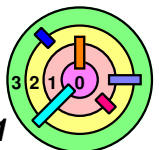
```
#include <stdio.h>
int X;

void subr(int i) {
    printf("i = %d\n", i);
}
```

subr.c

```
% gcc -o prog main.c subr.c
```

- main.c is compiled into main.o
- subr.c is compiled into subr.o
- ld is then invoked to combine them into prog
 - ld knows where to find printf()
 - prog can be loaded into memory through one of the **exec** system calls



A Slight Revision

```
extern int X;
int *aX = &X;

int main( ) {
    void subr(int);
    int y = *aX;
    subr(y);
    return(0);
}
```

main.c

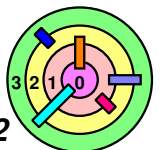
```
#include <stdio.h>
int X;

void subr(int i) {
    printf("i = %d\n", i);
}
```

subr.c

```
% gcc -o prog main.c subr.c
```

- how does ld decides what needs to be done?
- main.c contains undefined references to X and subr ()
 - instructions for doing this are provided in main.o
- later on, when the actual locations for these are determined,
 - ld will modify them when main.o is copied into prog



A Slight Revision

```
extern int X;
int *aX = &X;

int main( ) {
    void subr(int);
    int y = *aX;
    subr(y);
    return(0);
}
```

main.c

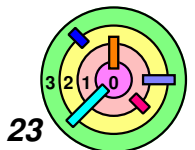
```
#include <stdio.h>
int X;

void subr(int i) {
    printf("i = %d\n", i);
}
```

subr.c

```
% gcc -o prog main.c subr.c
```

- main.o must contains a list of external symbols, along with their types, and instructions for updating this code

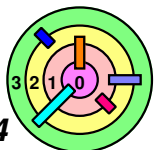


main.s

Offset	Op	Arg
0:	.data	; what follows is initialized data
0:	.globl	aX ; aX is global: it may be used ; by others
0:	aX:	
0:	.long	X
4:		
0:	.text	; offset restarts; what follows is ; text (read-only code)
0:	.globl	main
0:	main:	
0:	pushl	%ebp ; save the frame pointer
1:	movl	%esp,%ebp ; point to
3:	subl	\$4,%esp ; make space
6:	movl	aX,%eax ; put conten
11:	movl	(%eax),%eax ; put *X
13:	movl	%eax,-4(%ebp) ; stor
16:	pushl	-4(%ebp) ; push y on
19:	call	subr
24:	addl	\$4,%esp ; remove y f
27:	movl	\$0,%eax ; set return
31:	movl	%ebp, %esp ; restore
33:	popl	%ebp ; pop frame pointer
35:	ret	

```
extern int X;
int *aX = &X;

int main( ) {
    void subr(int);
    int y = *aX;
    subr(y);
    return(0);
}
```



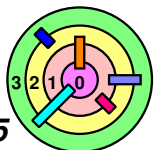
main.s

Offset	Op	Arg
0:	.data	; what follows is initialized data
0:	.globl	aX ; aX is global: it may be used ; by others
0:	aX:	
0:	.long	X
4:		
0:	.text	; offset restarts; w ; text (read-only co
0:	.globl	main
0:	main:	
0:	pushl	%ebp ; save the frame pointer
1:	movl	%esp,%ebp ; point to
3:	subl	\$4,%esp ; make space
6:	movl	aX,%eax ; put conten
11:	movl	(%eax),%eax ; put *X
13:	movl	%eax,-4(%ebp) ; stor
16:	pushl	-4(%ebp) ; push y on
19:	call	subr
24:	addl	\$4,%esp ; remove y f
27:	movl	\$0,%eax ; set return
31:	movl	%ebp, %esp ; restore
33:	popl	%ebp ; pop frame pointer
35:	ret	

What follows goes into the **data** segment

```
extern int X;
int *aX = &X;

int main( ) {
    void subr(int);
    int y = *aX;
    subr(y);
    return(0);
}
```



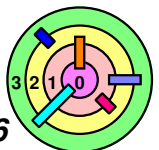
main.s

Offset	Op	Arg
0:	.data	; what follows is initialized data
0:	.globl	aX ; aX is global: it may be used ; by others
0:	aX:	
0:	.long	X
4:		
→ 0:	.text	; offset restarts; w ; text (read-only co
0:	.globl	main
0:	main:	
0:	pushl	%ebp ; save the frame pointer
1:	movl	%esp,%ebp ; point to
3:	subl	\$4,%esp ; make space
6:	movl	aX,%eax ; put conten
11:	movl	(%eax),%eax ; put *X
13:	movl	%eax,-4(%ebp) ; stor
16:	pushl	-4(%ebp) ; push y on
19:	call	subr
24:	addl	\$4,%esp ; remove y f
27:	movl	\$0,%eax ; set return
31:	movl	%ebp, %esp ; restore
33:	popl	%ebp ; pop frame pointer
35:	ret	

What follows goes into the **text** segment

```
extern int X;
int *aX = &X;

int main( ) {
    void subr(int);
    int y = *aX;
    subr(y);
    return(0);
}
```



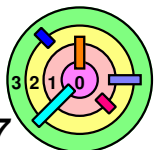
main.s

Offset	Op	Arg
0:	.data	; what follows is initialized data
0:	.globl	aX ; aX is global: it may be used ; by others
0:	aX:	
0:	.long	X
4:		
0:	.text	; offset restarts; w ; text (read-only co
0:	.globl	main
0:	main:	
0:	pushl	%ebp ; save the frame pointer
1:	movl	%esp,%ebp ; point to
3:	subl	\$4,%esp ; make space
6:	movl	aX,%eax ; put conten
11:	movl	(%eax),%eax ; put *X
13:	movl	%eax,-4(%ebp) ; stor
16:	pushl	-4(%ebp) ; push y on
19:	call	subr
24:	addl	\$4,%esp ; remove y f
27:	movl	\$0,%eax ; set return
31:	movl	%ebp, %esp ; restore
33:	popl	%ebp ; pop frame pointer
35:	ret	

offset got restarted because
segments are relocatable

```
extern int X;
int *aX = &X;

int main( ) {
    void subr(int);
    int y = *aX;
    subr(y);
    return(0);
}
```



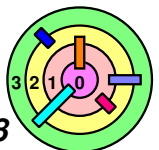
main.s

Offset	Op	Arg
0:	.data	; what follows is initialized data
→ 0:	.globl	aX ; aX is global: it may be used ; by others
0:	aX:	
0:	.long	X
4:		
0:	.text	; offset restarts; w ; text (read-only co
→ 0:	.globl	main
0:	main:	
0:	pushl	%ebp ; save the frame pointer
1:	movl	%esp,%ebp ; point to
3:	subl	\$4,%esp ; make space
6:	movl	aX,%eax ; put conten
11:	movl	(%eax),%eax ; put *X
13:	movl	%eax,-4(%ebp) ; stor
16:	pushl	-4(%ebp) ; push y on
19:	call	subr
24:	addl	\$4,%esp ; remove y f
27:	movl	\$0,%eax ; set return
31:	movl	%ebp, %esp ; restore
33:	popl	%ebp ; pop frame pointer
35:	ret	

`.global` directive means that the symbol mentioned is defined *here* and is *exported*
 ➡ i.e., can be referenced by other modules
 ➡ `aX` and `main` are global

```
extern int X;
int *aX = &X;

int main( ) {
    void subr(int);
    int y = *aX;
    subr(y);
    return(0);
}
```



main.s

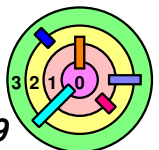
Offset	Op	Arg
0:	.data	; what follows is initialized data
0:	.globl	aX ; aX is global: it may be used ; by others
0:	aX:	
0:	.long	X
4:		
0:	.text	; offset restarts; w ; text (read-only co
0:	.globl	main
0:	main:	
0:	pushl	%ebp ; save the frame pointer
1:	movl	%esp,%ebp ; point to
3:	subl	\$4,%esp ; make space
6:	movl	aX,%eax ; put conten
11:	movl	(%eax),%eax ; put *X
13:	movl	%eax,-4(%ebp) ; stor
16:	pushl	-4(%ebp) ; push y on
19:	call	subr
24:	addl	\$4,%esp ; remove y f
27:	movl	\$0,%eax ; set return
31:	movl	%ebp, %esp ; restore
33:	popl	%ebp ; pop frame pointer
35:	ret	

aX is 4 bytes long and put
the value of x here

☞ x will remain unresolved

```
extern int X;
int *aX = &X;

int main( ) {
    void subr(int);
    int y = *aX;
    subr(y);
    return(0);
}
```



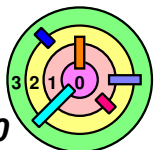
main.s

Offset	Op	Arg
0:	.data	; what follows is initialized data
0:	.globl	aX ; aX is global: it may be used ; by others
0:	aX:	
0:	.long	X
4:		
0:	.text	; offset restarts; w ; text (read-only co
0:	.globl	main
0:	main:	
0:	pushl	%ebp ; save the frame pointer
1:	movl	%esp,%ebp ; point to
3:	subl	\$4,%esp ; make space
6:	movl	aX,%eax ; put conten
11:	movl	(%eax),%eax ; put *X
13:	movl	%eax,-4(%ebp) ; stor
16:	pushl	-4(%ebp) ; push y on
19:	call	subr
24:	addl	\$4,%esp ; remove y f
27:	movl	\$0,%eax ; set return
31:	movl	%ebp, %esp ; restore
33:	popl	%ebp ; pop frame pointer
35:	ret	

these 3 places require
relocation

```
extern int X;
int *aX = &X;

int main( ) {
    void subr(int);
    int y = *aX;
    subr(y);
    return(0);
}
```



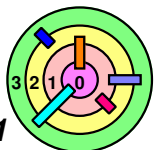
main.s

Offset	Op	Arg
0:	.data	; what follows is initialized data
0:	.globl	aX ; aX is global: it may be used ; by others
0:	aX:	
0:	.long	X
4:		
0:	.text	; offset restarts; w ; text (read-only co
0:	.globl	main
0:	main:	
0:	pushl	%ebp ; save the frame pointer
1:	movl	%esp,%ebp ; point to
3:	subl	\$4,%esp ; make space
6:	movl	aX,%eax ; put conten
11:	movl	(%eax),%eax ; put *X
13:	movl	%eax,-4(%ebp) ; stor
16:	pushl	-4(%ebp) ; push y on
19:	call	subr
24:	addl	\$4,%esp ; remove y f
27:	movl	\$0,%eax ; set return
31:	movl	%ebp, %esp ; restore
33:	popl	%ebp ; pop frame pointer
35:	ret	

this call is a PC-relative call
 ➡ what's stored at offset 20
 is not the absolute
 address of subr, but a
 relative address

```
extern int X;
int *aX = &X;

int main( ) {
    void subr(int);
    int y = *aX;
    subr(y);
    return(0);
}
```

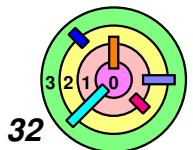


subr.s

Offset	Op	Arg
0:	.data	; what follows is initialized data
0:	printfarg:	
0:	.string	"i = %d\n"
8:		
0:	.comm	X, 4 ; 4 bytes in BSS is required ; for global X
4:		
0:	.text	; offset restarts; what follows is ; text (read-only code)
0:	.globl	subr
0:	subr:	
0:	pushl	%ebp ; save the frame pointer
1:	movl	%esp, %ebp ; point to new frame
3:	pushl	8(%ebp) ; push i onto stack
6:	pushl	\$printfarg ; push address of printfarg ; onto stack
11:	call	printf
16:	addl	\$8, %esp ; pop arguments
19:	movl	%ebp, %esp ; restore stack pointer
21:	popl	%ebp ; pop frame pointer
23:	ret	

```
#include <stdio.h>
int X;

void subr(int i) {
    printf("i = %d\n", i);
}
```



subr.s

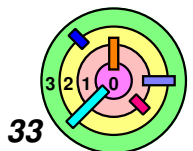
Offset	Op	Arg
0:	.data	; what follows is initialized data
0:	printfarg:	
0:	.string	"i = %d\n"
8:		
0:	.comm	X, 4 ; 4 bytes in BSS ; for global X
4:		
0:	.text	; offset restarts; w ; text (read-only co
0:	.globl	subr
0:	subr:	
0:	pushl	%ebp ; save the fra
1:	movl	%esp, %ebp ; point t
3:	pushl	8(%ebp) ; push i on
6:	pushl	\$printfarg ; push ad ; onto st
11:	call	printf
16:	addl	\$8, %esp ; pop argum
19:	movl	%ebp, %esp ; restore stack pointer
21:	popl	%ebp ; pop frame pointer
23:	ret	

this is how you create a string constant

- ▢ this one is 8 bytes long
- ▢ and local to this module (since it's not global)

```
#include <stdio.h>
int X;
```

```
void subr(int i) {
    printf("i = %d\n", i);
}
```



subr.s

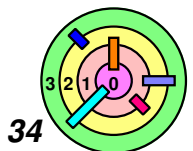
Offset	Op	Arg
0:	.data	; what follows is initialized data
0:	printfarg:	
0:	.string	"i = %d\n"
8:		
0:	.comm	X, 4 ; 4 bytes in BSS ; for global X
4:		
0:	.text	; offset restarts; w ; text (read-only co
0:	.globl	subr
0:	subr:	
0:	pushl	%ebp ; save the fra
1:	movl	%esp, %ebp ; point t
3:	pushl	8(%ebp) ; push i on
6:	pushl	\$printfarg ; push ad ; onto st
11:	call	printf
16:	addl	\$8, %esp ; pop argum
19:	movl	%ebp, %esp ; restore stack pointer
21:	popl	%ebp ; pop frame pointer
23:	ret	

this is how you create a string constant

- ▢ this one is 8 bytes long
- ▢ and local to this module (since it's not global)
- ▢ it is used here

```
#include <stdio.h>
int X;
```

```
void subr(int i) {
    printf("i = %d\n", i);
}
```



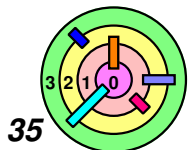
subr.s

Offset	Op	Arg
0:	.data	; what follows is initialized data
0:	printfarg:	
0:	.string	"i = %d\n"
8:		
0:	.comm	X, 4 ; 4 bytes in BSS ; for global X
4:		
0:	.text	; offset restarts; w ; text (read-only co
0:	.globl	subr
0:	subr:	
0:	pushl	%ebp ; save the fra
1:	movl	%esp, %ebp ; point t
3:	pushl	8(%ebp) ; push i on
6:	pushl	\$printfarg ; push ad ; onto st
11:	call	printf
16:	addl	\$8, %esp ; pop argum
19:	movl	%ebp, %esp ; restore stack pointer
21:	popl	%ebp ; pop frame pointer
23:	ret	

4 bytes is required in the **bss** segment for this global variable

```
#include <stdio.h>
int X;

void subr(int i) {
    printf("i = %d\n", i);
}
```



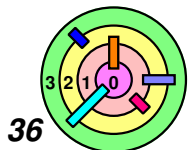
subr.s

Offset	Op	Arg
0:	.data	; what follows is initialized data
0:	printfarg:	
0:	.string	"i = %d\n"
8:		
0:	.comm	X, 4 ; 4 bytes in BSS ; for global X
4:		
0:	.text	; offset restarts; w ; text (read-only co
→ 0:	.globl	subr
0:	subr:	
0:	pushl	%ebp ; save the fra
1:	movl	%esp, %ebp ; point t
3:	pushl	8(%ebp) ; push i on
6:	pushl	\$printfarg ; push ad ; onto st
11:	call	printf
16:	addl	\$8, %esp ; pop argum
19:	movl	%ebp, %esp ; restore stack pointer
21:	popl	%ebp ; pop frame pointer
23:	ret	

subr is a global symbol
exported from here

```
#include <stdio.h>
int X;

void subr(int i) {
    printf("i = %d\n", i);
}
```



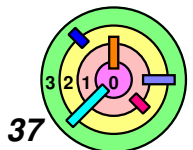
subr.s

Offset	Op	Arg
0:	.data	; what follows is initialized data
0:	printfarg:	
0:	.string	"i = %d\n"
8:		
0:	.comm	X, 4 ; 4 bytes in BSS ; for global X
4:		
0:	.text	; offset restarts; w ; text (read-only co
0:	.globl	subr
0:	subr:	
0:	pushl	%ebp ; save the fra
1:	movl	%esp, %ebp ; point t
3:	pushl	8(%ebp) ; push i on
→ 6:	pushl	\$printfarg ; push ad ; onto st
→ 11:	call	printf
16:	addl	\$8, %esp ; pop argum
19:	movl	%ebp, %esp ; restore stack pointer
21:	popl	%ebp ; pop frame pointer
23:	ret	

relocation is required for
printf and printfarg

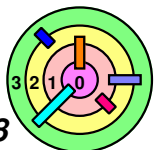
```
#include <stdio.h>
int X;

void subr(int i) {
    printf("i = %d\n", i);
}
```



Object Files

- ➡ An object file describes what's in the data, bss, and text segments in separate sections
- ➡ Along with each section is a list of:
 - ▬ global symbols
 - ▬ undefined symbols
 - ▬ instructions for relocation
 - these instructions indicate
 - ◆ which locations within the section must be modified
 - ◆ which symbol's value is used to modify the location
 - a symbol's value is the address that is ultimately determined for it
 - typically, this address is added to the location being modified
- ➡ To inspect an object file on Unix
 - ▬ **nm** - list symbols from object files
 - ▬ **objdump** - display information from object files



subr.s

Offset	Op	Arg
0:	.data	; what follows is initialized data
0:	printfarg:	
0:	.string	"i = %d\n"
8:		
0:	.comm	X, 4 ; 4 bytes in BSS ; for global X
4:		
0:	.text	; offset restarts; w ; text (read-only co
0:	.globl	subr
0:	subr:	
0:	pushl	%ebp ; save the fra
1:	movl	%esp, %ebp ; point t
3:	pushl	8(%ebp) ; push i onto stack
→ 6:	pushl	\$printfarg ; push address of string ; onto stack
→ 11:	call	printf
16:	addl	\$8, %esp ; pop arguments from stack
19:	movl	%ebp, %esp ; restore stack pointer
21:	popl	%ebp ; pop frame pointer
23:	ret	

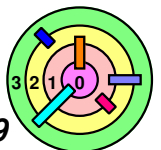
relocation is required for:

▢ printfarg at offset 7

▢ printf at offset 12

```
#include <stdio.h>
int X;
```

```
void subr(int i) {
    printf("i = %d\n", i);
}
```



subr.o

Data:

Size: 8

Contents: "i = %d\n"

bss:

Size: 4

Global: X, offset 0

Text:

Size: 24

Global: subr, offset 0

Undefined: printf

Relocation:

offset 7, size 4, value: addr of printfarg

offset 12, size 4, value: PC-relative addr of
printf

Contents: [machine instructions]

relocation is required for:

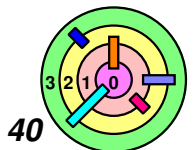
▢ printfarg at offset 7

▢ printf at offset 12

```
#include <stdio.h>
```

```
int X;
```

```
void subr(int i) {  
    printf("i = %d\n", i);  
}
```



subr.o

Data:

Size: 8

Contents: "i = %d\n"

bss:

Size: 4

→ Global: X, offset 0

Text:

Size: 24

→ Global: subr, offset 0

Undefined: printf

Relocation:

offset 7, size 4, value: addr of printfarg

offset 12, size 4, value: PC-relative addr of
printf

Contents: [machine instructions]

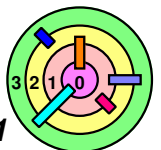
X and subr are exported

→ needed in main.o

```
#include <stdio.h>
```

```
int X;
```

```
void subr(int i) {  
    printf("i = %d\n", i);  
}
```



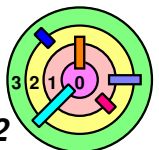
main.s

Offset	Op	Arg
0:	.data	; what follows is initialized data
0:	.globl	aX ; aX is global: ; by others
0:	aX:	
0:	.long	X
4:		
0:	.text	; offset restarts; w ; text (read-only co
0:	.globl	main
0:	main:	
0:	pushl	%ebp ; save the fram
1:	movl	%esp,%ebp ; point to
3:	subl	\$4,%esp ; make space
6:	movl	aX,%eax ; put conten
11:	movl	(%eax),%eax ; put *X }
13:	movl	%eax,-4(%ebp) ; stor
16:	pushl	-4(%ebp) ; push y onto stack
19:	call	subr
24:	addl	\$4,%esp ; remove y from stack
27:	movl	\$0,%eax ; set return value to 0
31:	movl	%ebp, %esp ; restore stack pointer
33:	popl	%ebp ; pop frame pointer
35:	ret	

these 2 places remained unresolved

```
extern int X;
int *aX = &X;

int main( ) {
    void subr(int);
    int y = *aX;
    subr(y);
    return(0);
}
```



main.o

Data:

Size: 4

Global: aX, offset 0

Undefined: X

Relocation: offset 0, size 4, value:

Contents: 0x00000000

these 2 places remained
unresolved

→ they are noted in main.o

bss:

Size: 0

Text:

Size: 36

Global: main, offset 0

Undefined: subr

Relocation:

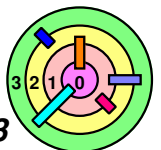
offset 7, size 4, value: addr of aX

offset 20, size 4, value: PC-relative
addr of subr

Contents: [machine instructions]

```
extern int X;
int *aX = &X;

int main( ) {
    void subr(int);
    int y = *aX;
    subr(y);
    return(0);
}
```



subr.o

Data:

Size: 8

Contents: "i = %d\n"

bss:

Size: 4

Global: X, offset 0

Text:

Size: 24

Global: subr, offset 0

Undefined: printf

Relocation:

offset 7, size 4, value: addr of printfarg

offset 12, size 4, value: PC-relative addr of
printf

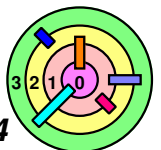
Contents: [machine instructions]

printf remained unresolved

```
#include <stdio.h>
```

```
int X;
```

```
void subr(int i) {  
    printf("i = %d\n", i);  
}
```



printf.o

Data:

Size: 1024

Global: StandardFiles

Contents: ...

bss:

Size: 256

Text:

Size: 12000

Global: printf, offset 100

...

Undefined: write

Relocation:

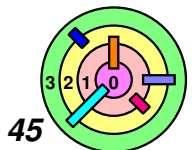
offset 211, value: addr of StandardFiles

offset 723, value: PC-relative addr of write

Contents: [machine instructions]

assume that printf.o looks like this

→ write is unresolved



write.o

Data:

Size: 0

bss:

Size: 4

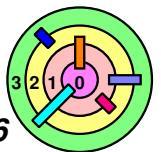
Global: errno, offset 0

Text:

Size: 16

Contents: [machine
instructions]

and write.o looks like this



startup function

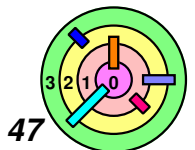
Data:
Size: 0

bss:
Size: 0

Text:
Size: 36
Undefined: main
Relocation:
offset 21, value: main
Contents: [machine
instructions]

every C program contains a
startup routine that is called
first

- it calls `main()`
- if `main()` returns, it calls `exit()`
- our example is incomplete



prog

Text

main	4096
subr	4132
printf	4156
write	16156
startup	16172

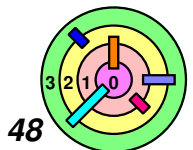
Data

aX	16384
printfargs	16388
StandardFiles	16396

BSS

X	17420
errno	17680

- ▢ this is how `ld` might set things up
- ▢ main does not start at location 0
 - first "page" is typically made inaccessible so that references to null pointers will fail (get SIGSEG)



prog

Text

main	4096
subr	4132
printf	4156
write	16156
startup	16172

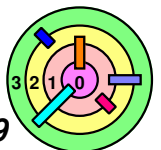
Data

aX	16384
printfargs	16388
StandardFiles	16396

BSS

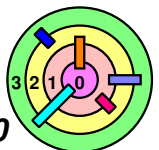
X	17420
errno	17680

- due to the use of "**pages**", the data segment needs to start at a page boundary (i.e., multiple of page size)
 - this way, the text segment can be made read-only while the data and bss segments read-write
 - here we assume that pages start at 4096, 8192, 12288, 16384, etc.



Virtual Memory Basics

- ➡ A process has, say, a 32-bit address space
 - that's 4GB of memory
- ➡ Our prog process, when it starts, only needs about 16KB for text+data+bss
 - plus more for stack
- ➡ Allocating 4GB of memory will be a huge waste
- ➡ Solution: *page table* in *virtual memory*
 - OS allocate *pages* of *physical memory* at a time
 - a page is 4KB in many systems
 - a page corresponds to physical memory that can be located (or "*mapped*") anywhere in virtual memory
 - one level of *indirection* to get to the physical memory
 - the hardware makes this transparent
- ➡ We will spend a lot of time talking about virtual memory (Ch 7)



Virtual Memory Basics

Text

main	4096
subr	4132
printf	4156
write	16156
startup	16172

Data

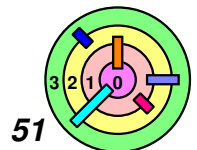
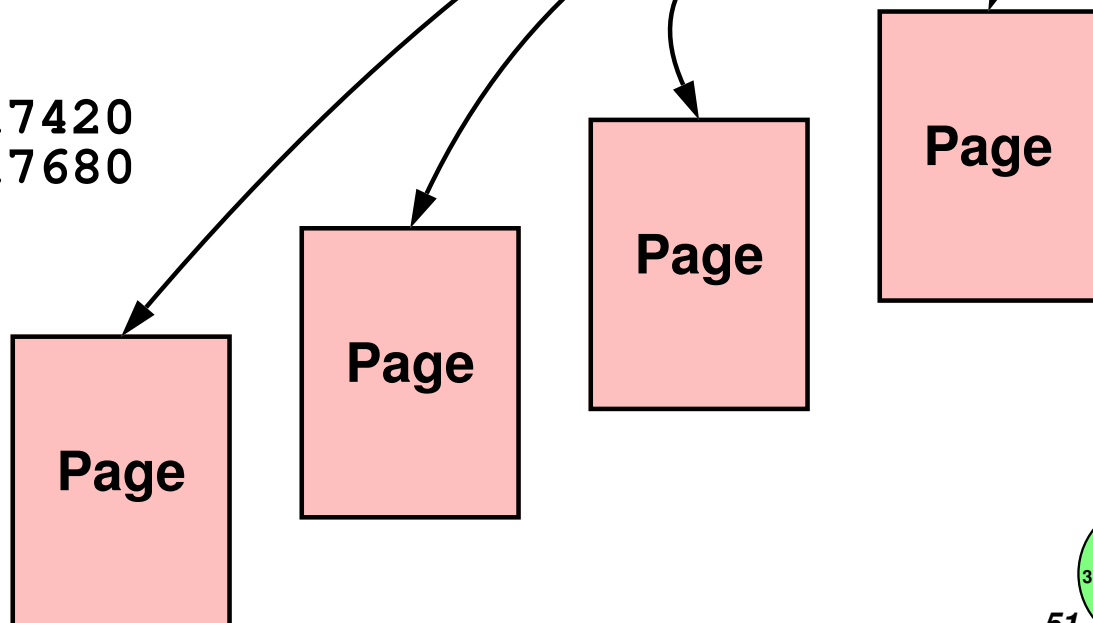
aX	16384
printfargs	16388
StandardFiles	16396

BSS

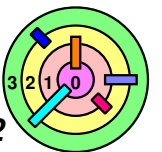
X	17420
errno	17680

Page Table

<i>Start</i>	<i>Access</i>	<i>Physical Addr</i>
0	-	-
4096	R	•
8192	R	•
12288	R	•
16384	R/W	•

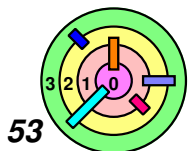


3.5 Booting



Boot

- ➡ Came from the idiomatic expression, "to pull yourself up by your bootstraps"
 - without the help of others
 - it's a difficult situation
- ➡ In OS
 - load its OS into memory
 - which kind of means that you need an OS in memory to do it
- ➡ Solution
 - load a tiny OS into memory
 - known as the *bootstrap loader*
 - then again, who loads this tiny OS into memory?
 - ◆ how about first loading a tiny bootstrap loader?

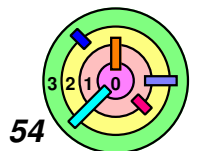


PDP-8



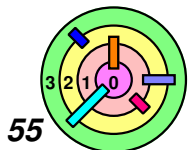
toggle
switches

- ➡ How about manually put into memory a simple bootstrap loader?
- approach taken by PDP-8
 - "toggles in" the program
 - read OS from paper tape

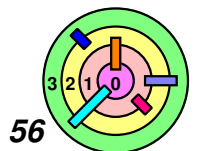


PDP-8 Boot Code

```
07756 6032 KCC
07757 6031 KSF
07760 5357 JMP .-1
07761 6036 KRB
07762 7106 CLL RTL
07763 7006 RTL
07764 7510 SPA
07765 5357 JMP 7757
07766 7006 RTL
07767 6031 KSF
07770 5367 JMP .-1
07771 6034 KRS
07772 7420 SNL
07773 3776 DCA I 7776
07774 3376 DCA 7776
07775 5356 JMP 7756
07776 0000 AND 0
07777 5301 JMP 7701
```



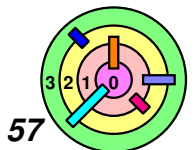
VAX-11/780



VAX-11/780 Boot

- ➡ Separate "console computer"
 - LSI-11
 - hard-wired to always run the code contained in its on-board read-only memory
 - then read boot code (i.e., the bootstrap loader) from floppy disk
 - then load OS from root directory of first file system on primary disk

- ➡ Code on floppy disk (the bootstrap loader) would handle:
 - *disk device*
 - *on-disk file system*
 - it needs the right *device driver*
 - it needs to know how the disk is setup
 - what sort of *file system* is on the disk
 - how the disk is *partitioned*
 - ◆ a disk may hold multiple and different file systems, each in a separate partition

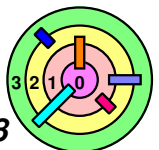


Configuring the OS



Early Unix

- OS statically linked to contain all needed device drivers
 - device drivers were statically linked to the OS
- all device-specific info included with drivers
- disk drivers contained partitioning description
- therefore, the following actions may all require compiling a new version of the OS:
 - adding a new device
 - replacing a device
 - modifying disk-partitioning information



Configuring the OS



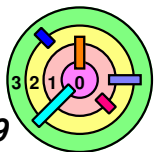
Later Unix

- OS statically linked to contain all needed device drivers
- at boot time, OS would probe to see which devices were present and discover device-specific info
- partition table in first sector of each disk



Even later Unix

- allowed device drivers to be dynamically loaded into a running system



IBM PC

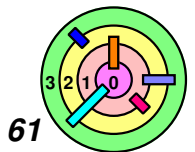


Issues



Open architecture

- although MS-DOS was distributed in binary form only
- large market for peripherals, most requiring special drivers
- how to access boot device?
- how does OS get drivers for new devices?



The Answer: BIOS



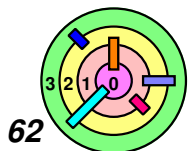
Basic Input-Output System (*BIOS*)

- code stored in read-only memory (ROM)
- configuration data in non-volatile RAM (NVRAM)
 - such as *CMOS*
- including set of boot-device names
- the BIOS provides three primary functions
 - power-on self test (*POST*)
 - ◇ so it knows *where* to load the boot program, etc. into
 - load and transfer control to boot program
 - provide *drivers* for all devices



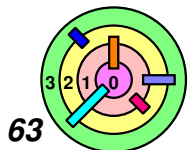
Main BIOS on motherboard

- supplied as a chip on the "motherboard"
- contains everything necessary to perform the above 3 functions
- additional BIOSes on other boards
 - provide access to additional devices

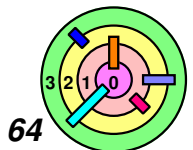
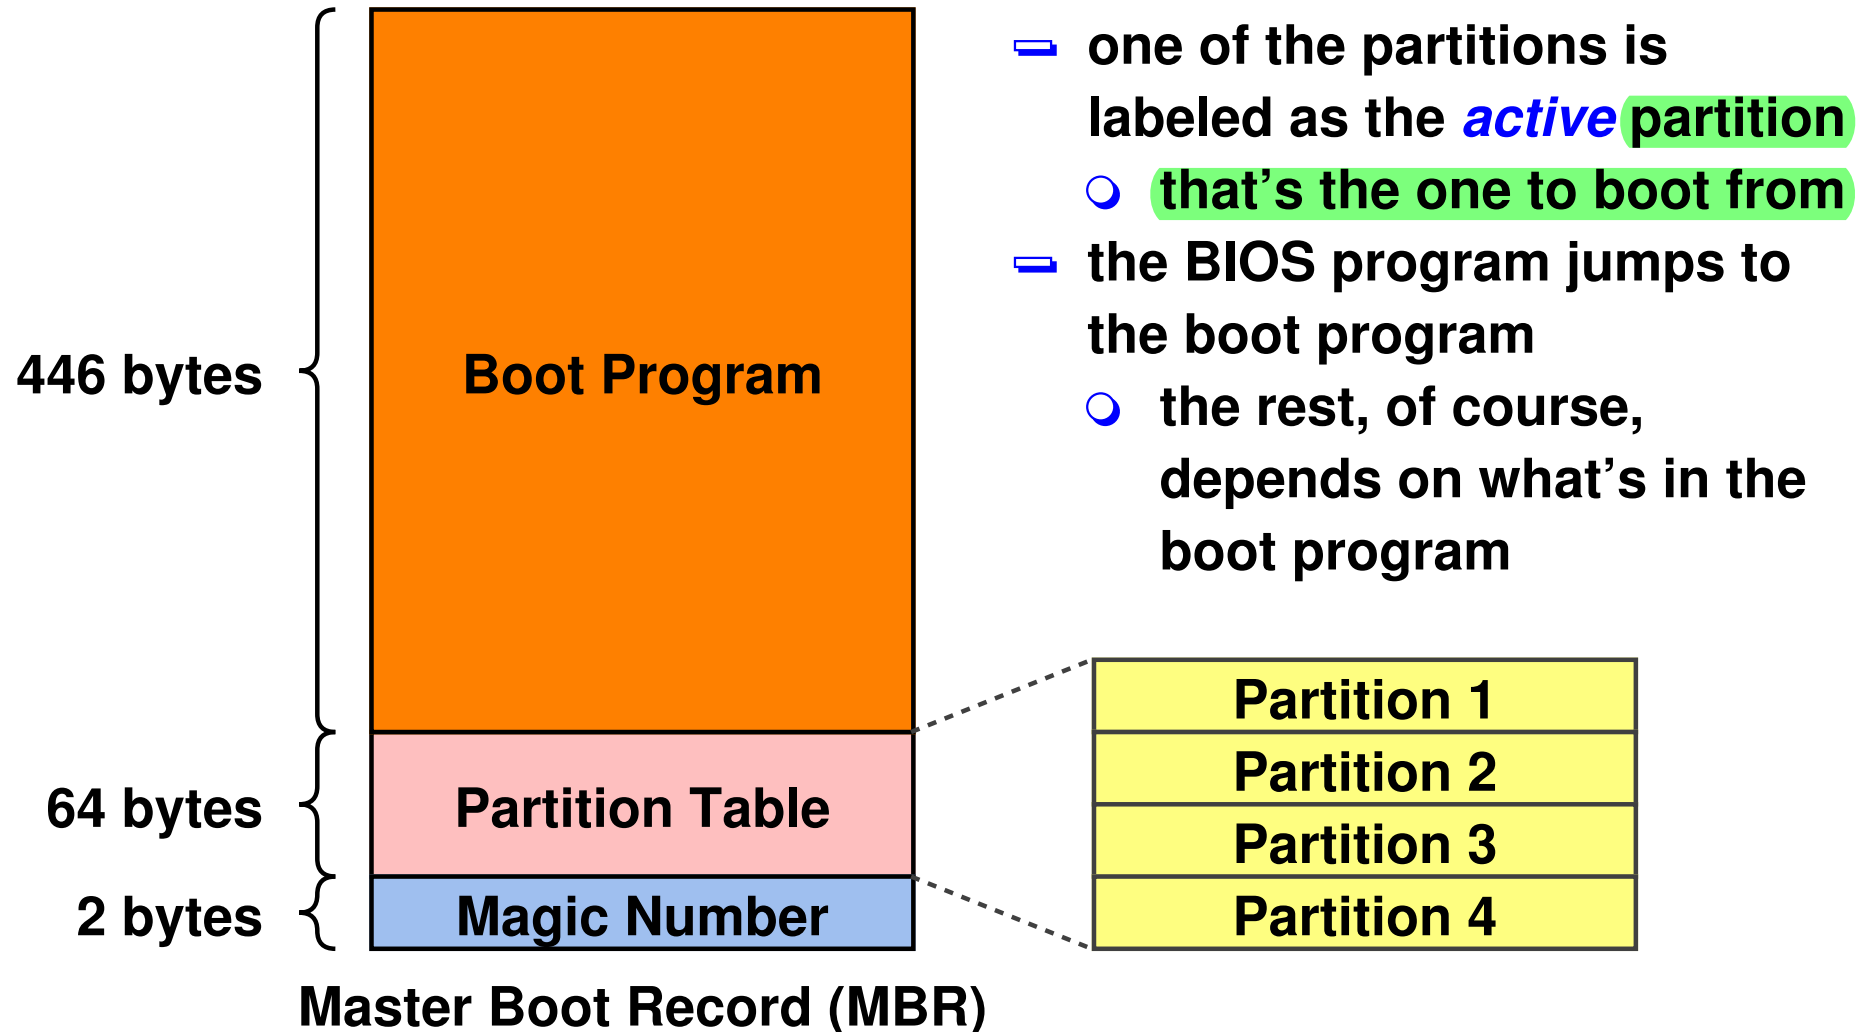


POST

- ➡ On power-on, CPU executes BIOS code
 - located in last 64KB of first megabyte of address space
 - starting at location 0xf0000
 - CPU is **hard-wired to start executing at 0xffff0** on startup
 - ◆ the last 16 bytes of this region
 - ◆ jump to POST
- ➡ POST
 - **initializes hardware**
 - **counts memory locations**
 - **by testing for working memory**
- ➡ Next step is to find a boot device
 - **the CMOS is configured with a boot order**
- ➡ Next step is to **load the Master Boot Record (*MBR*) from the first sector of the boot device**, if it's a floppy/diskette
 - or cylinder 0, head 0, section 1 of a hard disk (Ch 6)



Getting the Boot Program

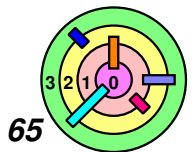


MS-DOS Boot Program



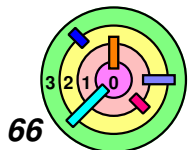
Find the active partition

- load the first sector from it
 - which contains the "volume boot program"
- pass control to that program
 - which then load the OS from that partition



Linux Booting (1)

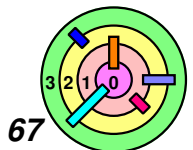
- ➡ Two stages of booting provided by one of:
- *lilo* (Linux Loader)
 - uses sector numbers of kernel image
 - therefore, must be modified if a kernel image moves
 - *grub* (Grand Unified Boot Manager)
 - understands various file systems
 - can find a kernel image given a file system *path name*
 - both allow dual (or greater) booting
 - select which system to boot from menu
 - perhaps choice of Linux or Windows
- ➡ The next step is for the kernel to *configure* itself



Linux Booting (2)

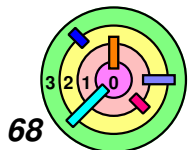
- | | | |
|--|---|---|
| assembler code
(startup_32) | { | ➡ Kernel image is compressed |
| | | ➡ step 1: set up stack, clear BSS, uncompress kernel, then transfer control to it |
| | | |
| assembler code
(different startup_32) | { | ➡ <i>Process 0</i> is created |
| | | ➡ step 2: set up initial page tables, turn on <i>address translation</i> (Ch 7) |
| | | ➡ process 0 knows how to handle some aspects of paging |
| C code
(start_kernel) | { | ➡ Do further initialization |
| | | ➡ step 3: initialize rest of kernel, create the "init" process (i.e., <i>process 1</i> , which is the ancestor of all other user processes) |
| | | ➡ invoke the <i>scheduler</i> |

➡ Note: weenix is not exactly Linux



BIOS Device Drivers

- ➡ Originally, the B^IO provided drivers for all devices
 - ▬ OS would call BIOS-provided code whenever it required services of a device driver
- ➡ These drivers sat in low memory and provided minimal functionality
 - ▬ later systems would copy them into primary memory
 - ▬ even later systems would provide their own drivers
 - ▬ nevertheless, B^IO drivers are still used for booting
 - how else can you do it?



Beyond BIOS

➡ BIOS

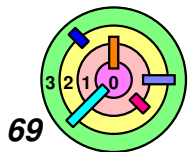
- designed for 16-bit x86 of mid 1980s
- not readily extensible to other architectures

➡ Open Firmware

- designed by Sun
- portable
- drivers, boot code in Forth
 - compiled into bytecode

➡ Intel developed a replacement for BIOS called *EFI (Extensible Firmware Interface)*

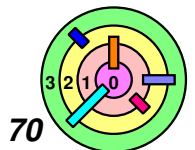
- also uses bytecode



Ch 4: Operating-System Design

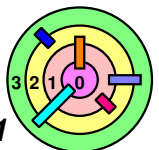
Bill Cheng

<http://merlot.usc.edu/cs402-f13>



OS Design

- ➡ We will now look at how OSes are constructed
 - what goes into an OS
 - how they interact with each other
 - how is the software structured
 - how performance concerns are factored in
- ➡ We will introduce new components in this chapter
 - scheduling (Ch 5)
 - file systems (Ch 6)
 - virtual memory (Ch 7)
- ➡ We will start with a simple hardware configuration
 - what OS is needed to support this
- ➡ Applications views the OS as the "computer"
 - the OS needs to provide a *consistent* and *usable interface*
 - while being *secure* and *efficient*
 - that's a pretty tall order!



OS Design

- ➡ Our goal is to build a general-purpose OS
- can run a variety of applications
 - some are interactive
 - many use network communication
 - all read/write to a file system
 - it's like most general-purpose OSes
 - Linux
 - Solaris
 - FreeBSD
 - Mac OS X
 - Chromium OS (has a Linux kernel)
 - Windows (the only one that's not directly based on Unix)
 - all these OSes are quite similar, functionally!
they all provide:
 - processes
 - threads
 - file systems
 - network protocols with similar APIs
 - user interface with display, mouse, keyboard
 - access control based on file ownership and that file owners can control

