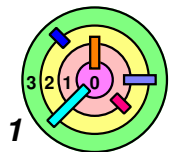



Housekeeping (Lecture 11 - 10/2/2013)

- ➡ Warmup #2 due at 11:45pm this Friday, 10/4/2013
 - if you have code from a previous semester, be very careful and **not copy any code from it**
 - it's best if you just get rid of it
- ➡ **Grading guidelines** is the **ONLY** way we will grade and we will grade on `nunki.usc.edu` in our grading account (which you don't have access to)
 - it's a good idea to run your code against the grading guidelines
- ➡ After submission, make sure you **Verify Your Submission**
- ➡ Have you installed **Ubuntu 11.10** on your laptop/desktop?
- ➡ Do you have partners for kernel assignments?
 - work with your potential partners for warmup 2
 - again, work at high level and must **not** share code



Housekeeping (Lecture 11 - 10/2/2013)


- ➡ Please do not address your post to the class Google Group to me (or the TA or the course producer)
 - if you post to the class Google Group, I will assume that it's not addressed to us (even if you copy us in the e-mail)
- ➡ We will not respond to posts to the class Google Group about the kernle assignments immediately
 - we will wait for at least 2 hours
 - this gives your classmates an opportunity to earn extra credits



Kernel Programming Assignments

Bill Cheng

<http://merlot.usc.edu/cs402-f13>



Kernel Programming Assignments



Tom Doeppner's *weenix* source and binary code

- provided as `weenix-assignment-1.0.7.tar.gz`
- incomplete
- contains code like:


```
NOT_YET_IMPLEMENTED("PROCS: bootstrap");
```

assignment name

function name

- your job is to implement these functions by replacing these lines with your code
 - ◆ please replace them *in-place*
- to look for such code:

```
grep PROCS: kernel/*.c
grep PROCS: kernel/**/*.c
grep PROCS: kernel/**/*.c
grep PROCS: kernel/**/*.c
```



Download and Setup

```
% gunzip -c weenix-assignment-1.0.7.tar.gz | \
    tar xvf -
% cd weenix-assignment-1.0.7/weenix
% make clean
% make
% ./weenix -n
```

- if all goes well, you should see tons of stuff fly by and the following at the bottom of the console:


```
Not yet implemented: PROCS: bootstrap, file
    main/kmain.c, line 127
panic in main/kmain.c:129 bootstrap():
    weenix returned to bootstrap()!!! BAD!!!
```

Kernel Halting.



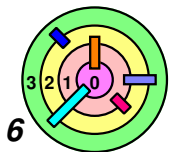
Make sure you have tried the above this weekend

- any problem, let us know **NOW**



Documentation

- ➡ The `weenix` documentation is in `doc/latex/documentation.pdf`
 - introduces `weenix` to you
 - detailed instructions on assignments
 - you must read it thoroughly
- ➡ We are doing three of the assignments
 - *Processes and Threads (PROCS)*
 - *Virtual File System (VFS)*
 - *Virtual Memory (VM)*
- ➡ We are *not* doing two of the assignments
 - Drivers (DRIVERS)
 - System V File System (S5FS)
 - these are done for you and they are compiled and provided as libraries
 - `kernel/libdrivers.a` and `kernel/libS5fs.a`



Compilation and Configuration

➡ **Config.mk controls what gets compiled and configured into the kernel**


- ▬ for PROCS, use the original Config.mk
 - set DRIVERS to 1 to complete this assignment
- ▬ for VFS, set DRIVERS and VFS to 1
- ▬ for VM, set DRIVERS, VFS, S5FS, and VM to 1
 - set VM to 0 at the beginning to get kernel/mm/pframe.c working
 - then set VM to 1 to work on the rest of the assignment
 - set DYNAMIC to 1 in the end if everything is working

➡ **Modify Config.mk first, then do:**

```
% make clean
% make
% ./weenix -n
```

➡ **Use the debugger *right away!***

```
% ./weenix -n -d gdb
```



Debugging

➡ Use the debugger *right away!*

```
% ./weenix -n -d gdb
```

— unfortunately, when you do the above, you won't see kernel debugging messages

```
% ./weenix -n -d gdb
```

➡ To see kernel debugging messages AND debug the kernel, do:

— set GDB_WAIT to 1 in Config.mk then recompile kernel


— run weenix under gdb with:

```
% ./weenix -n -d gdb -w
```

— unfortunately, if you want to run *without gdb*, weenix will *hang*

- you have to set GDB_WAIT back to 0 and recompile if you want to run weenix without gdb

- you should set GDB_WAIT back to 0 when you submit your assignment for grading



Submissions

➡ Electronic submissions only with `bsubmit`

➡ Processes and Threads (PROCS)

— need to fill out `procs-README.txt`, it's your assignment documentation

○ this is where you should include

1) how to split the points (in terms of percentages and must sum to 100%)

2) brief justification about the split

`% make procs-submit`

`tar cvzf procs-submit.tar.gz \`


`Config.mk \`

`procs-README.txt \`

`...`

○ you must add the required information about what you have implement in `procs-README.txt` as well

— `scp` or `sftp` `procs-submit.tar.gz` to `nunki/aludra` and run `bsubmit` there



Submissions



Virtual File System (VFS)

- need to fill out `vfs-README.txt`

```
% make vfs-submit
```

```
...
```



Virtual Memory (VM)

- need to fill out `vm-README.txt`


```
% make vm-submit
```

```
...
```



Submit source code only

- we will deduct 2 points if you submit binary files
- we will deduct 2 points if you submit unmodified files
- we will deduct 2 points if you do not keep the same directory structure




Verify Your Submission

- ➡ Assume that in your home directory, you have
- a *pristine* weenix-assignment-1.0.7.tar.gz
 - your submission file, e.g., procs-submit.tar.gz


- ➡ Do the following to verify your submission

```
% rm -rf /tmp/xyzzzy
% mkdir /tmp/xyzzzy
% cd /tmp/xyzzzy
% gunzip -c ~/weenix-assignment-1.0.7.tar.gz | \
    tar xvf -
% cd weenix-assignment-1.0.7/weenix
% gunzip -c ~/procs-submit.tar.gz | tar xvf -
% make clean
% make
% ./weenix -n
```



Grading Guidelines


- ➡ For *every function* you need to implement, you need to think about its pre-conditions and post-conditions
 - you need to add `dbg()` and `KASSERT()` statements to show:
 - 1) you understand what a function is suppose to do
 - 2) what the important pre-conditions and post-conditions are
 - see grading guidelines for required `KASSERT()` statements
 - right after `KASSERT()`, prefix string in `dbg()` with `"(GRADING# X.Y)"`, where X and Y are from grading guidelines and # is the kernel assignment number
- ➡ You need to show us that you have thoroughly verified that your code works
 - turn in additional code or procedure and document how to enable/use these tests
 - it's okay if this is similar to our test code
- ➡ Run your code with additional test code from us
 - we will make suggestions in the grading guidelines



Backing Up Your Work & Collaboration


- ➡ You *have to* have a plan to *backup your code* and backup routinely
 - if you lose your work, there's nothing we can do
- ➡ Ubuntu 11.10 comes with "Ubuntu One"
 - free (up to 5GB) personal cloud
 - set one up for your group
 - all group members can sync with Ubuntu One
 - create a directory called `cs402backup` under `~/Ubuntu One`
- ➡ One simple way to backup your work
 - at the end of each day, do (for example, if today is 10/01/13):

```
% make procs-submit
% mv procs-submit.tar.gz procs-100213.tar.gz
% cp procs-100213.tar.gz ~/Ubuntu\ One/cs402backup
```
 - make sure only one person does this, or you will wipe out your partner's backup
 - can do backups more often if you'd like
 - `procs-100213-[time].tar.gz`




Backing Up Your Work & Collaboration

- ➡ You should use `git` to collaborate among project partners
 - read "Pro Git" (a free online book, one of our textbooks)
- ➡ But you need a shared repository in the cloud
 - there are free `git` repositories on the web
 - unfortunately, most of the free ones are required to be visible by the world - you must **not** use these
 - if you can find one that can be setup such that only you and your partners can see it, then you can use it
 - maybe you can use **Ubuntu One** as your repository
 - the course producers will share their experience
- ➡ If you have two people working on the same file and then update the repository one after another
 - `git` will attempt to merge the changes, but it may not be what you want
 - maybe it's best to **coordinate** and not have two people modifying the same file




Early and Late Policies

➡ Similar to warmup projects




Extra Credit

- ➡ You can get extra credit for posting good, useful, and insightful answers to the class Google Group in response to questions posted by other students regarding *kernel* programming assignments
- the maximum number of extra credit points you can get is **10** points for each of the kernel assignments (on a 100-point scale)




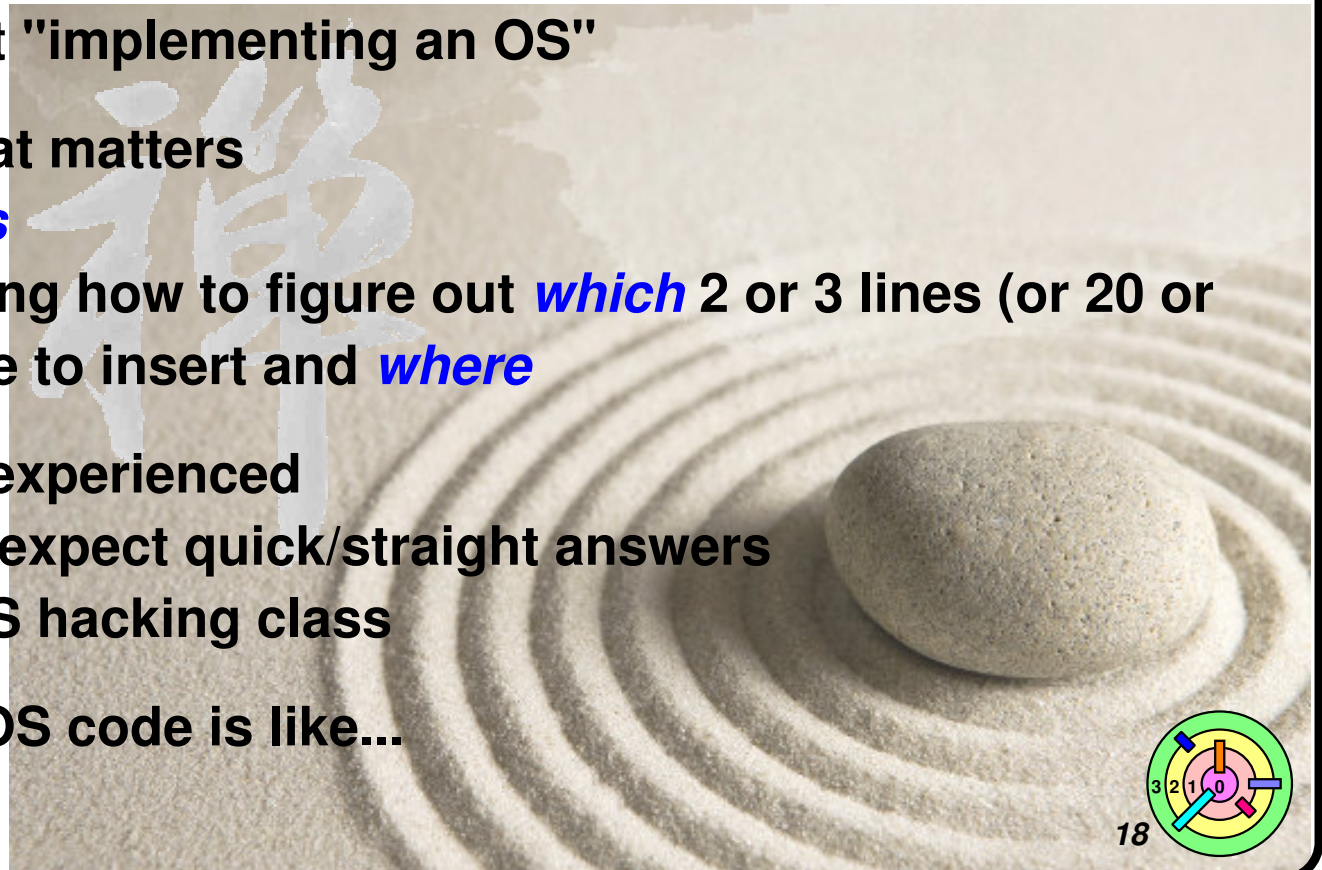
How Do You Start?

- ➡ Definitely start with the documentation
- ➡ Read code, read lots and lots of code
 - try things out and see what happens (debugging statements)
 - you need to *absorb* other people's code, make sense of it
 - that's what OS hacking (in the good sense) is all about
 - it's *not* about "implementing an OS"
- ➡ It's the *process* that matters
 - *not* the *answers*
 - it's about learning how to figure out *which* 2 or 3 lines (or 20 or 30 lines) of code to insert and *where*
- ➡ So, it needs to be experienced
 - you should not expect quick/straight answers
 - this is not an OS hacking class
- ➡ Learning to write OS code is like...




How Do You Start?

- ➡ Definitely start with the documentation
- ➡ Read code, read lots and lots of code
 - ▬ try things out and see what happens (debugging statements)
 - ▬ you need to *absorb* other people's code, make sense of it
 - ▬ that's what OS hacking (in the good sense) is all about
 - it's *not* about "implementing an OS"
- ➡ It's the *process* that matters
 - ▬ *not* the *answers*
 - ▬ it's about learning how to figure out *which* 2 or 3 lines (or 20 or 30 lines) of code to insert and *where*
- ➡ So, it needs to be experienced
 - ▬ you should not expect quick/straight answers
 - ▬ this is not an OS hacking class
- ➡ Learning to write OS code is like...
 - ▬ *Zen*




Getting Help

- ➡ If you have questions about the kernel assignments
- read documentation, textbook, lecture slides, read more code
 - post your questions to the class Google Group
 - we will not immediately answer these questions to give your classmates an opportunity to earn extra credit points
 - ◆ we will wait 2 hours
 - if you send us questions about the assignments directly, we may simply forward your post to the class Google Group

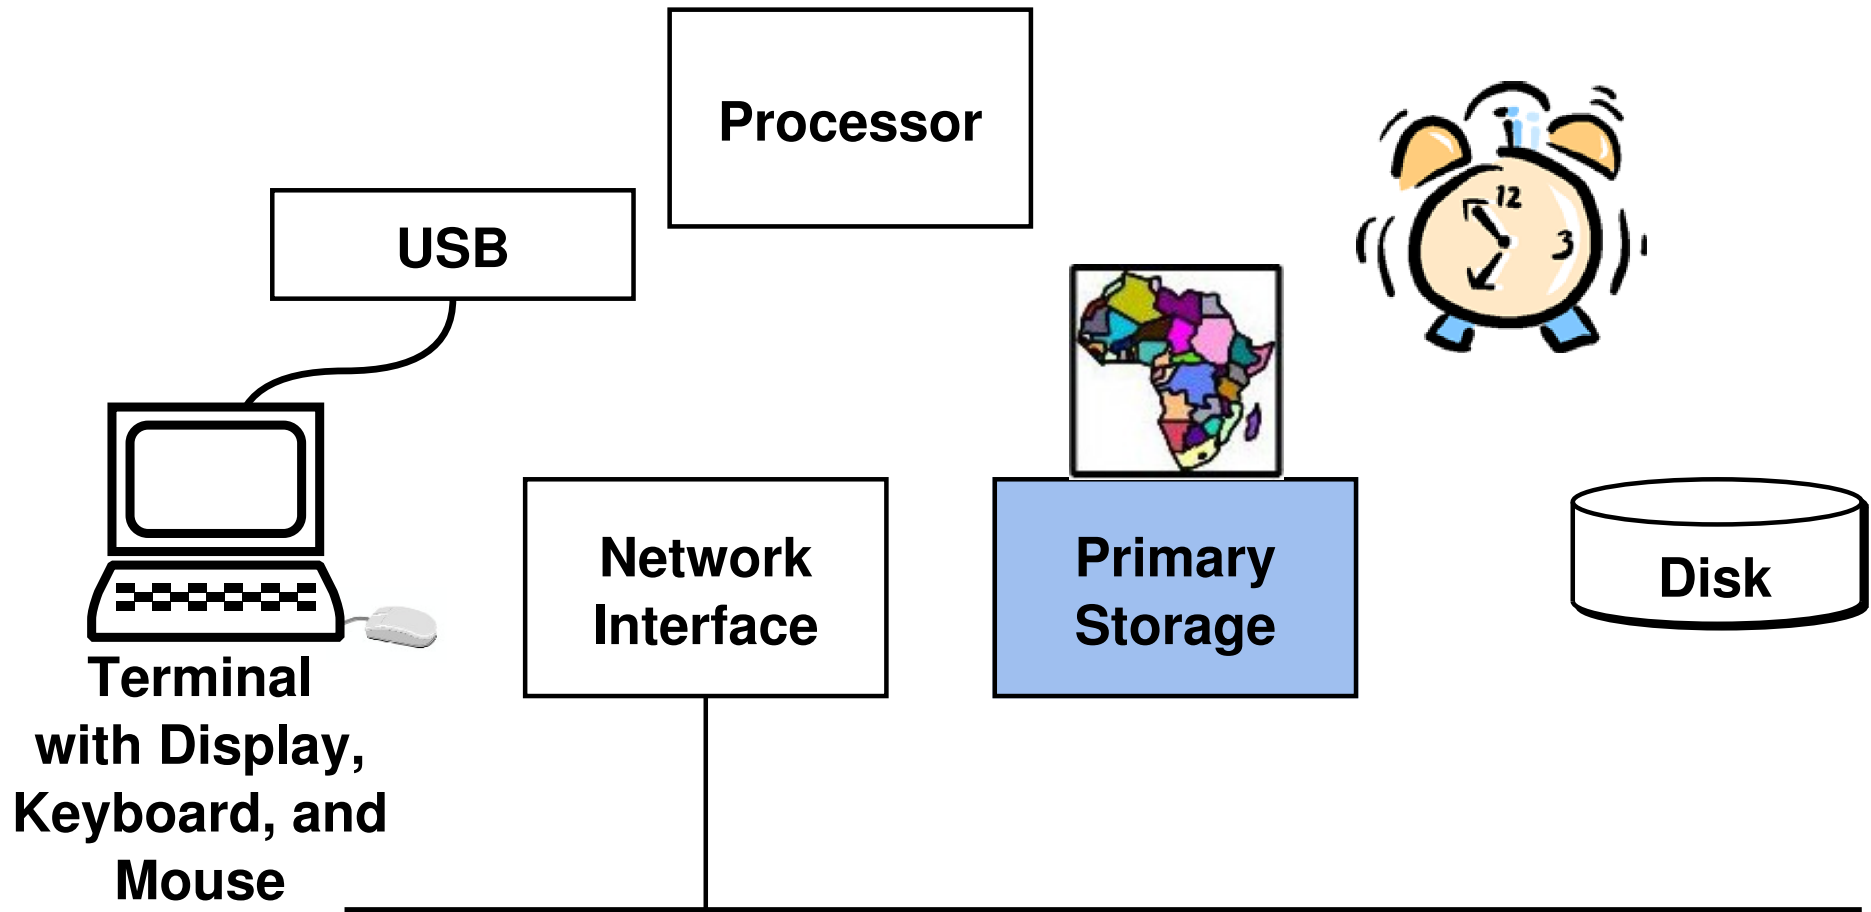


OS Design Issues


- ➡ **Performance**
 - efficiency of application
- ➡ **Modularity**
 - tradeoffs between modularity and performance
- ➡ **Device independence**
 - for new devices, don't need to write a new OS
- ➡ **Security/Isolation**
 - isolate OS from application




Simple Configuration



- ➡ Early 1980s OS, so we can focus on the basic OS issues
- no support for bit-mapped displays and mice
 - generally less efficient design


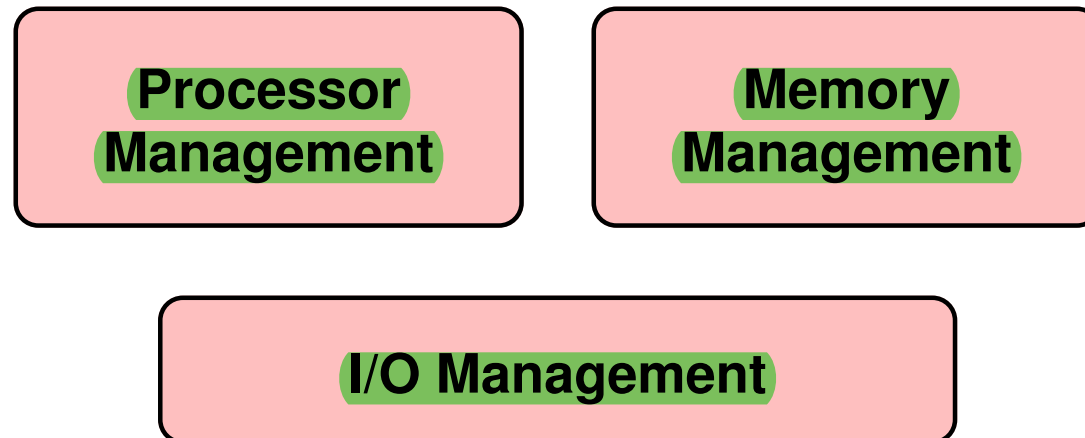


OS Components


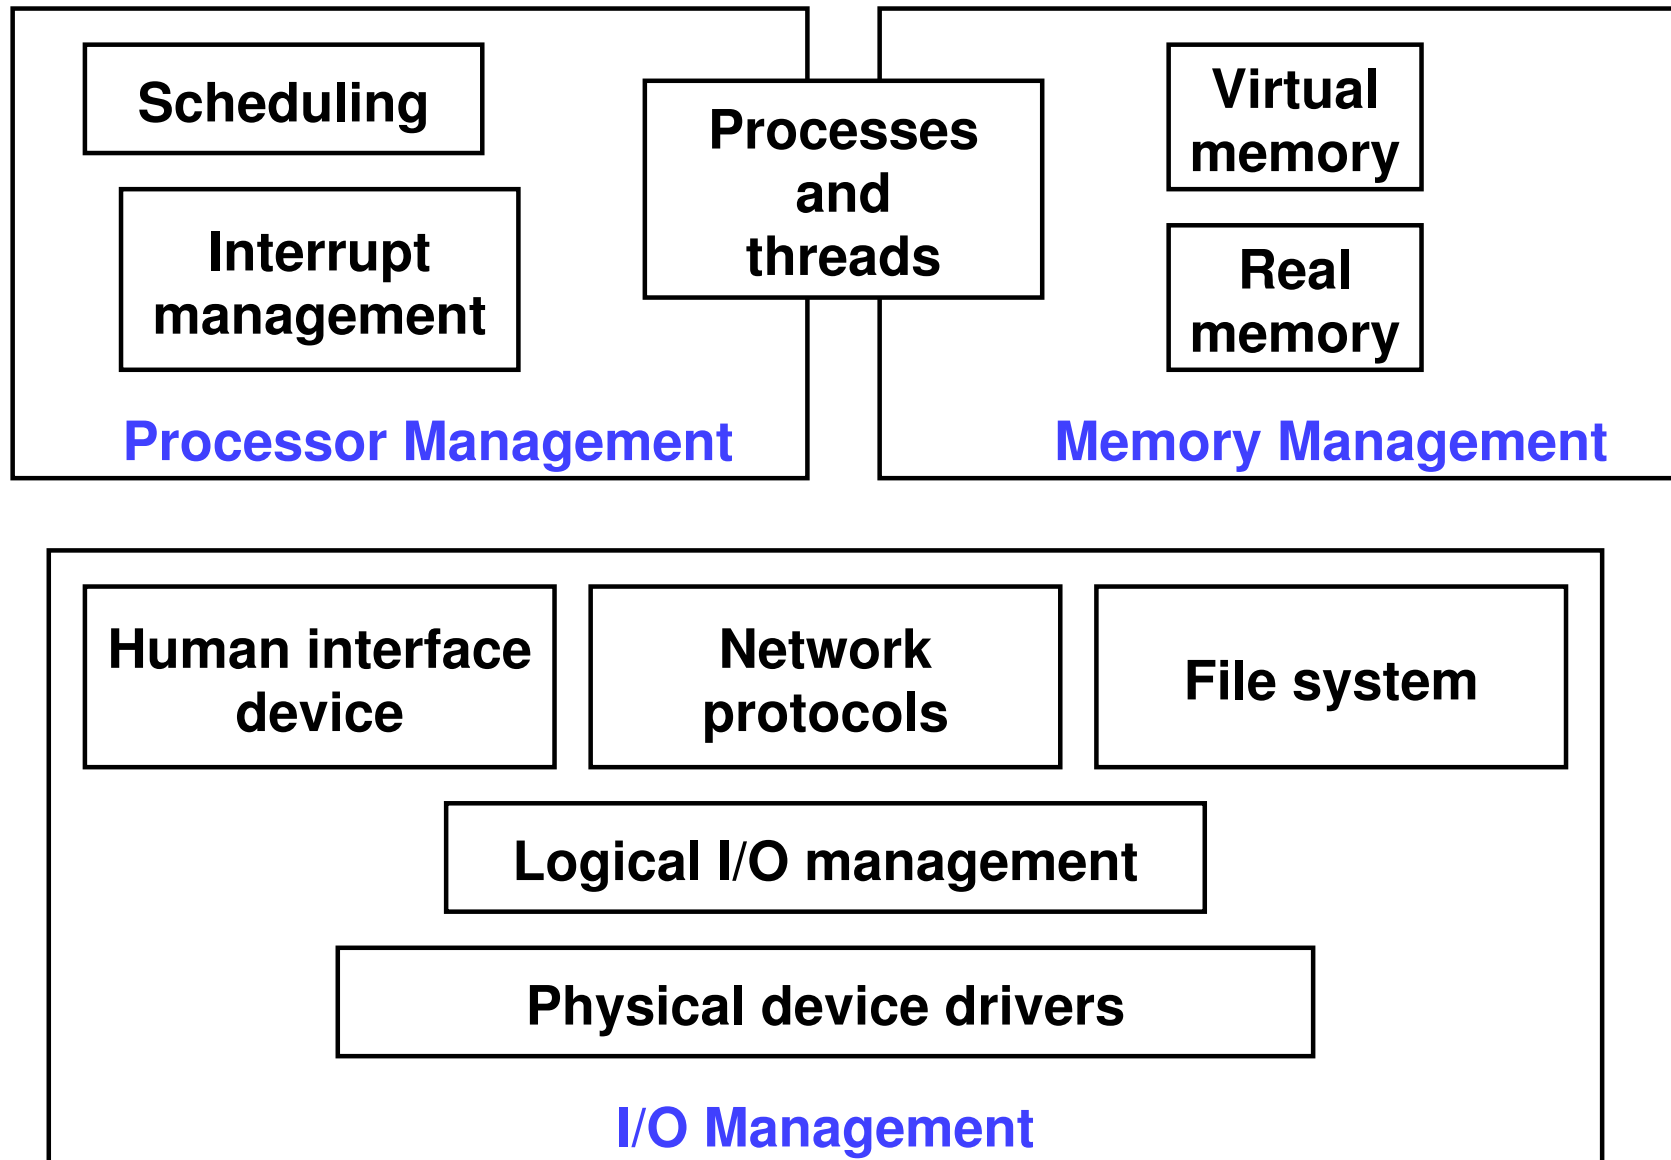


Applications


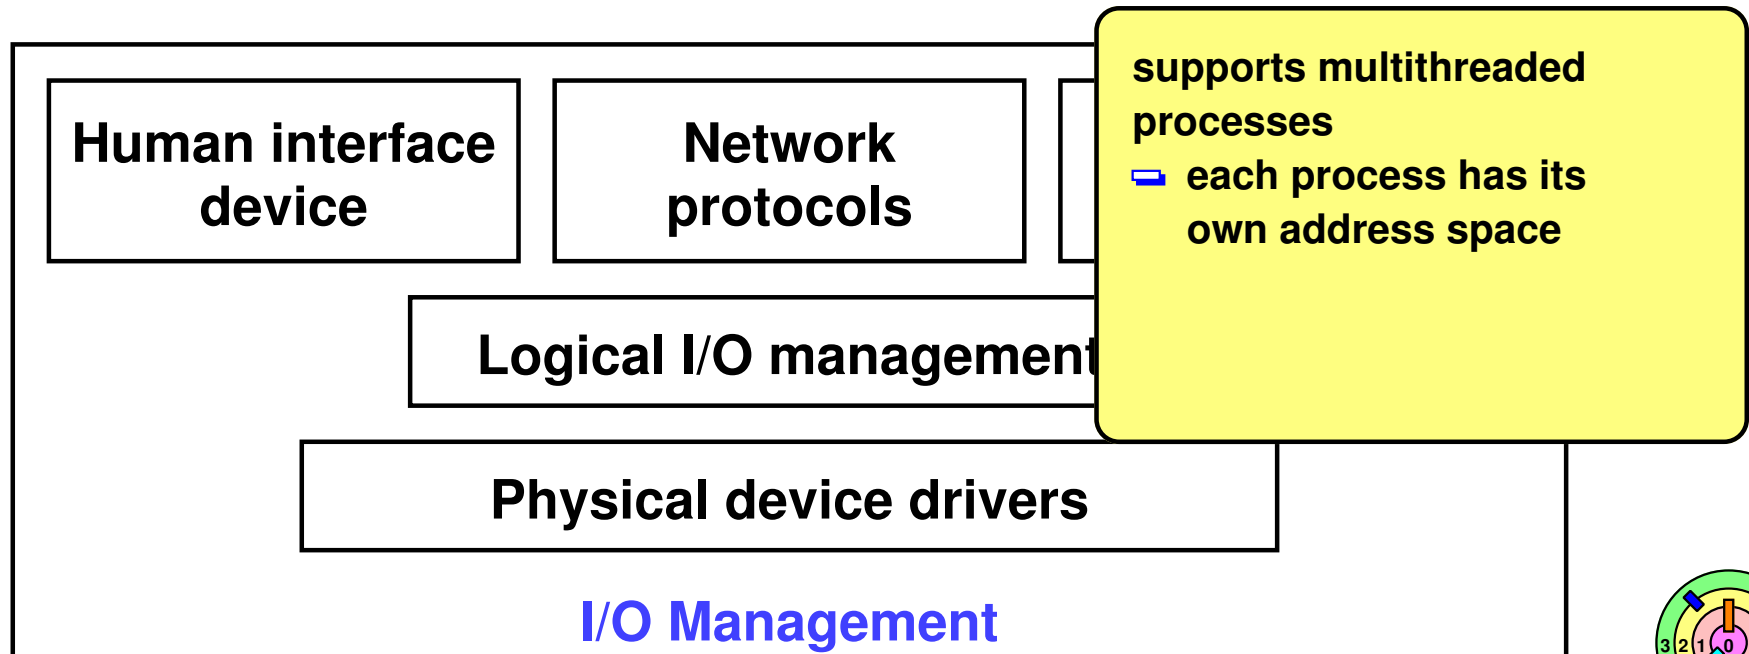
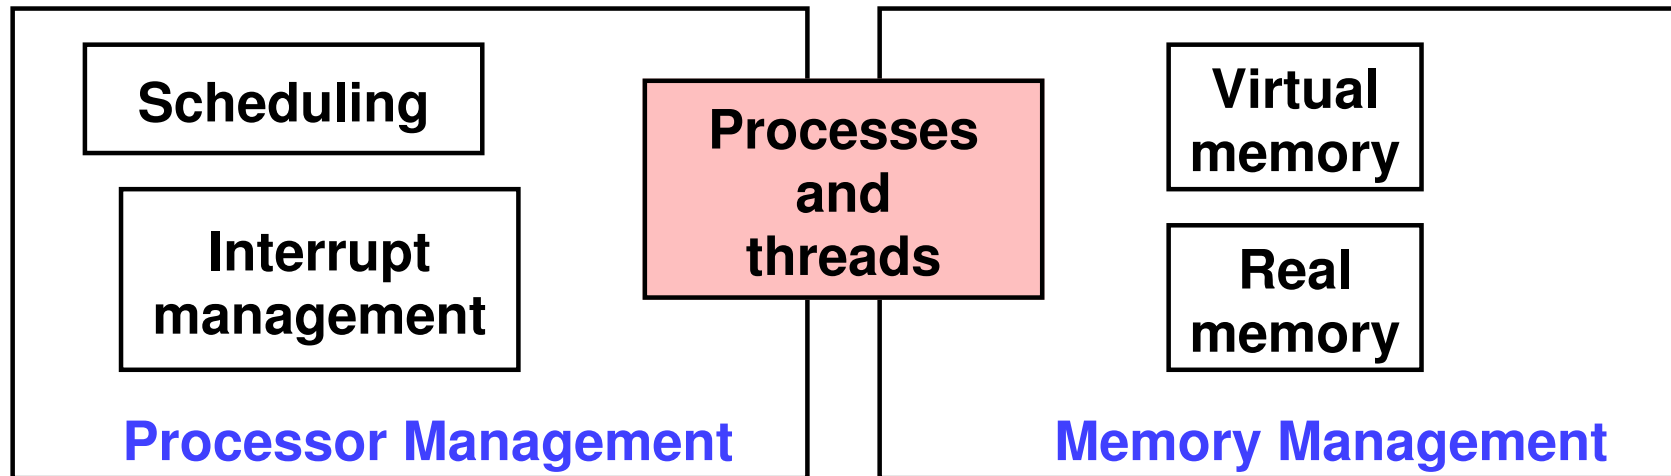
OS




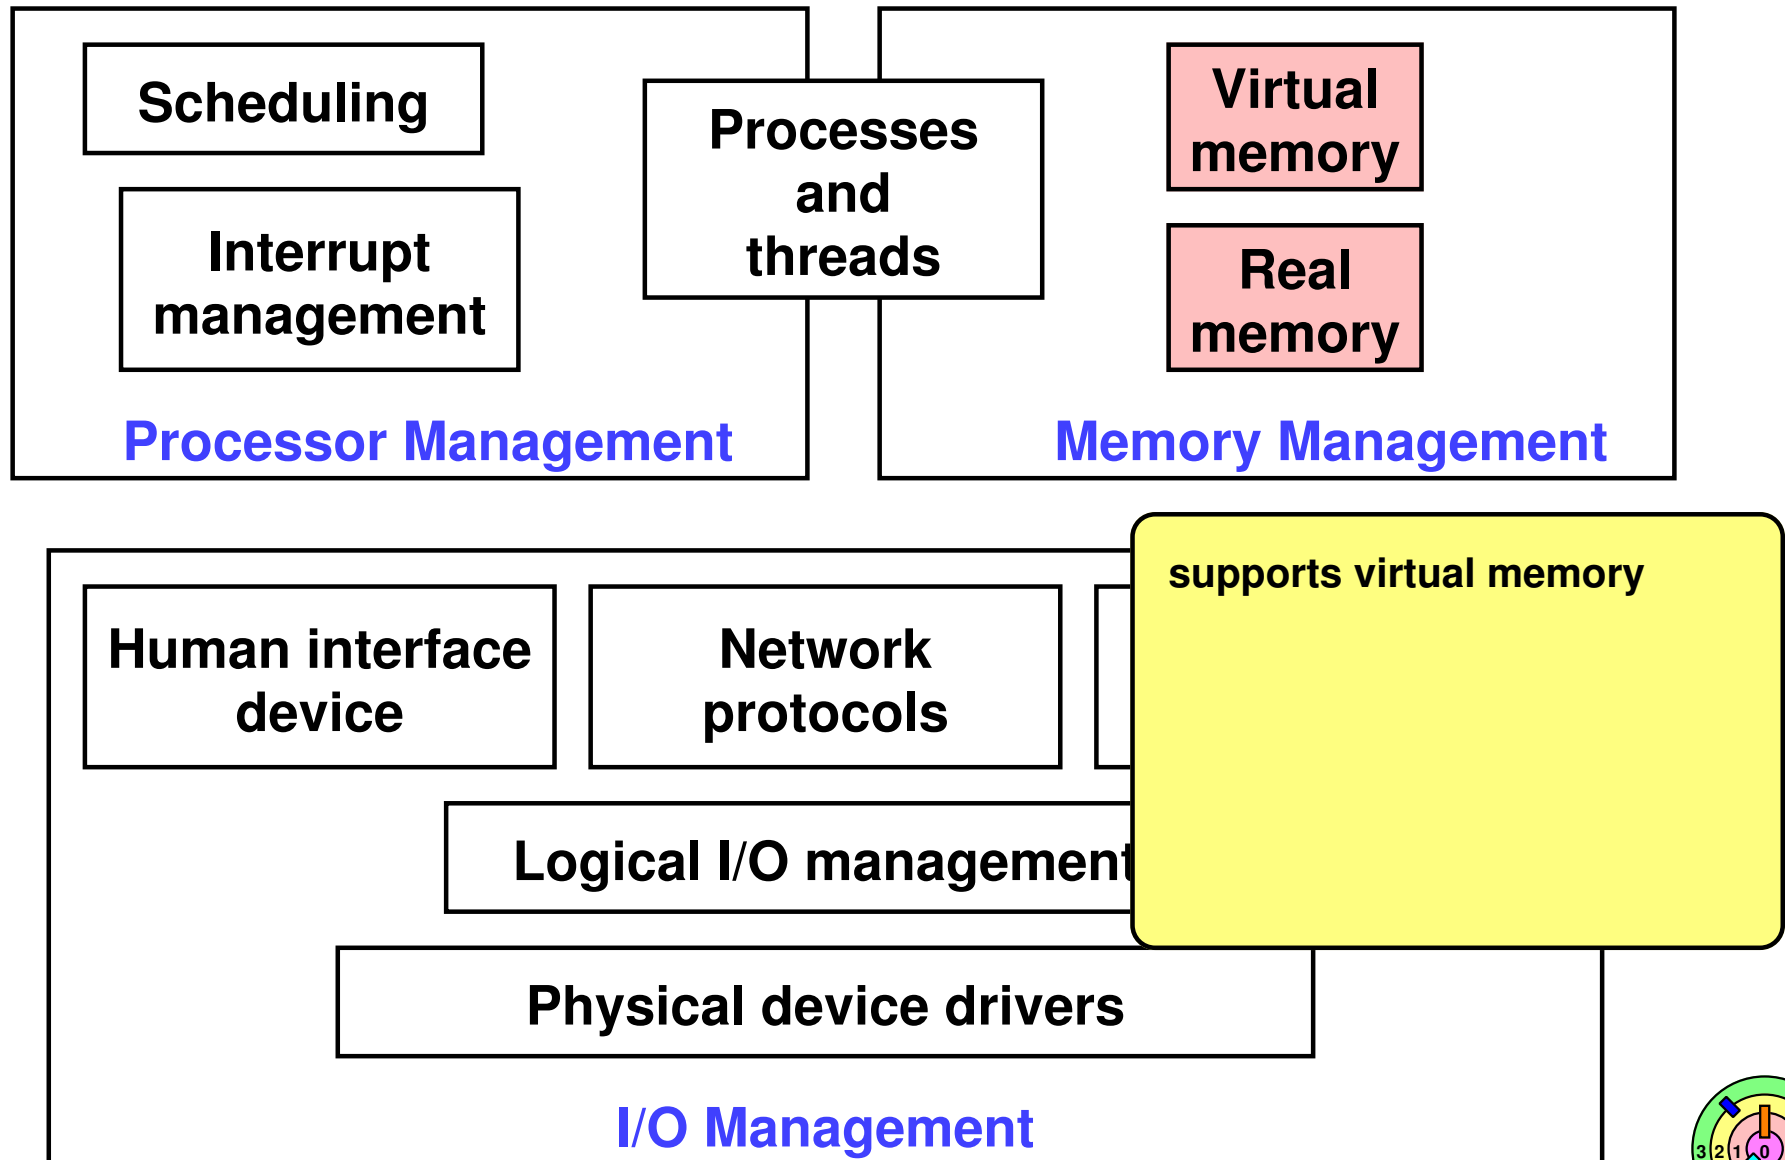
OS Components




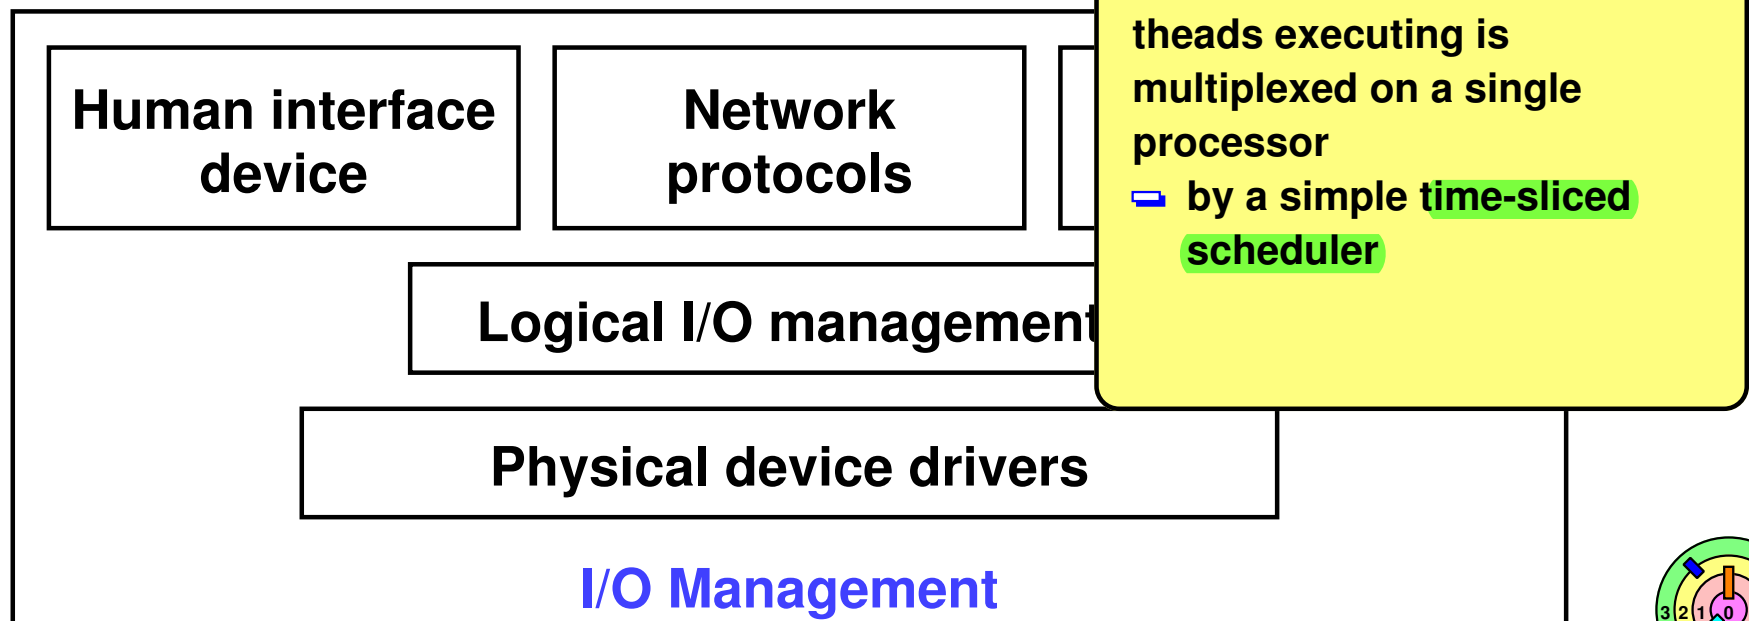
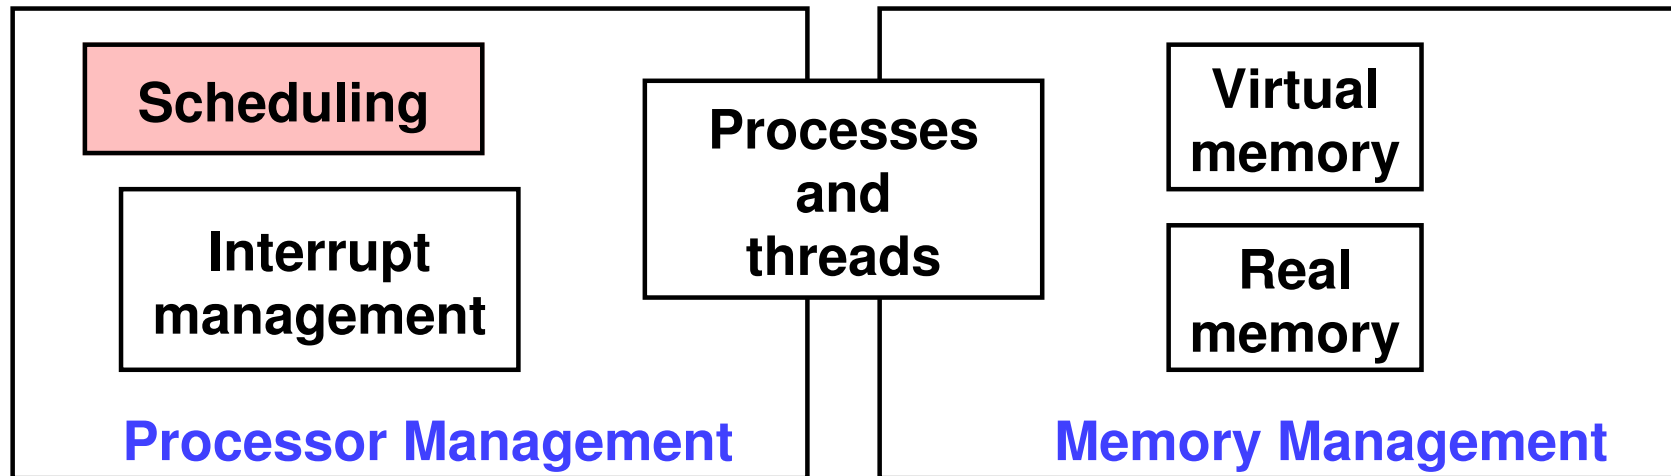
OS Components




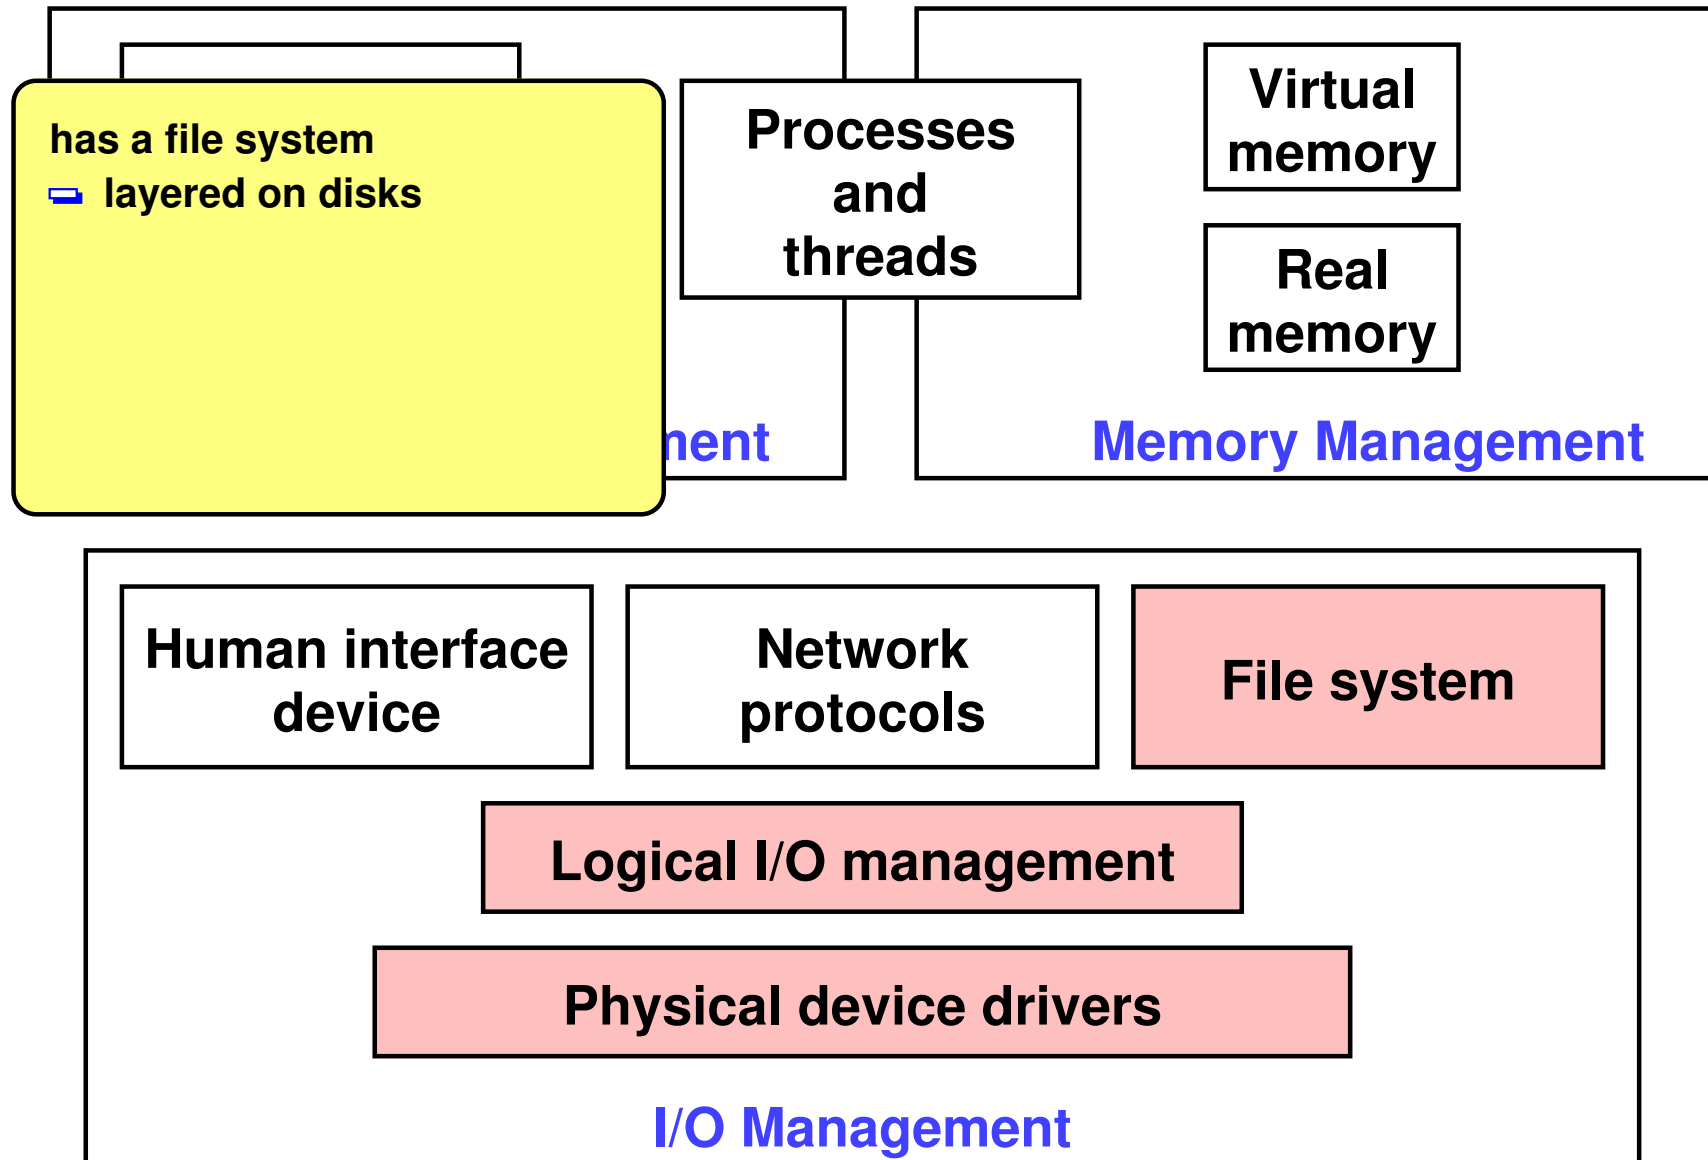
OS Components




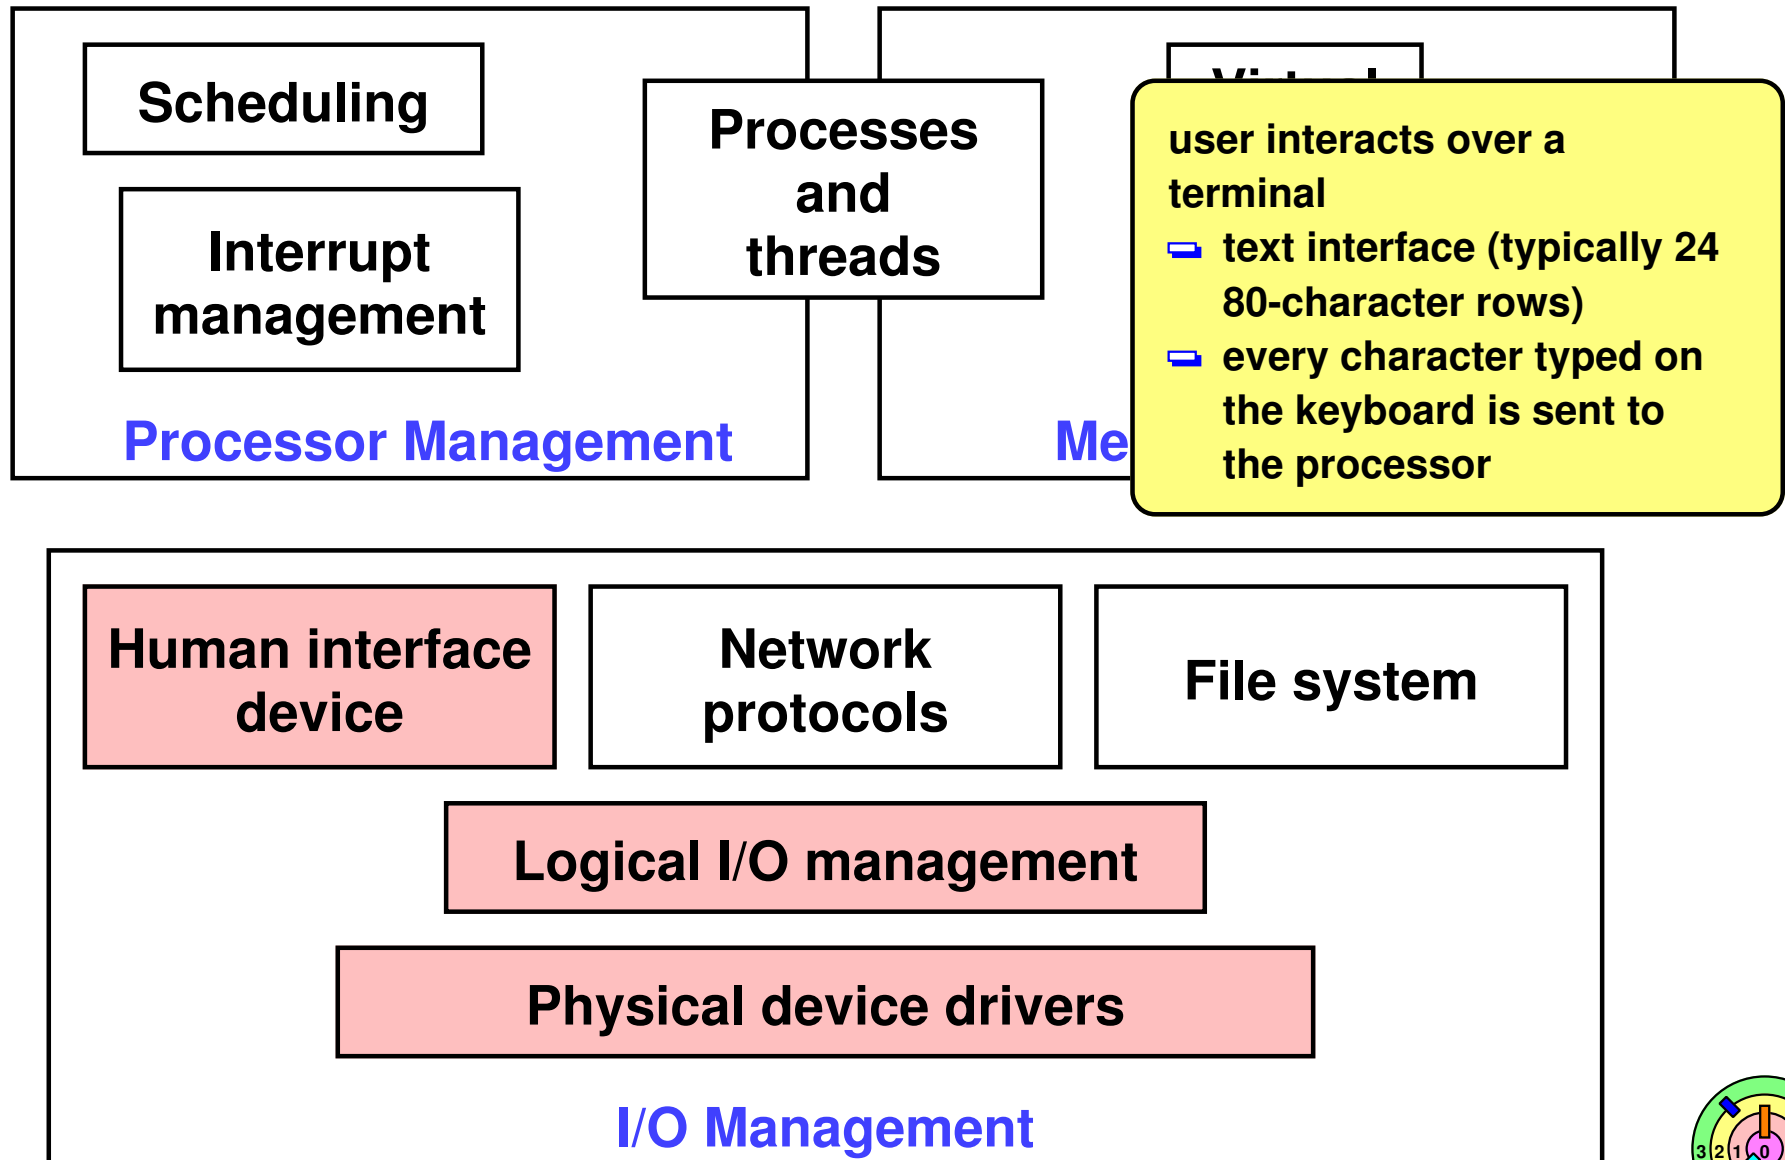
OS Components




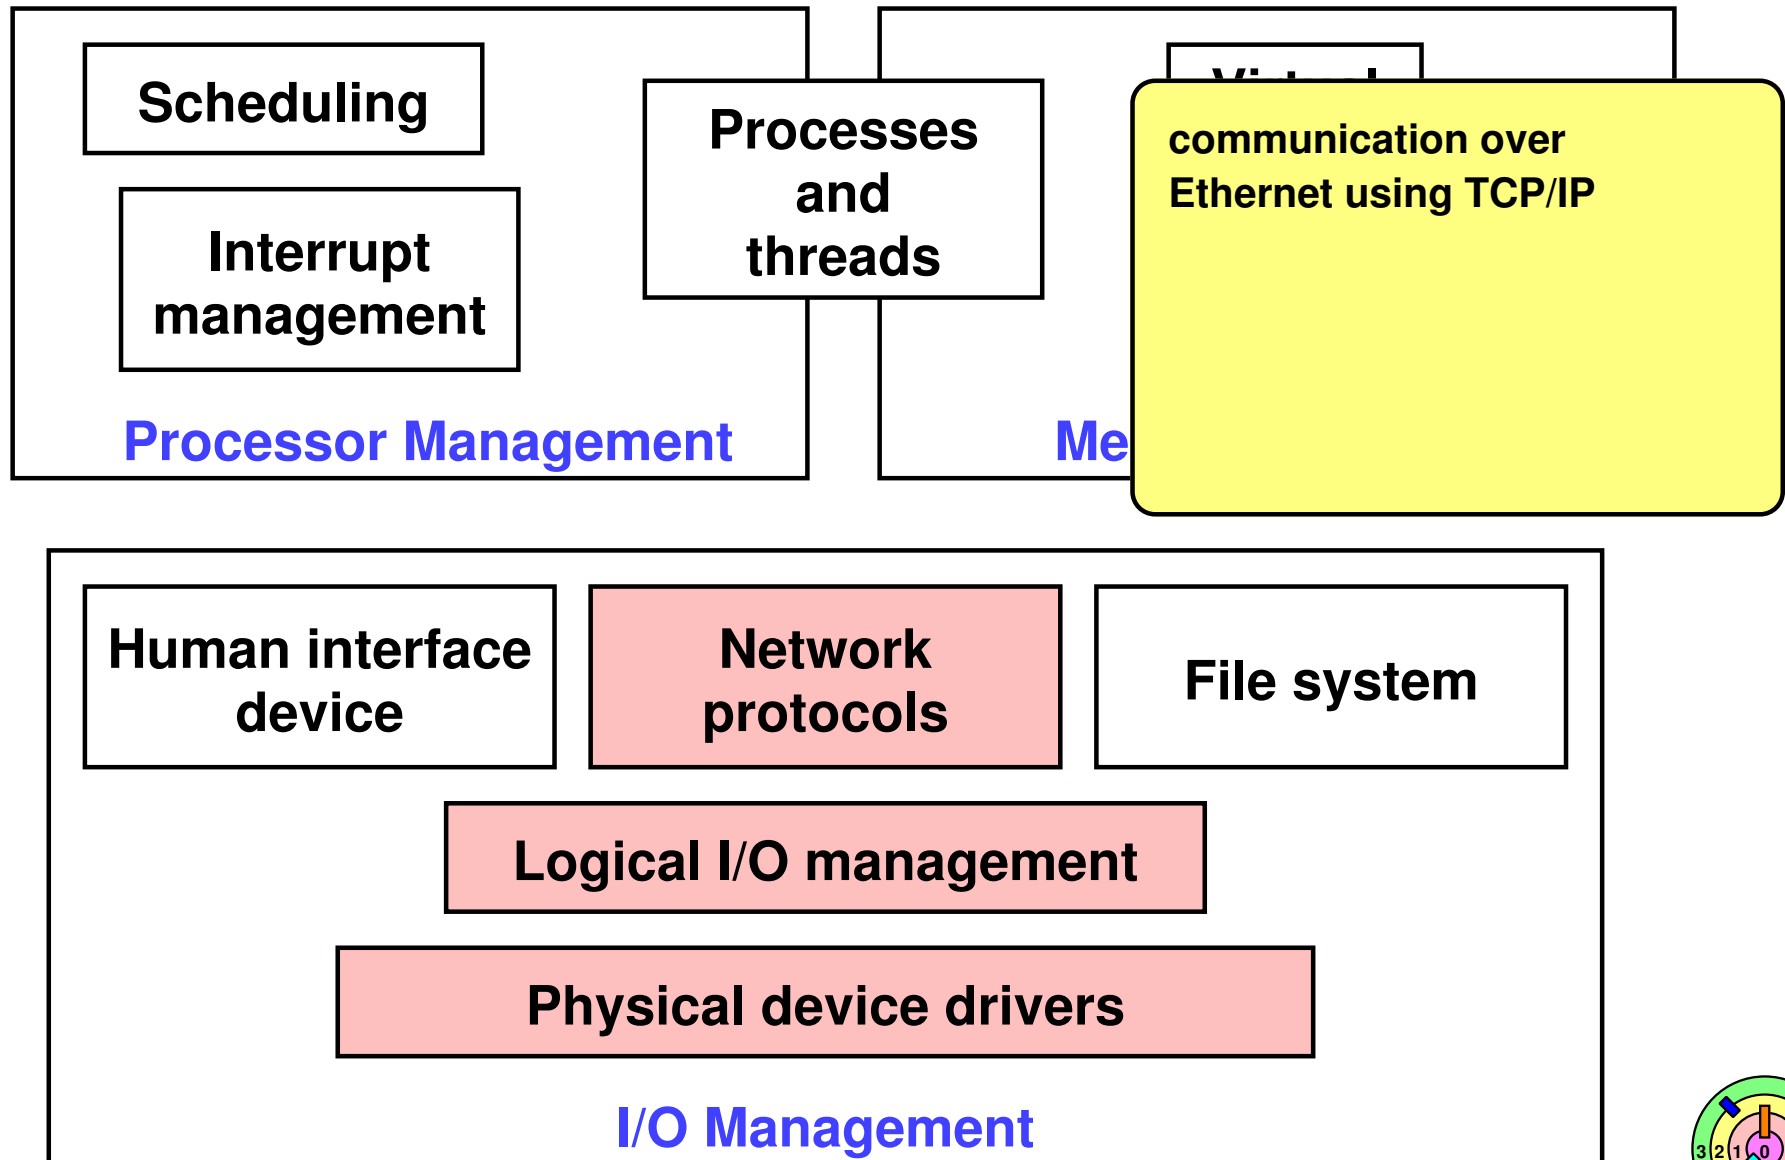
OS Components



OS Components




OS Components



Applications

- ➡ How does an *application* view these OS components?
- ➡ From an application program's point of view, our system has:
 - processes with threads
 - a file system
 - terminals (with keyboards)
 - a network connection
- ➡ Need more details on these components... Need to look at:
 - how applications use them
 - how this affects the design of the OS



Processes And File Systems




The purpose of a *process*

- holds an *address space*
- holds a group of *threads* that execute within that address space
- holds a collection of references to *open files* and other "*execution context*"




Address space:

- set of addresses that threads of the process can usefully reference
- more precisely, it's the content of these addressable locations
 - text, data, bss, dynamic, stack regions and what's in them



Processes And File Systems

- ➡ Design issue:
 - ➡ how should the OS *initialize* these address space regions?
- ➡ Simple approach: copy their contents from the file system to the process address space (as part of the *exec* operation)
 - ➡ quite wasteful (both in space and time) for the text region since it's read-only data
 - should *share* the text region
 - ➡ what about data regions? they can potentially be written into
 - can also *share* a portion of a data region if that portion is never modified



Remember This?

➡ **Virtual Memory** (using *memory maps*)



Text
 main 4096
 subr 4132
 printf 4156
 write 16156
 startup 16172

Data
 aX 16384
 printfargs 16388
 StandardFiles 16396

BSS
 X 17420
 errno 17680

Page Table

Start	Access	Physical Addr
0	-	-
4096	R	•
8192	R	•
12288	R	•
16384	R/W	•



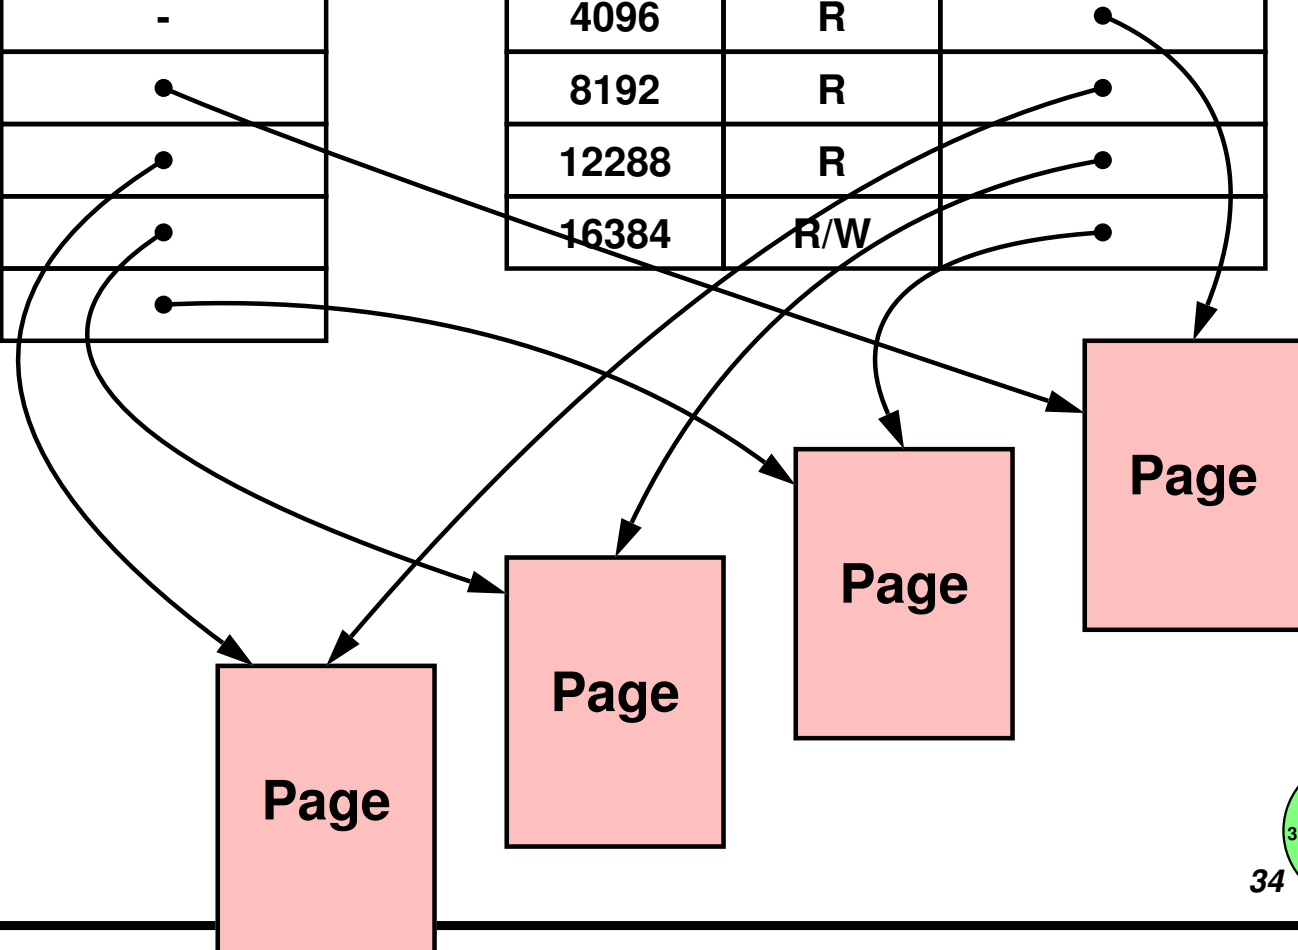
Processes Can Share Memory Pages

Process 2 Page Table

<i>Start</i>	<i>Access</i>	<i>Physical Addr</i>
0	-	-
4096	R	•
8192	R	•
12288	R	•
16384	R/W	•

Process 1 Page Table

<i>Start</i>	<i>Access</i>	<i>Physical Addr</i>
0	-	-
4096	R	•
8192	R	•
12288	R	•
16384	R/W	•




Processes And File Systems

➡ For the text region, why bother copying the executable file into the address space in the first place?

— can just *map* the file into the address space (Ch 7)

- *mapping* let the OS *tie* the regions of the address space to the file system
- address space and files are divided into pieces, called *pages*
- if several processes are executing the same program, then at most one copy of that program's text page is in memory at once

— *text regions* of all processes running this program are setup, using hardware address translation facilities, to share these pages



Processes Can Share Memory Pages

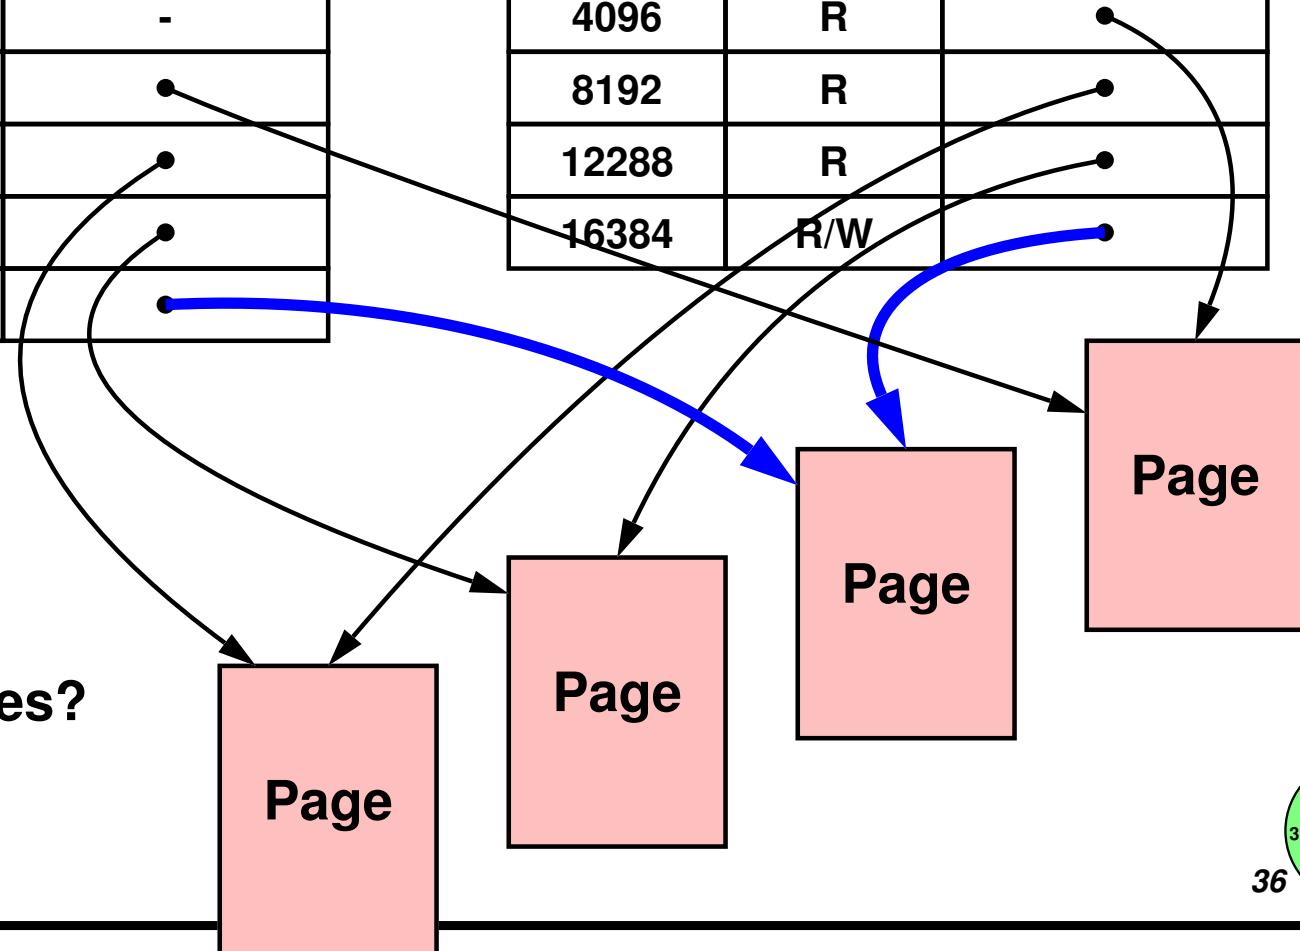
Process 2 Page Table

<i>Start</i>	<i>Access</i>	<i>Physical Addr</i>
0	-	-
4096	R	•
8192	R	•
12288	R	•
16384	R/W	•

Process 1 Page Table


<i>Start</i>	<i>Access</i>	<i>Physical Addr</i>
0	-	-
4096	R	•
8192	R	•
12288	R	•
16384	R/W	•

— can we really
share *data*
segment pages?



Processes And File Systems

- ➡ *Data regions* of all processes running this program *initially* refer to pages of memory containing a copy of the *initial* data region
- ➡ this type of mapping is *private mapping*
 - when does each process really need a private copy of such a page?
 - when data is *modified* by a process, it gets a *new* and *private* copy of the initial page



Processes And File Systems



Copy-on-write (COW):

- a process gets a **private** copy of the page after a thread in the process performs a **write** for the **first time**
 - the basic idea is that only those pages of memory that are modified are copied




The BSS and stack regions use a special form of **private mapping**

- their pages are initialized, with zeros; **copy-on-write**
 - known as **anonymous objects** in `weenix`




If a bunch of processes share a file

- we can also map the file into the address space
- in this case, **the mapping is shared**
- when one process modifies a page, no private copy is made
 - instead, the original page itself is modified
 - **everyone gets the changes**
 - and changes are written back to the file
- ◆ more on issues in Ch 6




Block I/O vs. Sequential I/O

- ➡ Mapping files into address space is one way to perform I/O on files
 - block/page is the basic unit
 - this is referred to as *block I/O*
- ➡ Some devices cannot be mapped into the address space
 - e.g., receiving characters typed into the keyboard, sending a message via a network connection
 - need a more traditional approach using explicit system calls such as `read()` and `write()`
 - this is referred to as *sequential I/O*
- ➡ It also makes sense to be able to read a file like reading from the keyboard
 - similarly, a program that produces lines of text as output should be able to use the same code to write output to a file or write it out to a network connection
 - makes life easier! (and make code more robust)




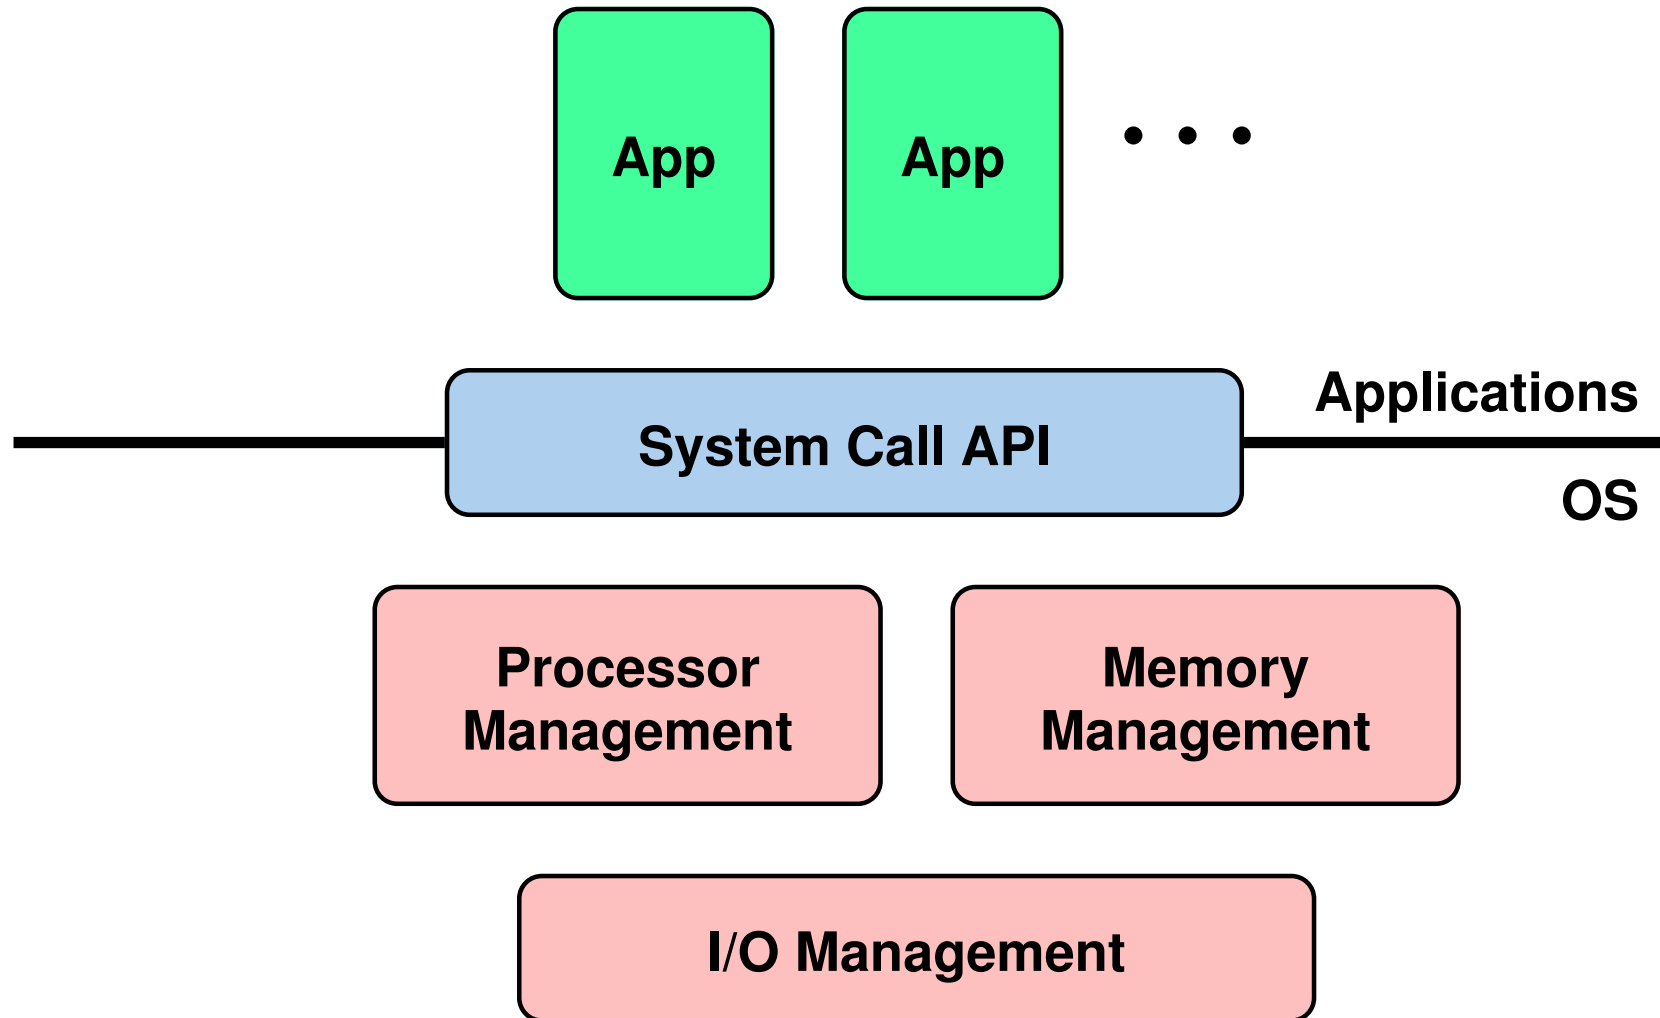
A Simple System: To Be Discussed

- ➡ What is the functionality of the components?
- ➡ What are the key data structures?
- ➡ How is the system broken up into modules?
- ➡ To what extent is the system extensible?
- ➡ What parts run in the OS kernel in privileged mode? What parts run as library code in user applications? What parts run as separate applications?
- ➡ In which execution contexts do the various activities take place?
 - e.g., thread context vs. interrupt context




System Call API

- ➡ Backwards compatibility is an important issue
- try not to change it much (to make the developers happy)



4.1 A Simple System (Monolithic Kernel)

- ➡ *A Framework for Devices*
- ➡ Processes & Threads
- ➡ Storage Management
- ➡ Low-level Kernel (will come back to talk about this after Ch 7)




Computer Terminal



➡ VT100

A "tty"



Devices




Challenges in supporting devices

- device independence
- device discovery



Device naming

- two choices
 - independent name space (i.e., named independently from other things in the system)
 - devices are named as files



A Framework for Devices



Device driver:

- every device is identified by a device "number", which is actually a pair of numbers
 - a *major device number* - identifies the device driver
 - a *minor device number* - device index for all devices managed by the same device driver




Special entries were created in the file system to refer to devices

- usually in the `/dev` directory
 - e.g., `/dev/disk1`, `/dev/disk2` each marked as a *special file*
 - ◆ a *special file* does not contain data
 - ◆ it refers to devices by their major and minor device numbers
 - ◆ if you do `ls -l`, you can see the device numbers

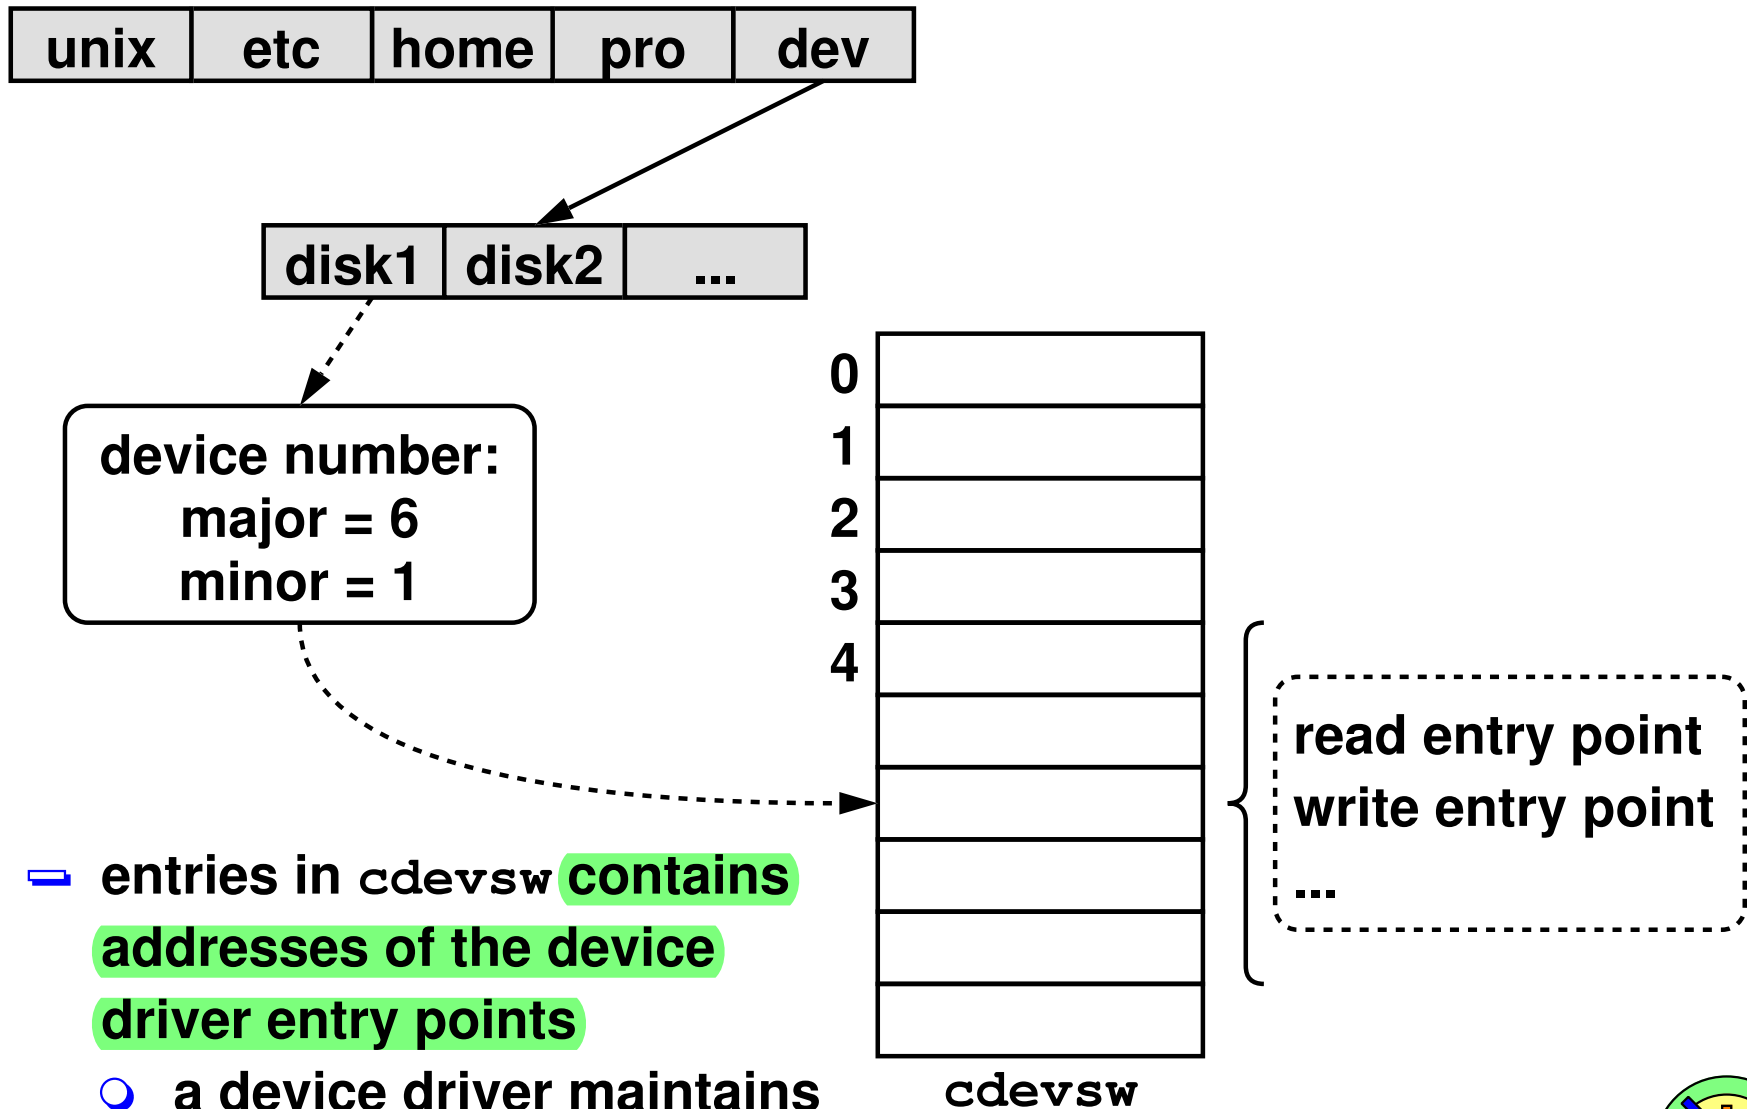


Data structure in the early Unix systems


- statically allocated array in the kernel called `cdevsw` (*character device switch*)



Finding Devices

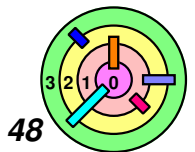


- entries in `cdevsw` contains addresses of the device driver entry points
- a device driver maintains its own data structure




Device Drivers in Early Unix Systems

- ➡ The kernel was statically configured to contain device-specific information such as:
 - interrupt-vector locations
 - locations of device-control registers on whatever bus the device was attached to
- ➡ Static approach was simple, but cannot be easily extended
 - a kernel must be custom configured for each installation



Device Probing

- ➡ First step to improve the old way
 - allow the devices to to be found and automatically configured when the system booted
 - (still require that a kernel contain all necessary device drivers)
- ➡ Each device driver includes a *probe routine*
 - invoked at boot time
 - probe the relevant buses for devices and configure them
 - including identifying and recording interrupt-vector and device-control-register locations
- ➡ This allowed one kernel image to be built that could be useful for a number of similar but not identical installations
 - boot time is kind of long
 - impractical as the number of supported devices gets big



Device Probing



What's the right thing to do?

Step 1: **discover the device** without the benefit of having the relevant device driver in the kernel

Step 2: **find the needed device drivers and dynamically link them into the kernel**

— but how do you achieve this?



Solution: **use meta-drivers**

— a meta-driver handles a particular kind of bus

— e.g., USB (Universal Serial Bus)

- **a USB meta-driver is installed into the kernel**

- any device that goes onto a USB (Universal Serial Bus) must know how to interact with the USB meta-driver via the *USB protocol*


- once a connected device is identified, **system software would select the appropriate device driver and load into the kernel**

- what about applications? how can they reference dynamically discovered devices?




Discovering Devices

- ➡ So, you plug in a new device to your computer on a particular bus
 - OS would notice
 - find a device driver
 - what kind of device is it?
 - where is the driver?
 - assign a name, but how is it chosen?
 - multiple similar devices, but how does application choose?
- ➡ In some Linux systems, entries are added into `/dev` as the kernel discovers them
 - lookup the names from a database of names known as `devfs`
 - downside of this approach is that device naming conventions not universally accepted
 - what's an application to do?
 - some current Linux systems use `udev`
 - user-level application assigns names based on rules provided by an administrator




Discovering Devices

- ➡ What about the case where different devices acted similarly?
 - e.g., touchpad on a laptop and USB mouse
 - how should the choice be presented to applications?
- ➡ Windows has the notion of *interface classes*
 - a device can register itself as members of one or more such classes
 - an application can *enumerate* all currently connected members of such a class and choose among them (or use them all)




4.1 A Simple System (Monolithic Kernel)

- ➡ A Framework for Devices
- ➡ *Processes & Threads*
- ➡ Storage Management
- ➡ Low-level Kernel (will come back to talk about this after Ch 7)

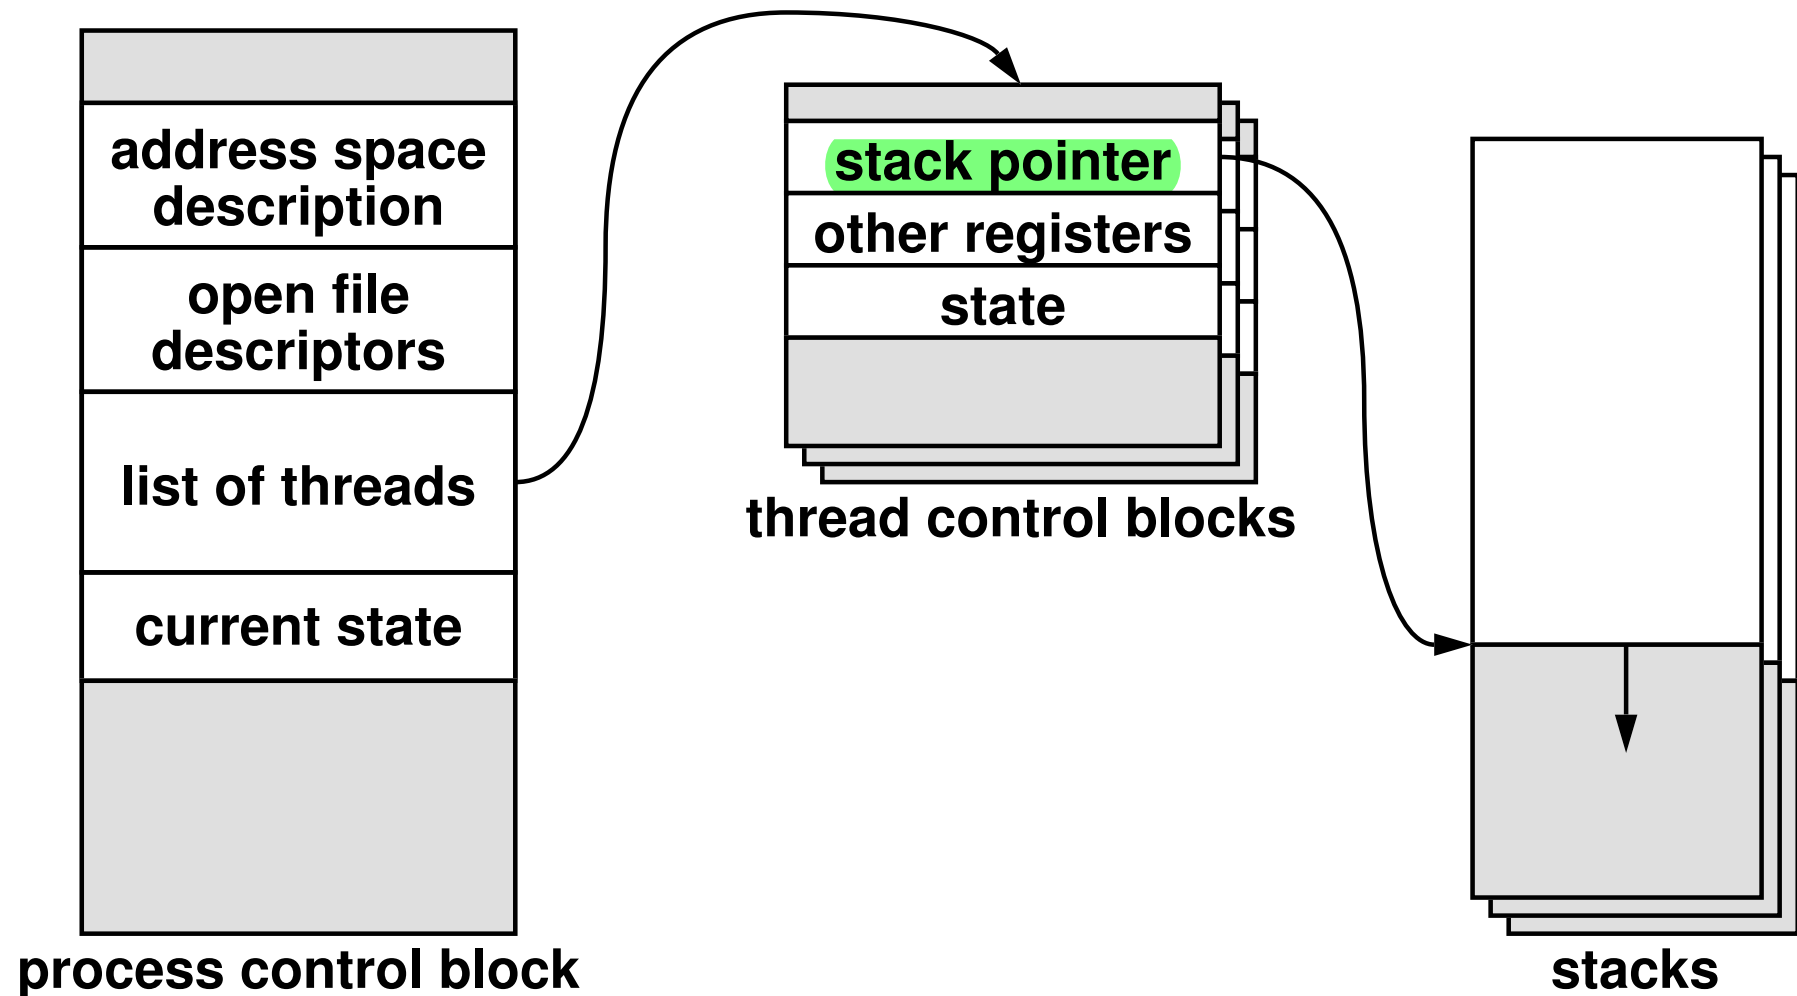


Processes and Threads

- ➡ A process is:
- ➡ a holder for an *address space*
 - ➡ a collection of other information shared by a set of *threads*
 - ➡ a collection of references to *open files* and other "execution context"
- ➡ As discussed in Ch 1, processes related APIs include
- ➡ `fork()`, `exec()`, `wait()`, `exit()`



Processes and Threads



Note: all these are relevant to your Kernel Assignment 1

