
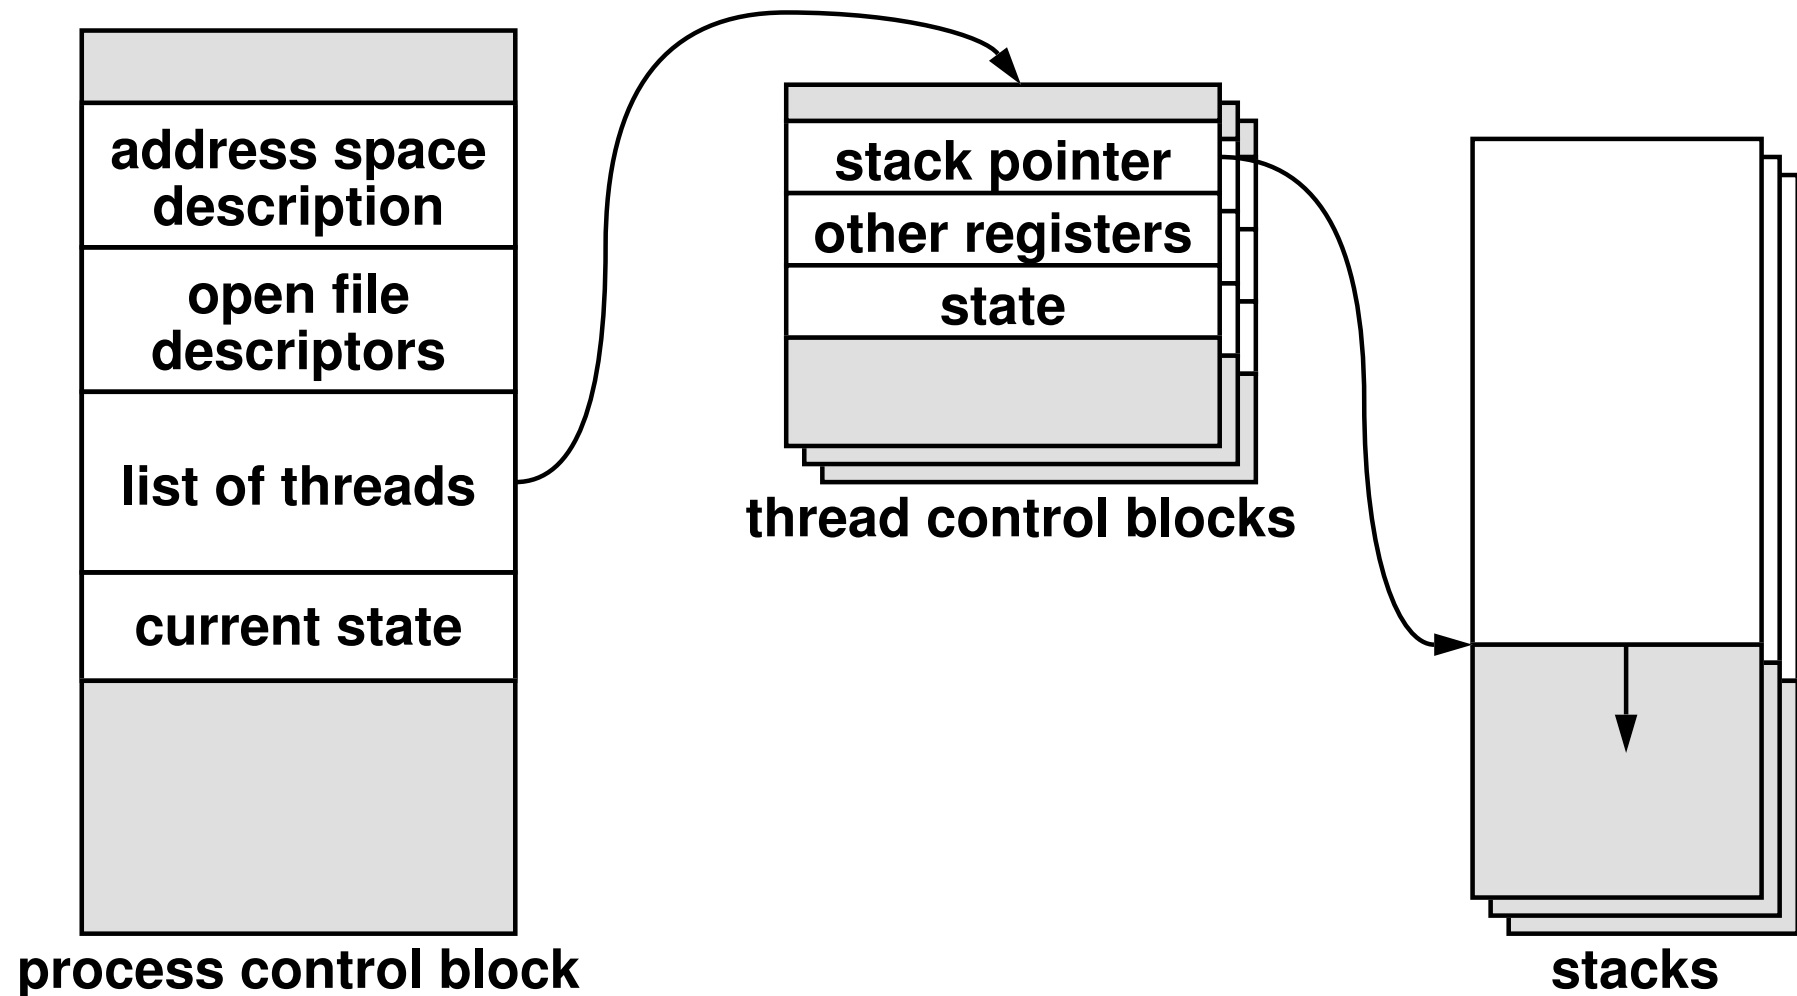


Housekeeping (Lecture 12 - 10/7/2013)

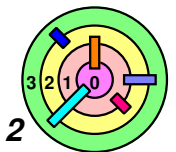
- ➡ Kernel #1 due at 11:45pm on Friday, 10/25/2013
 - if you have code from a previous semester, be very careful and *not copy any code from it*
 - it's best if you just get rid of it
- ➡ Any system issue, please get it resolved NOW
 - come to office hours to get help
- ➡ We did pretty well with kernel group forming
 - there is only one student in each section who does not have partners
 - I'm still waiting for the student to let me know if I should let the class know who they are so you can contact him/her directly
 - even if your team already has 4 students, you can add this student to your team!
- ➡ Post questions about the kernel assignments to class Google Group
 - extra credit for posting good responses



Processes and Threads





Note: all these are relevant to your Kernel Assignment 1



Process Life Cycle

➡ Pretty simple


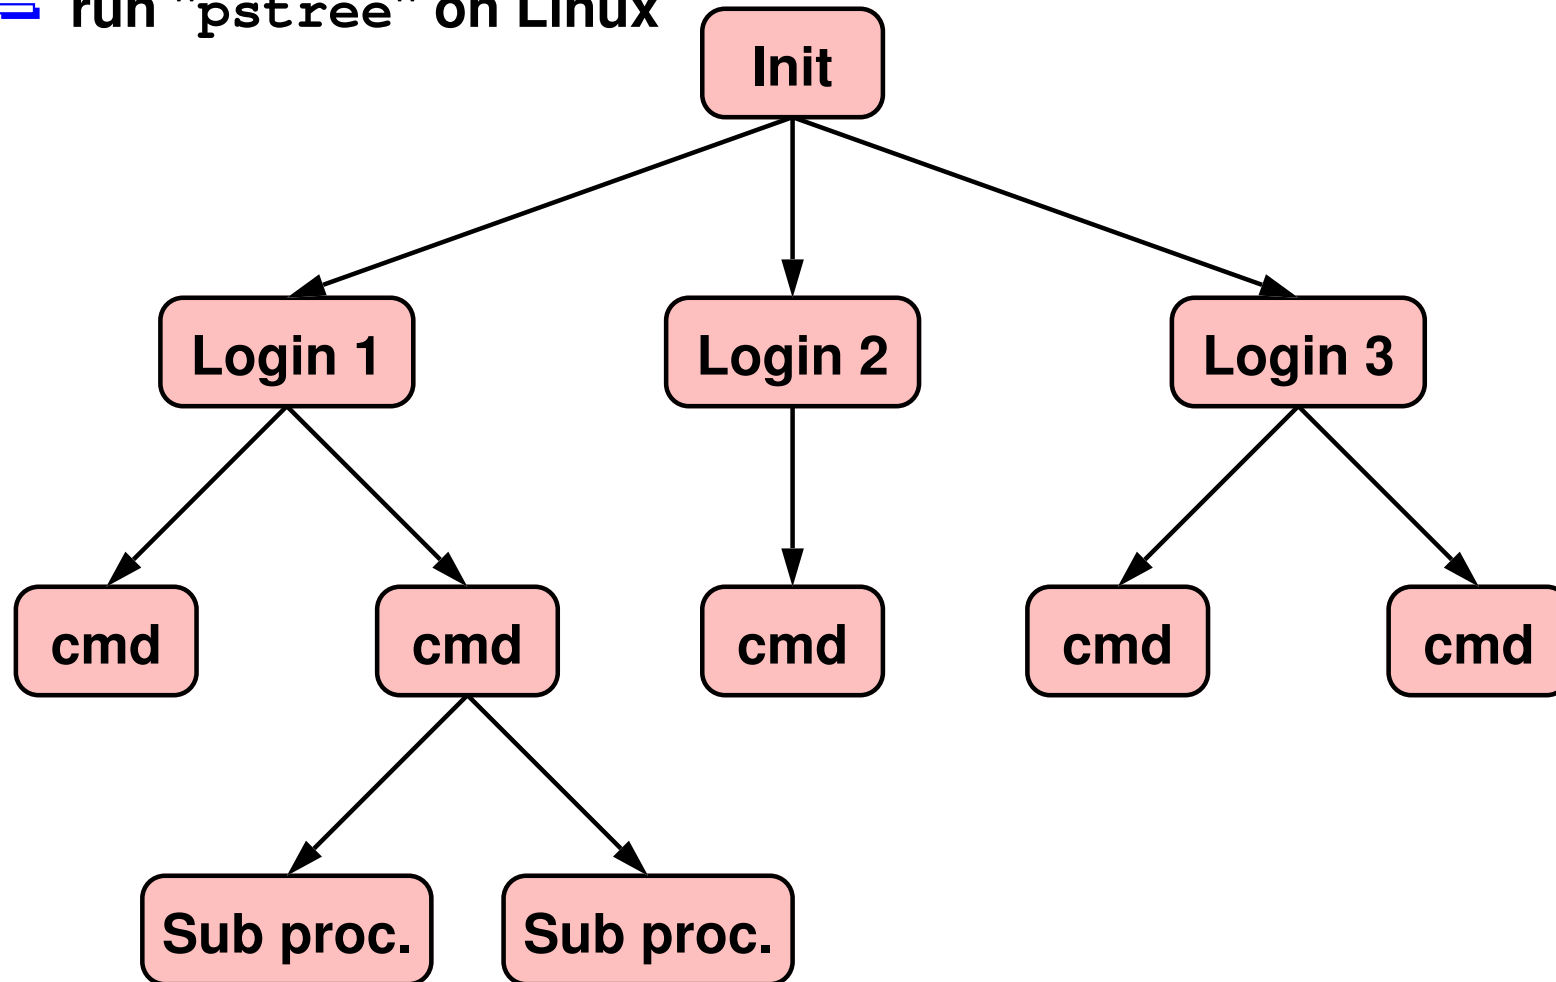


Process Relationships (1)

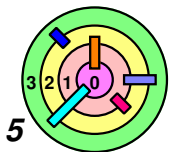
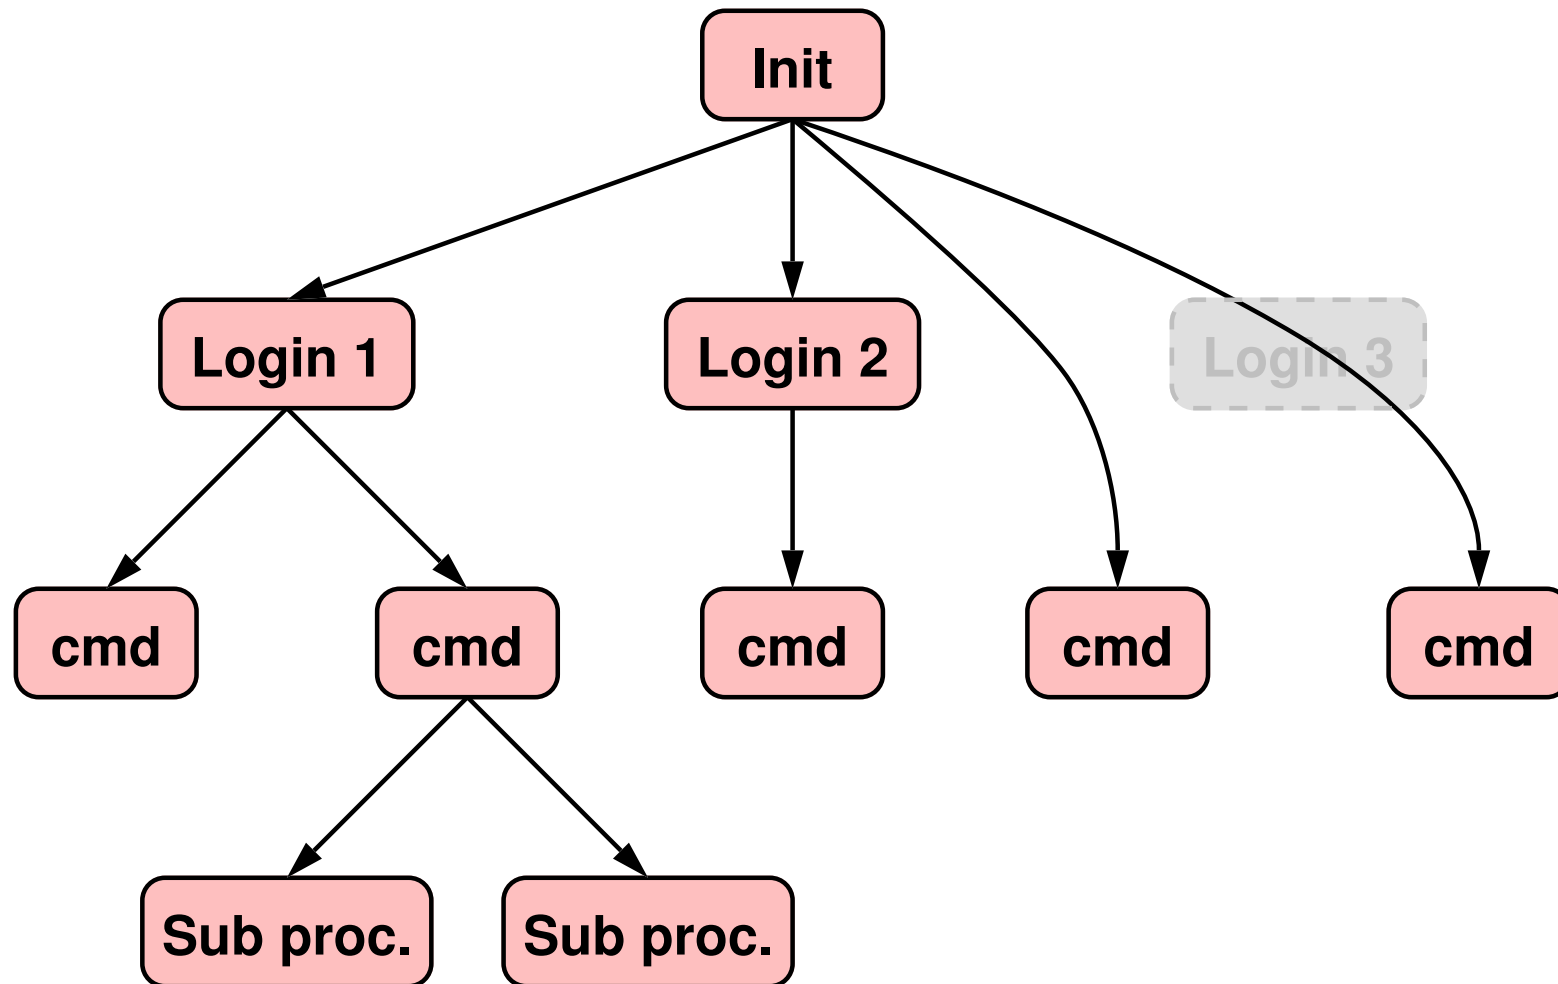


Process hierarchy

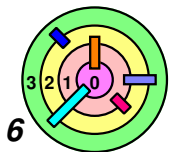
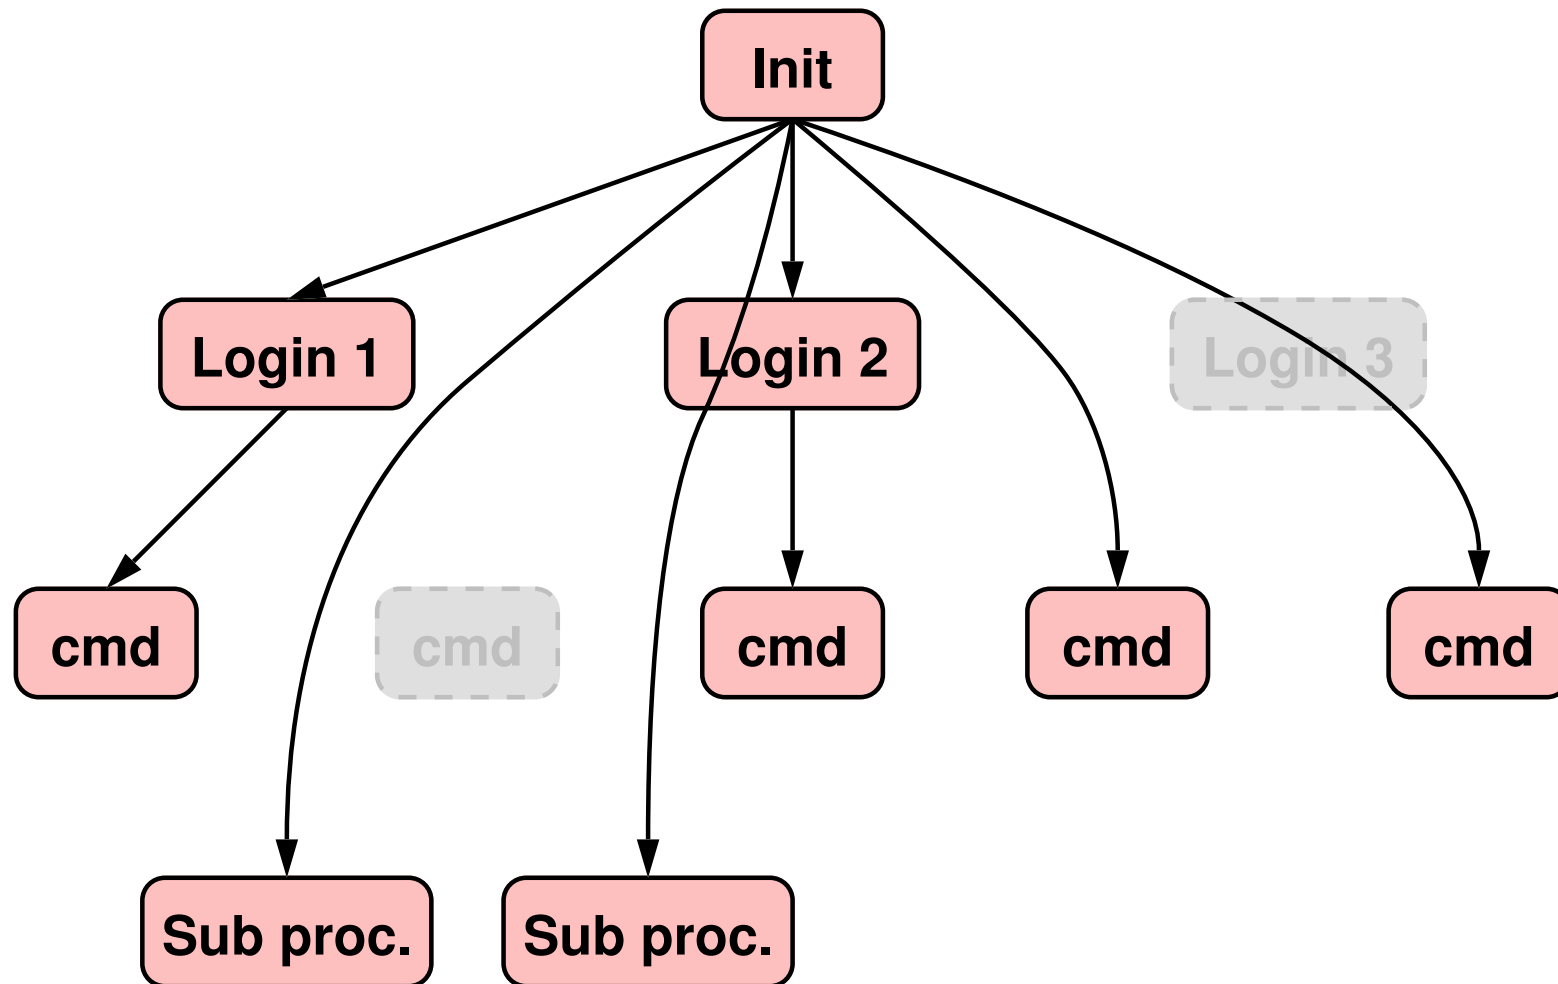
= run "pstree" on Linux




Process Relationships (2)



Process Relationships (3)



Fork and Threads



Or



Solaris uses the 2nd approach

— expensive to fork a process




Problem with 1st approach


— thread 1 called `fork()` and thread 2 has a mutex locked

○ who will unlock the mutex?


— **POSIX solution is to provide a way to unlock all mutex before `fork()`**




Thread Life Cycle




— a thread starts in the *runnable* state




Thread Life Cycle




- a thread starts in the *runnable* state
- the *scheduler* switches a thread's state from runnable to running



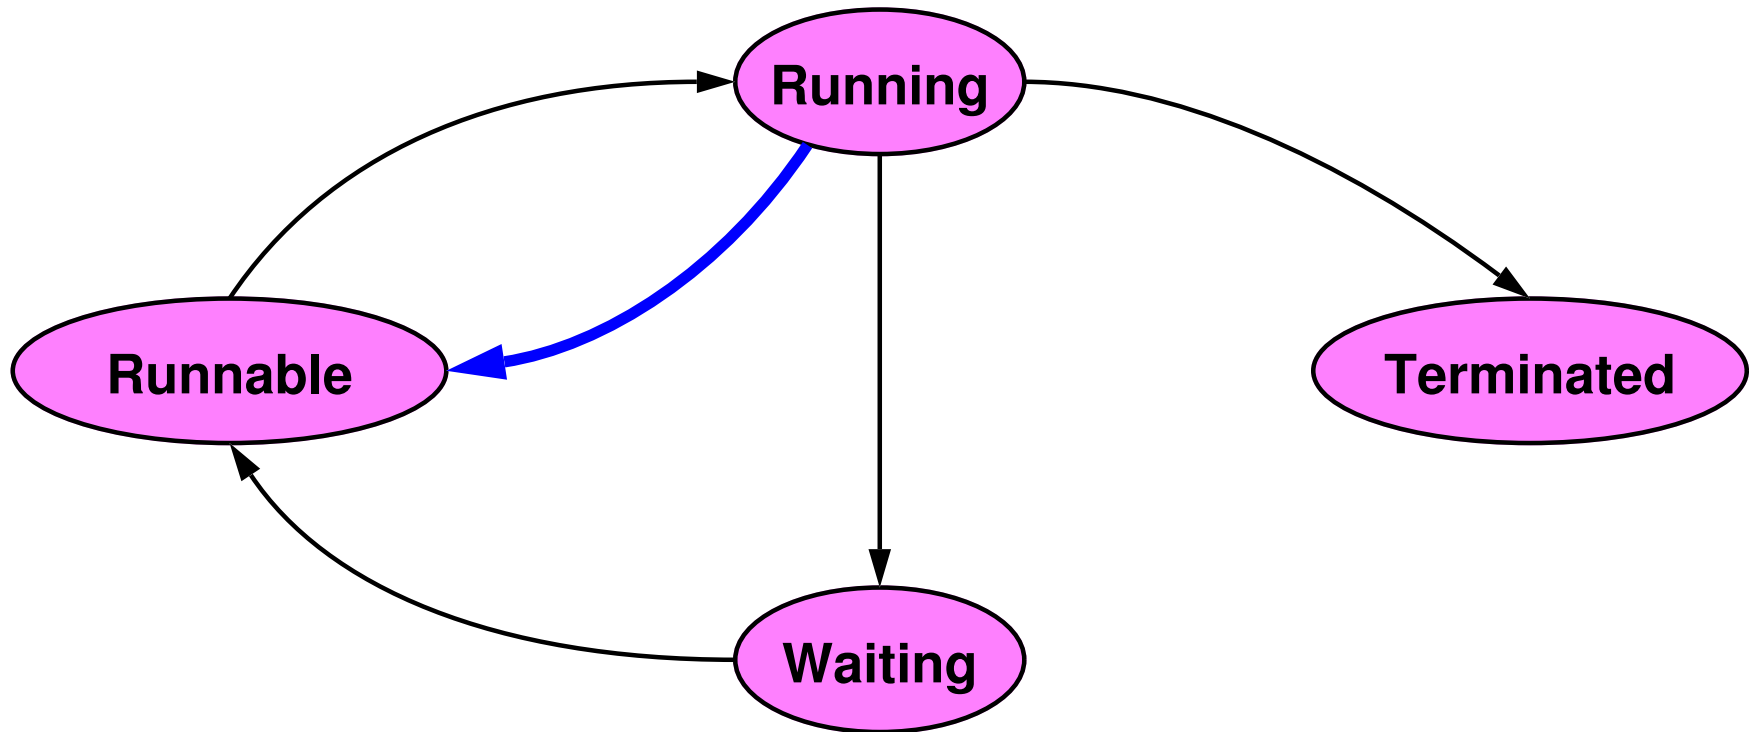
Thread Life Cycle




- a thread starts in the *runnable* state
- the *scheduler* switches a thread's state from runnable to running
- a thread goes from running to waiting when a *blocking call* is made




Thread Life Cycle




- a thread starts in the *runnable* state
- the *scheduler* switches a thread's state from runnable to running
- a thread goes from running to waiting when a *blocking call* is made
- the *scheduler* switches a thread's state from running to runnable when the thread used up its execution quantum




Thread Life Cycle




- a thread get *unblocked* by the action of another thread or by an interrupt handler



Thread Life Cycle




- a thread get **unblocked** by the action of another thread or by an interrupt handler
- in order for a thread to enter the terminated state, it **has to be in the running state just before that**
 - what if `pthread_cancel()` is invoked when the thread is not in the running state?




Thread Life Cycle

- ➡ Does `pthread_exit()` delete the thread (completely) that calls it?
 - no
- ➡ What's left in the thread after it calls `pthread_exit()`?
 - its thread control block
 - needs to keep thread ID and return code around
 - its stack
 - how can a thread *delete its own stack*? no way!




Thread Life Cycle

- ➡ Who is deleting the *thread control block* and freeing up the thread's *stack* space?
- ➡ If a thread is not detached
 - it can be taken care of in the `pthread_join()` code
 - the thread that calls `pthread_join()` does the clean up
- ➡ If a thread is detached (our simple OS does not support this)
 - can do this is one of two ways
 - 1) use a special *reaper thread*
 - ◆ basically doing `pthread_join()`
 - 2) queue these threads on a list and have other threads free them when it's convenient (e.g., when the scheduler schedule a thread to run)




4.1 A Simple System (Monolithic Kernel)

- ➡ A Framework for Devices
- ➡ Processes & Threads
- ➡ *Storage Management*
- ➡ Low-level Kernel (will come back to talk about this after Ch 7)




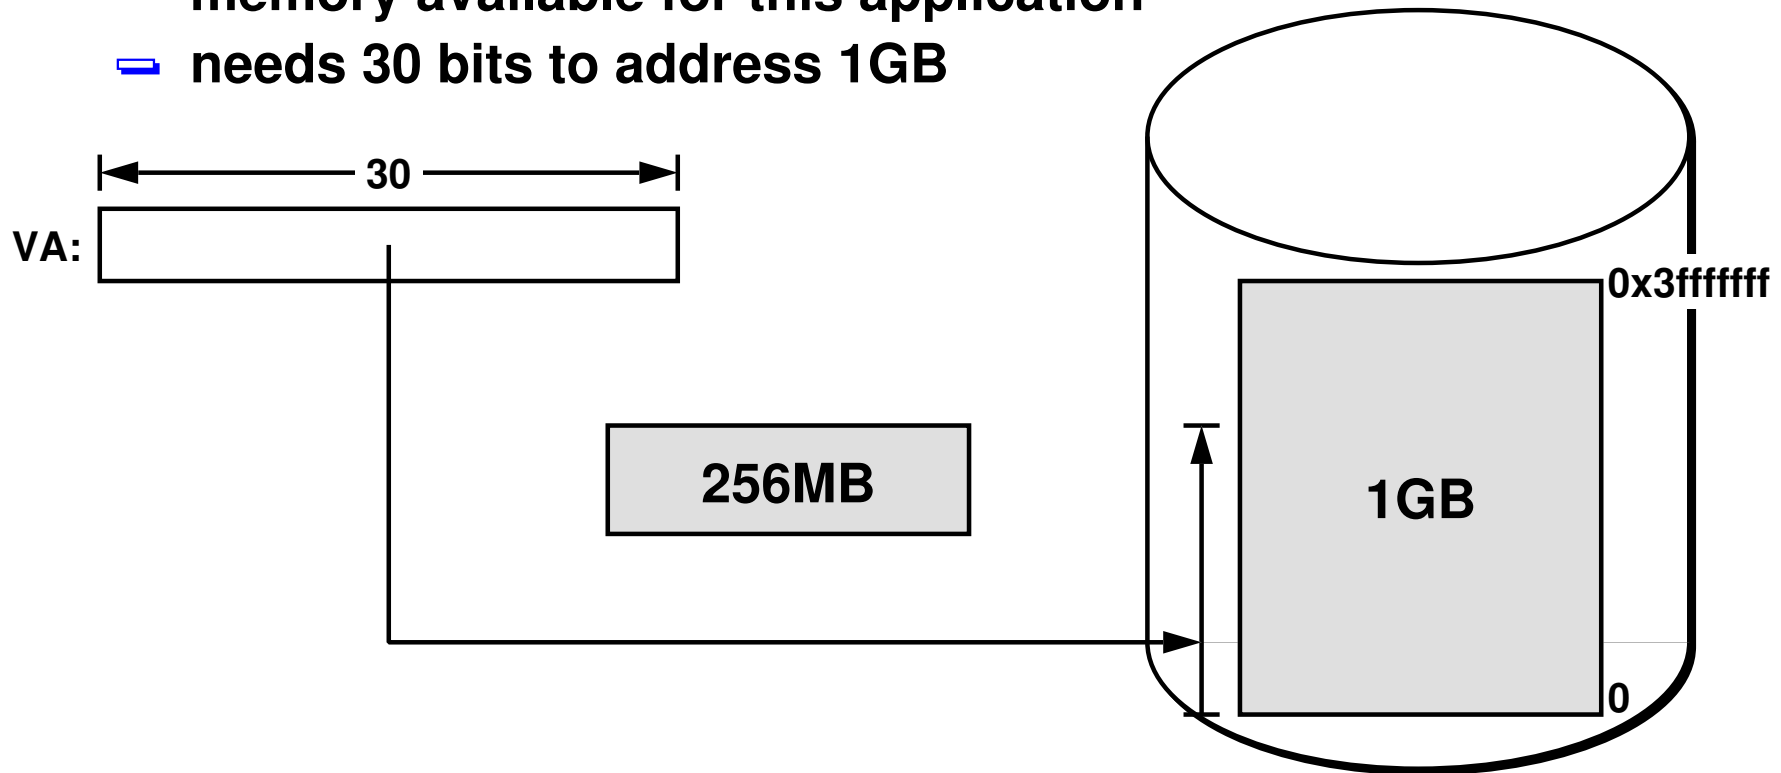
Storage Space

- ➡ Where to store data?
 - primary storage, i.e., physical memory
 - directly addressable
 - secondary storage, i.e., disk-based storage
- ➡ What would it take to support the idea of virtual memory, i.e., application's "view" of memory?
- ➡ An application only works with "virtual memory" (as far as an application is concerned, "virtual memory" is "real memory")
 - e.g., map a 1GB file into memory
 - this memory is *virtual memory*
 - can *allocate* 1GB of *virtual memory* while there's only 256MB of *physical memory*
 - the OS makes sure that real primary storage is available when necessary




Storage Space

- ➡ A simple example of virtual memory
- application needs 1GB but there is only 256MB of physical memory available for this application
 - needs 30 bits to address 1GB




Storage Space

- ➡ A simple example of virtual memory
- application needs 1GB but there is only 256MB of physical memory available for this application
 - needs 30 bits to address 1GB




- e.g., divide 1GB into 4 *pages*, 256MB each




Storage Space

- ➡ A simple example of virtual memory
- application needs 1GB but there is only 256MB of physical memory available for this application
 - needs 30 bits to address 1GB

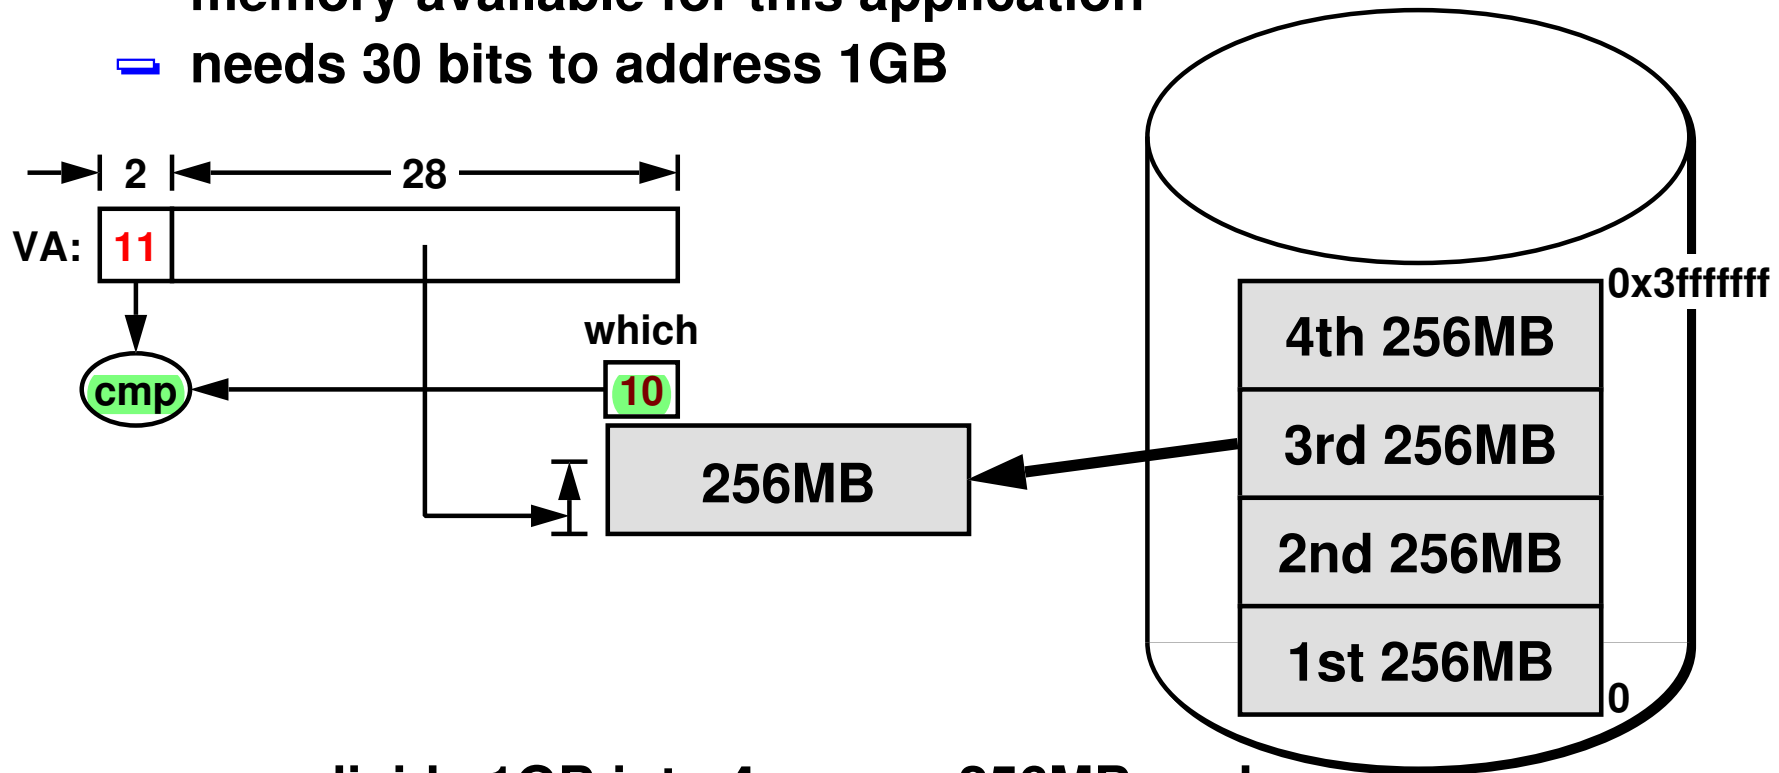


- e.g., divide 1GB into 4 *pages*, 256MB each
- the *first 2 bits* in the *virtual address* tell you which *page*
- the rest of the bits give you the *offset* within the *page*




Storage Space

- ➡ A simple example of virtual memory
- application needs 1GB but there is only 256MB of physical memory available for this application
 - needs 30 bits to address 1GB




- e.g., divide 1GB into 4 *pages*, 256MB each
- the *first 2 bits* in the *virtual address* tell you which *page*
- the rest of the bits give you the *offset* within the *page*
- **check to see if the right page is in *physical* memory**




Storage Space

- ➡ A simple example of virtual memory
 - ▬ application needs 1GB but there is only 256MB of physical memory available for this application
 - ▬ needs 30 bits to address 1GB




- ➡ If it's the wrong page that's in primary memory
 - ▬ accessing it will cause a *page fault*
 - ▬ during a page fault, OS brings the right page into real memory
 - ▬ then the thread is allow to proceed with accessing the memory




Storage Space

- ➡ A simple example of virtual memory
 - ▬ application needs 1GB but there is only 256MB of physical memory available for this application
 - ▬ needs 30 bits to address 1GB

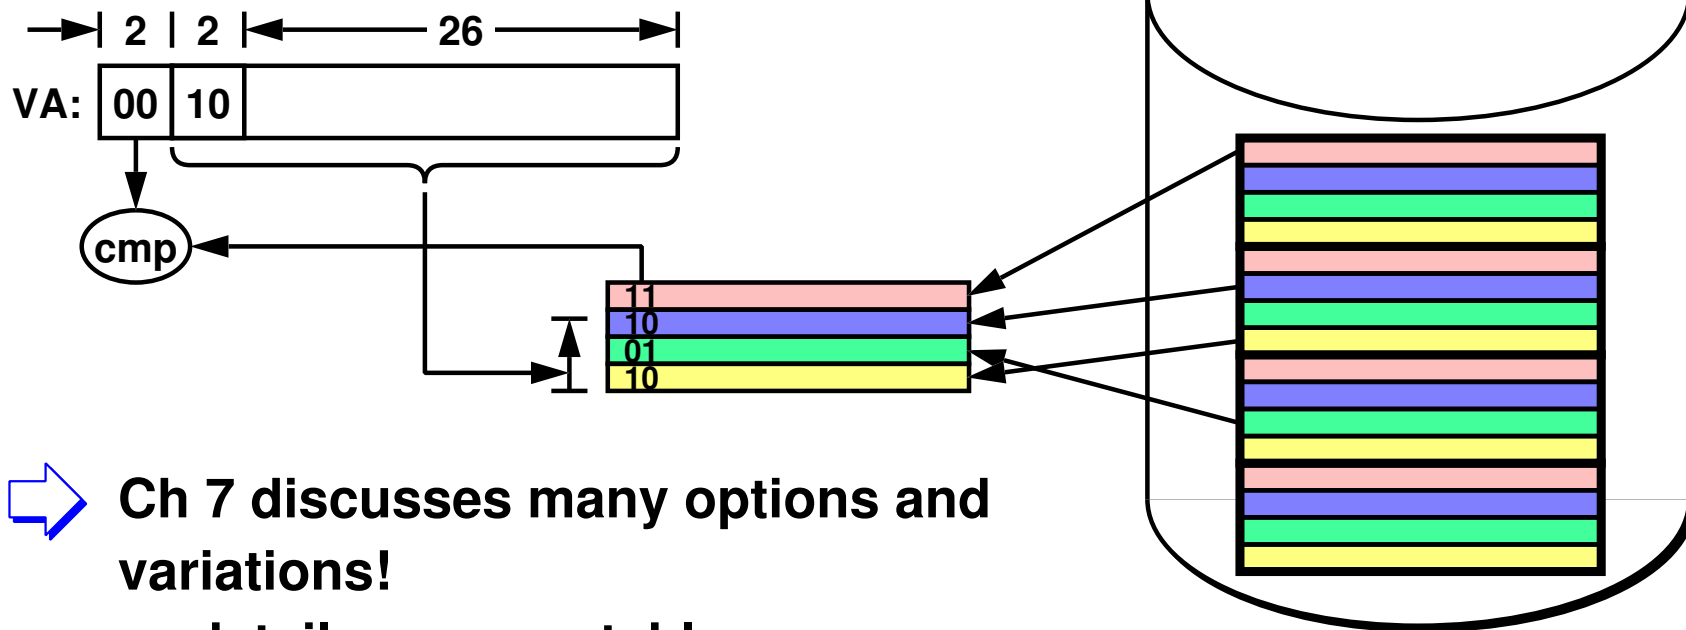


- ➡ If it's the wrong 256MB that's in primary memory
 - ▬ accessing it will cause a *page fault*
 - ▬ this "simple" approach has *really poor performance*
 - why just use 2 leading bits? different organizations?



Storage Space

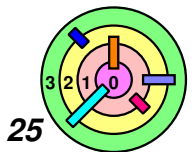
- ➡ A more complicated scheme with a smaller page size
 - compare to determine if there is a *hit* or not
 - can have even smaller page sizes



- ➡ Ch 7 discusses many options and variations!
 - details on page tables, translation look-aside buffers, etc.


Memory Management Concerns

- ➡ **Mapping** virtual addresses to real ones
- ➡ Determining which addresses are **valid**, i.e., refer to allocated memory, and which are not
- ➡ Keeping track of which real objects, if any, are mapped into each range of virtual addresses
- ➡ Deciding what should to keep in primary storage (RAM) and what to fetch from elsewhere




Segmentation Fault

- ➡ A valid virtual address must be ultimately *resolvable* by the OS to a location in the physical memory
- ➡ if it cannot be resolved, the virtual address is considered an *invalid* virtual address
 - ➡ referencing an invalid virtual address will cause a **segmentation fault** (the OS will deliver SIGSEG to the process)
 - the default action would be to terminate the process

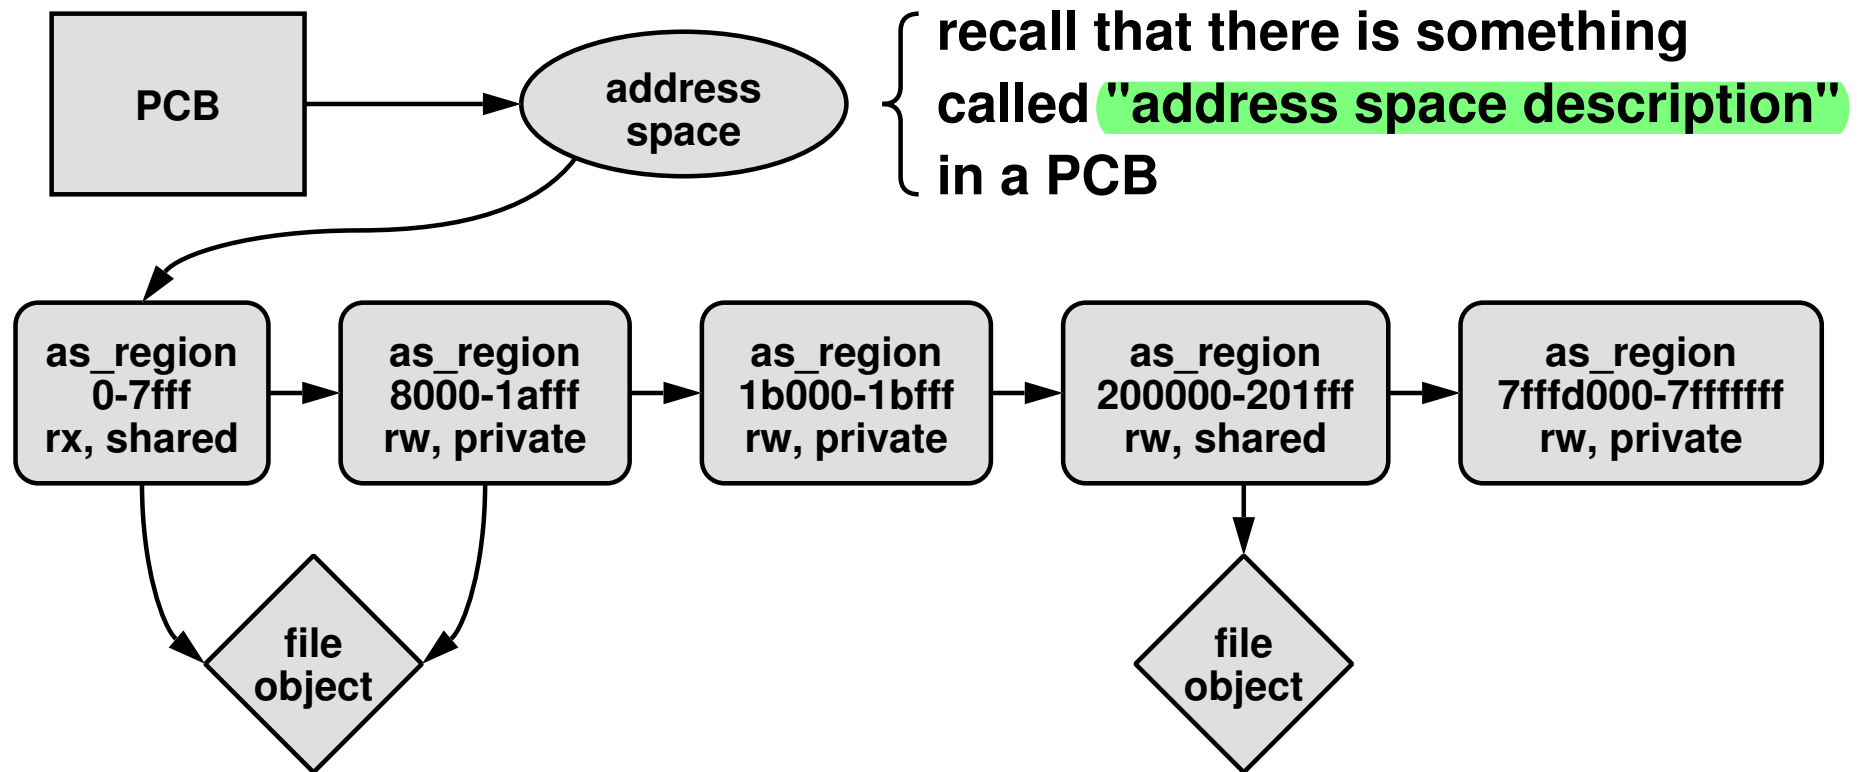


Hardware Memory Map

- ➡ In reality, the OS is too slow since *every* virtual address needs to be resolved
- some of the virtual memory mechanisms must be built into the *hardware*
 - in some cases, the hardware is given the complete *"map"* (i.e., mapping from virtual to physical address)
 - in some other cases, only a partial map is given to the hardware
 - in either case, OS needs to provide some map to the hardware and needs a *data structure* for the map
 - ◆ often referred as the *memory map*, or *mmap*




Address Space Representation

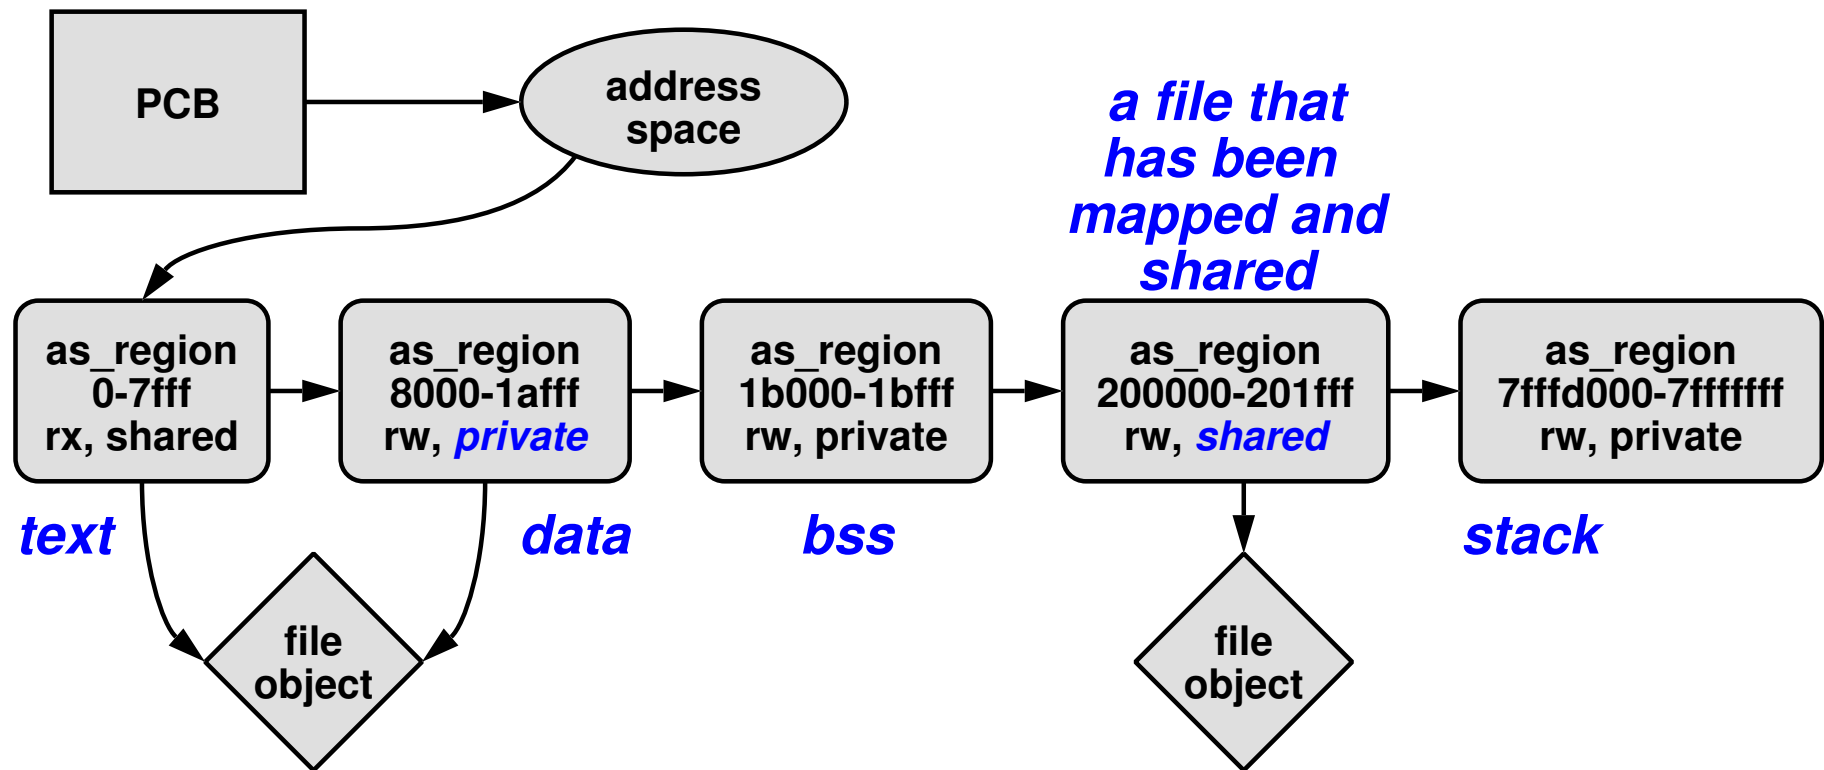


- ➡ **as_region** (address space region data structure) contains:
- = *start address*, *length*, *access permissions*, *shared* or *private*
 - = if mapped to a file, pointer to the corresponding file object


- ➡ This is related to Kernel Assignment 3 where you need to create and manage *address spaces* / *memory maps*



Address Space Representation




- ➡ In this example, text and data map portions of the same file
- ➡ **text** is marked read-execute and **shared**
 - ➡ data is marked read-write and **private** to mean that changes will be private, i.e., **will not affect other processes exec'ed from the same file**



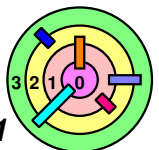
How OS Makes Virtual Memory Work?

- ➡ If a thread access a virtual memory location that's both in primary memory and mapped by the hardware's map
 - no action by the OS
- ➡ If a thread access a virtual memory location **that's not in primary memory or if the translation is not in the map**
 - a **fault** is occurred and the OS is invoked
 - OS **checks the `as_region` data structures to make sure the reference is valid**
 - ◆ if it's valid, the OS does whatever that's necessary to locate or create the object of the reference
 - ◆ find, or if necessary, make room for it in primary storage if it's not already there, and put it there
 - ◆ details in Ch 7
- ➡ Two issues need further discussion
 - how is the **primary storage** managed?
 - how are these objects managed in **secondary storage**?



How Is The Primary Storage Managed?

- ➡ Who needs primary memory?
 - application processes
 - terminal-handling subsystem
 - communication subsystem
 - I/O subsystem
- ➡ They *compete* for available memory
 - it's difficult to be "fair" (what does it even mean?)
- ➡ If primary memory is managed poorly
 - one subsystem can use up all the available memory
 - then other subsystem won't get to run
 - this many lead to OS crash when a subsystem runs out of memory
- ➡ If there are no mapped files, the solution can be simple
 - equally divide the primary memory among the participants
 - this way, they won't compete



In Reality, Have To Deal With Mapped Files



An example to demonstrate a dilemma

- one process is using all of its primary storage allocation
- it then maps a file into its address space and starts accessing that file
- should the memory that's needed to buffer this file be charged against the files subsystem or charged against the process?




If charged against the files subsystem

- if the newly mapped file takes up all the buffer space in the files subsystem, it's unfair to other processes



If charged against the process

- if other processes are sharing the same file, other processes are getting a free ride (in terms of memory usage)
- even worse, another process may increase the memory usage of this process (double unfair!)



In Reality, Have To Deal With Mapped Files



It's difficult to be *fair*


— it's difficult to even define what *fair* means



We will discuss some solutions in Ch 7


— for now, we use the following solution

- give each participant (processes, file subsystem, etc.) a minimum amount of storage
- leave some additional storage available for all to compete

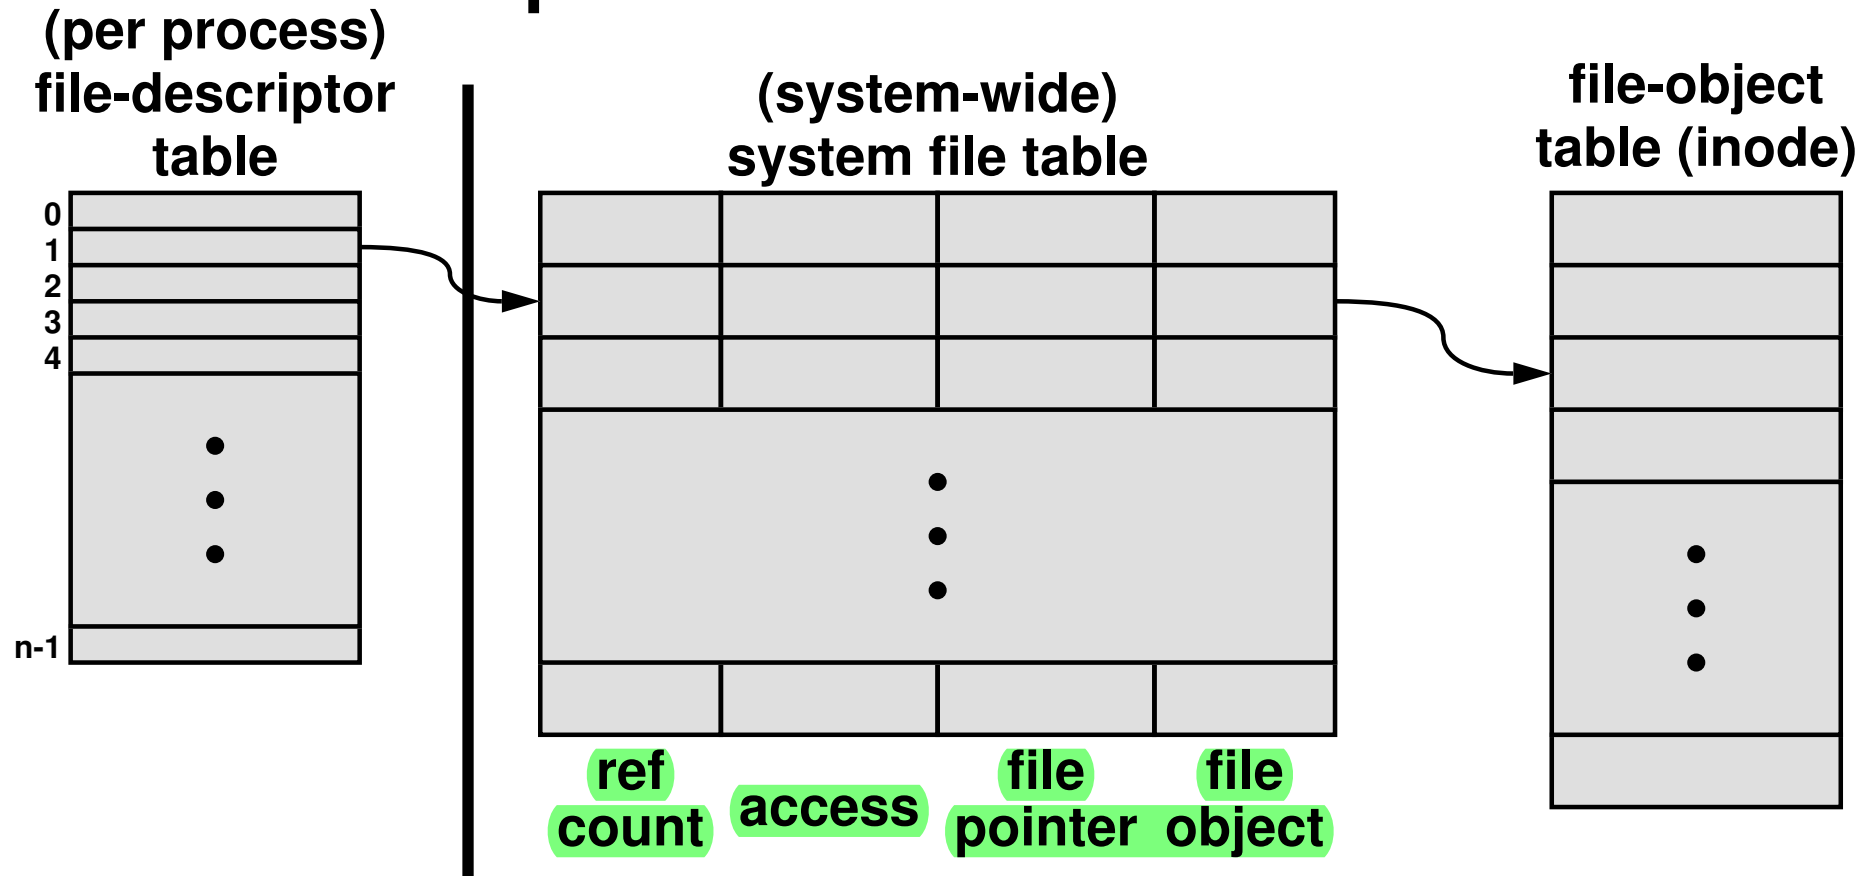


How Are Objects Managed In Secondary Storage?


- ➡ The *file system* is used to manage objects in secondary storage
- ➡ The file system is usually divided into two parts
 - ⇒ *file system independent*
 - supports the "file abstraction"
 - on Windows, this is called the "*I/O manager*"
 - on Unix, this is called the "*virtual file system (VFS)*"
 - ◆ Kernel Assignment 2
 - ⇒ *file system dependent*
 - on Windows, this is called the "file system"
 - on Unix, this is called the "actual file system"



Open-File Data Structures



- ➡ Each process has its own file-descriptor table
 - system file table and file-object table belongs to the kernel
- ➡ The *file object* forms the boundary between *VFS* and the actual file system (i.e., will point to device-dependent stuff)




File Object

- ➡ The file object is like an *abstract class* in C++
- subclasses of file object are the *actual* file objects

```
class FileObject {  
    unsigned short refcount;  
    ...  
    virtual int create(const char *, int, FileObject **);  
    virtual int read(int, void *, int);  
    virtual int write(int, const void *, int);  
    ...  
};
```

- ➡ But wait ...
- what's this about C++?
 - real operating systems are written in C ...
 - checkout the DRIVERS kernel documentation (we skipped this weenix assignment)




File Object in C

```
typedef struct {  
    unsigned short refcount;  
    struct file_ops *file_op;  
    /* function pointers */  
} FileObject;
```




- ➡ A file object uses an *array of function pointers*
 - this is how C implements *C++ polymorphism*
 - one for each operation on a file
 - where they point to is (actual) file system dependent
 - but the (virtual) interface is the same to higher level of the OS
- ➡ Loose coupling between the actual file system and storage devices
 - the actual file system is written to talk to the devices in a device-independent manner
 - i.e., using major and minor device numbers to reference the device and using standard interface provided by the device driver



File System Cache


- ➡ **Recently used blocks** in a file are kept in a **file system cache**
 - the primary storage holding these blocks might be mapped into one or more address spaces of processes that have this file mapped
 - blocks are available for immediate access by read and write system calls
- ➡ A simple **hash function** is used to locate file blocks in the cache
 - keyed by **inode number**



Ch 5: Processor Management


Bill Cheng

<http://merlot.usc.edu/cs402-f13>



Processor Management


- ➡ Threads *Implementation*
 - lock/mutex implementation on multiprocessors
- ➡ Interrupts
- ➡ Scheduling
- ➡ Linux/Windows Scheduler



5.1 Threads


Implementations

- ⇒ *Strategies*
- ⇒ A Simple Thread Implementation
- ⇒ Multiple Processors




Threads Implementation


- ➡ The ultimate goal of the OS is to support user-level applications
 - we will discuss various strategies for supporting threads
- ➡ Where are operations on threads implemented?
 - in the kernel?
 - or in user-level library?
- ➡ Approaches
 - one-level model (threads are implemented in the kernel)
 - variable-weight processes
 - two-level model (threads are implemented in user library)
 - $N \times 1$
 - $M \times N$
 - scheduler activations model




One-Level Model



User




Kernel



Processors

One-Level Model

- ➡ The simplest and most direct approach is the *one-level model*
 - all aspects of the *thread implementation* are *in the kernel*
 - i.e., all thread routines (e.g., `pthread_mutex_lock`) called by user code are all system calls
 - each *user thread* is mapped one-to-one to a *kernel thread*
- ➡ If a thread calls `pthread_create()`
 - it's a system call, so it traps into the kernel
 - the kernel creates a thread control block
 - associate it with the process control block
 - the kernel creates a kernel and a user stack for this thread
- ➡ What about `pthread_mutex_lock()`
 - why does it have to be done in the kernel?
 - it's not necessary to protect the threads from each other!
 - you definitely don't need the kernel to protect threads from each other




One-Level Model




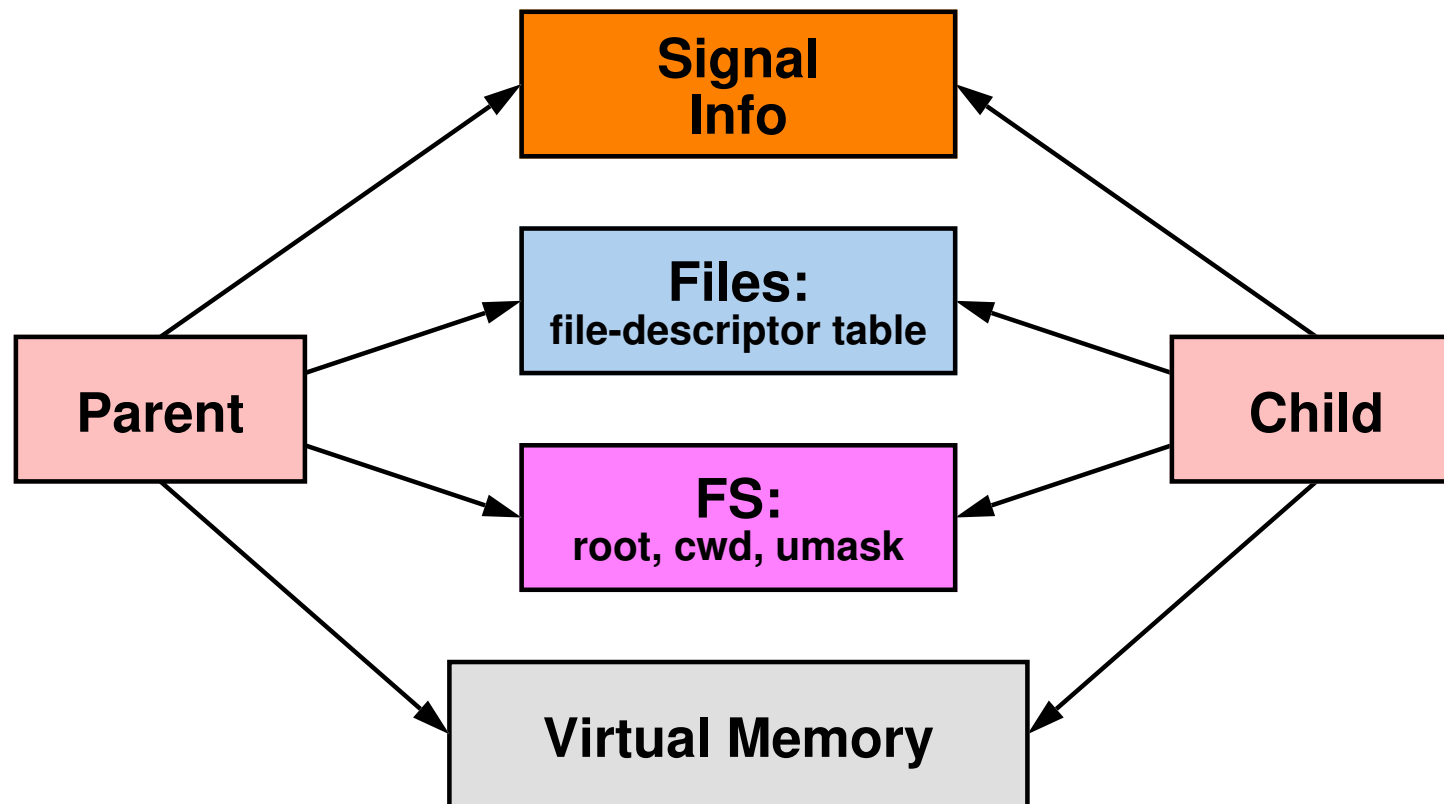
Problem: system calls are expensive

- if `pthread_mutex_lock` finds the mutex available, it should return quickly (and lock the mutex)**
 - if this can be done in user code, it can be 20 times faster (for the case where the mutex is available)**
 - in Win32 threads, an equivalent of a mutex is represented in a user-level data structure**
 - ◇ if such an object is not locked, it returns quickly**
 - ◇ if such an object is locked, it makes a system call and blocks in the kernel**




Variable-Weight Processes

- ➡ Variant of one-level model
- ➡ Portions of parent process selectively *copied* into or *shared* with child process
- ➡ Children created using `clone()` system call



Linux Threads (pre 2.6)




NPTL in Linux 2.6




Native POSIX-Threads Library

- full POSIX-threads semantics on improved variable-weight processes
- threads of a "process" form a *thread group*
 - `getpid()` returns process ID of first thread in group
 - any thread in group can wait for any other to terminate
 - signals to process delivered by kernel to any thread in group
- new kernel-supported synchronization construct: *futex* (fast mutex)
 - used to implement mutexes, semaphores, and condition variables




Two-Level Model

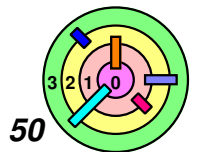
- ➡ In the two-level model, **a user-level library plays a major role**
 - what an user-level application perceives as a thread is implemented within user-level library code
- ➡ Two versions
 - single kernel thread (per user process)
 - multiple kernel threads (per user process)



Two-Level Model - One Kernel Thread




- ➡ This is one of the earliest ways of implementing threads
- threads are implemented entirely in the user level
 - thread control block, mutex in user space
 - thread stack allocated by user library code
 - mostly done on uniprocessors




Two-Level Model - One Kernel Thread

- ➡ Within a process, *user threads* are *multiplexed* not on the processor, but on a *kernel-supported thread*
 - ➡ the *OS multiplexes kernel threads* (or equivalently, processes) on the *processor*
- ➡ User thread creation
 - ➡ a stack and a thread control block is allocated
 - ➡ thread is put on a queue of runnable threads
 - wait for its turn to become the running thread
- ➡ Synchronization implementation
 - ➡ relative straightforward
 - ➡ e.g., mutex (one queue per mutex)
 - if a thread must block, it simply queues itself on a wait queue and calls context-switch routine to pass control to the first thread on the runnable queue



Two-Level Model - One Kernel Thread

- ➡ This is called the N-to-1 model
- ➡ Major advantage
 - no system calls (for thread-related APIs)!
- ➡ Major disadvantage
 - what if a thread makes a system call (for a non-thread-related API)?
 - it gets blocked in the kernel
 - no other user thread in the process can run



Coping ...

```

ssize_t read(int fd, void *buf, size_t count)
{
    ssize_t ret;
    while (1) {
        if ((ret = real_read(fd, buf, count)) == -1) {
            if (errno == EWOULDBLOCK) {
                sem_wait(&FileSemaphore[fd]);
                continue;
            }
        }
        break;
    }
    return(ret);
}

```




Solution is to have a non-blocking read() called real_read()

➤ real_read() either returns immediately with data in buf

➤ or returns immediately with an error code in errno


- EWOULDBLOCK means that a real read() would block, i.e., data is not ready to be read




Coping ...

```
ssize_t read(int fd, void *buf, size_t count)
{
    ssize_t ret;
    while (1) {
        if ((ret = real_read(fd, buf, count)) == -1) {
            if (errno == EWOULDBLOCK) {
                sem_wait(&FileSemaphore[fd]);
                continue;
            }
        }
        break;
    }
    return(ret);
}
```


- ➡ One semaphore for each open file
 - ➡ perhaps a signal handler will invoke `sem_post()` to when data is ready to be read
- ➡ Major drawback
 - ➡ only works for some I/O objects - not a general solution




Two-Level Model: Multiple Kernel Threads



- ➡ This is called the M-to-N model
- ➡ Implementation is similar to the two-level model with a single kernel thread
 - no system calls (for thread-related APIs)
 - if we don't have enough kernel threads per user process, we end up having the same problem with the N-to-1 model




Deadlock




➡ **Ex: two threads are communicating using a pipe (this is essentially a kernel implementation of the producer-consumer problem)**

- ➡ first user thread writes to a full pipe and get blocked in the kernel
 - first thread just happened to use the last kernel thread
 - 2nd thread wants to read the pipe to unblock the first thread, but cannot because no kernel thread left




Deadlock



- ➡ **Solaris solution: automatically create a new kernel thread**
= an obvious solution


Recap - Problems

- ➡ Two-level model **does not solve the I/O blocking problem**
 - if there are N kernel threads and if N user threads are blocked in I/O
 - no other user threads can make progress
- ➡ Another problem: *Priority Inversion*
 - user-level thread schedulers are not aware of the kernel-level thread scheduler
 - it may know the number of kernel threads
 - how can the user-level scheduler talk to the kernel-level scheduler?
 - people have tried this, but it's complicated
 - it's possible to have a higher priority user thread scheduled on a lower priority kernel thread and vice versa

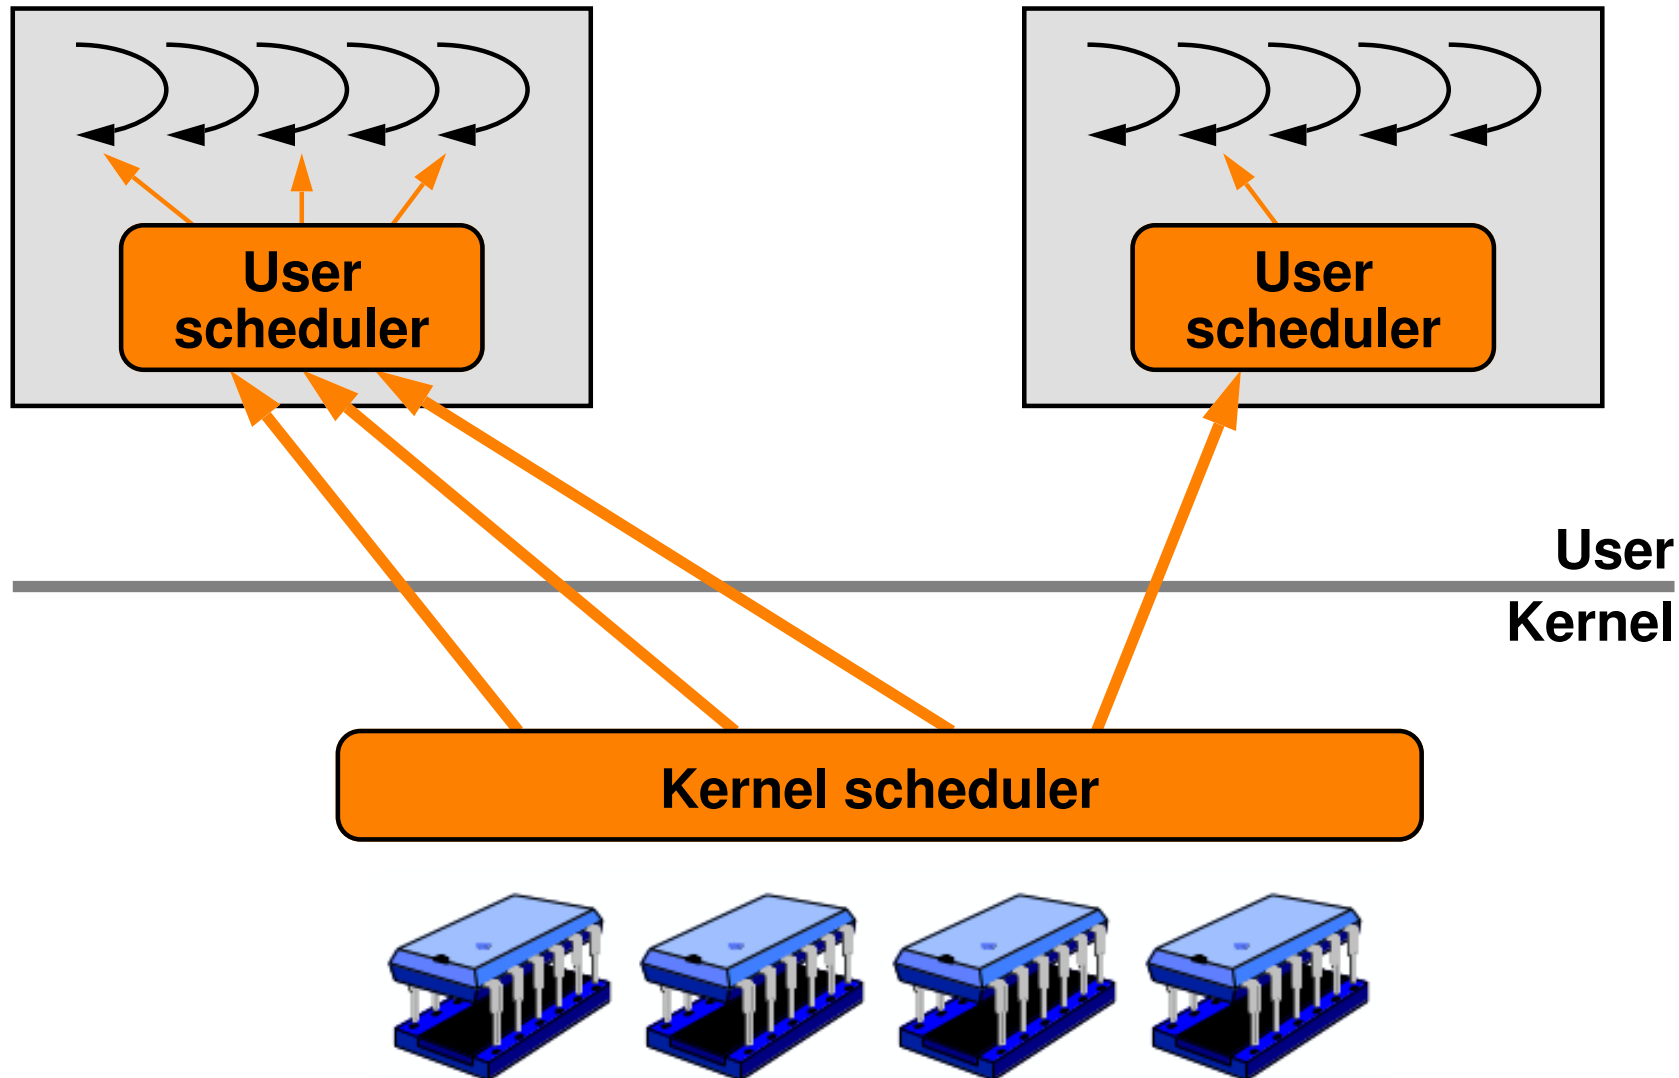


Scheduler Activations Model

- ➡ The scheduler activations model is radically different from the other models
- in other models, we think of the kernel as providing some kernel thread contexts
 - then multiplexing these contexts on processors using the kernel's scheduler
 - in scheduler activations model, we divvy up processors to processes, and processes determine which threads get to use these processors
 - the kernel should supply however many kernel contexts it finds necessary




Scheduler Activations Model Example

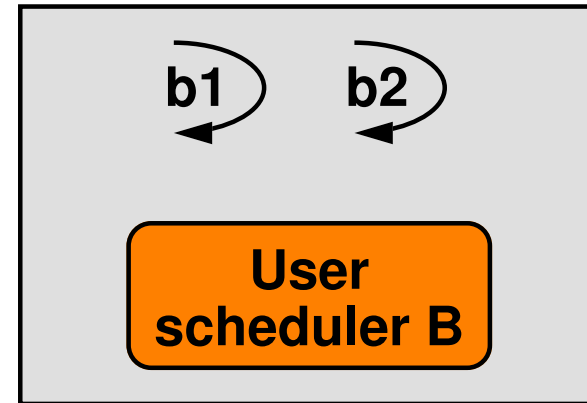
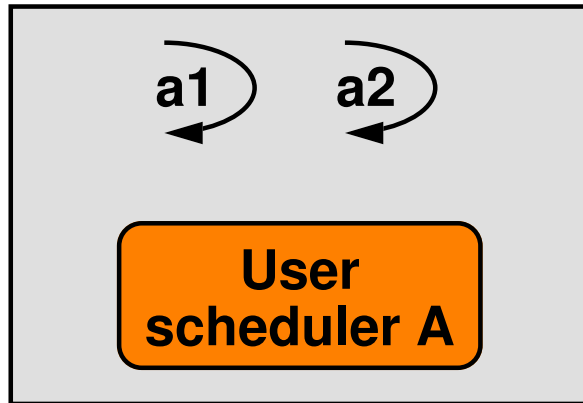


Scheduler Activations Model Example

- ➡ Let's say a process starts up running a single thread
 - kernel scheduler assigns a processor to the process
 - if the thread blocks, the process gives up the processor to the kernel scheduler
- ➡ Suppose the user program creates a new thread and parallelism is desired
 - code in user-level library notifies the kernel that it needs two processors
 - when a processor becomes available, the kernel creates a new kernel context
 - the kernel places an upcall to the user-level library, effectively giving it the processor
 - the user-level library code assigns this processor to the new thread




Scheduler Activations Model Example




User
Kernel

Kernel scheduler

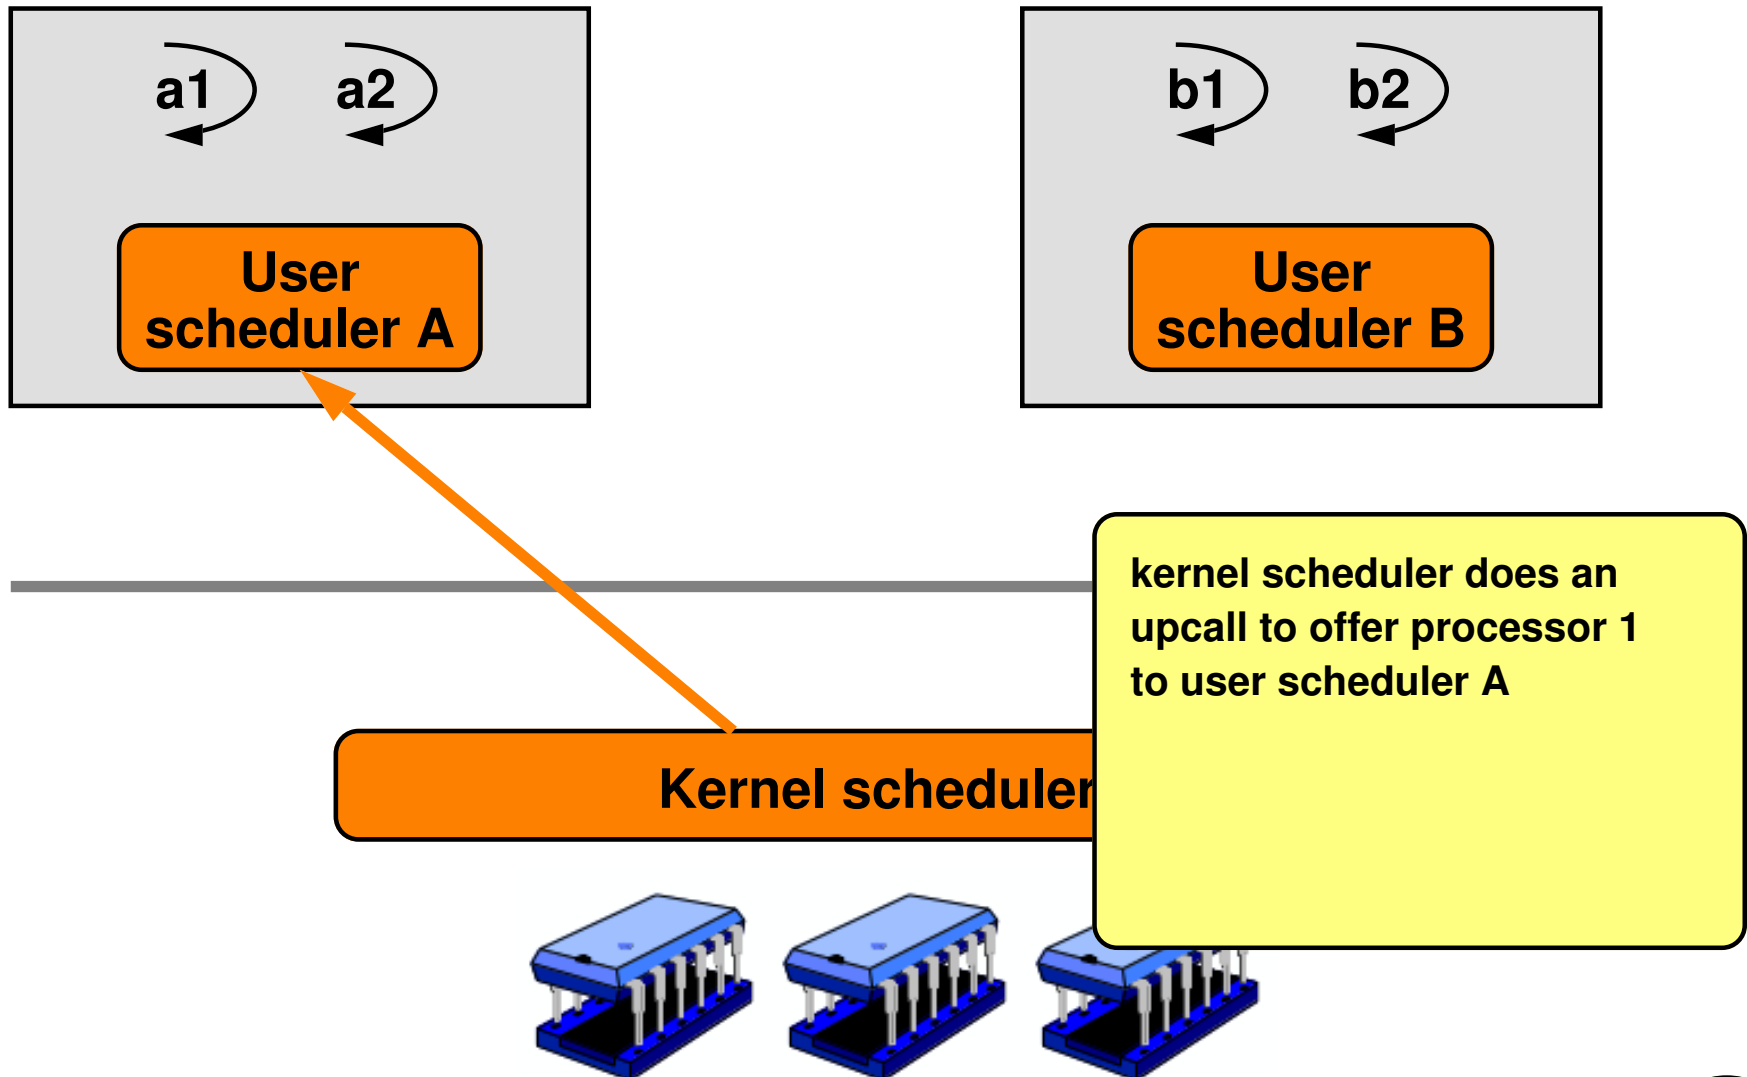


Kernel scheduler does not schedule threads



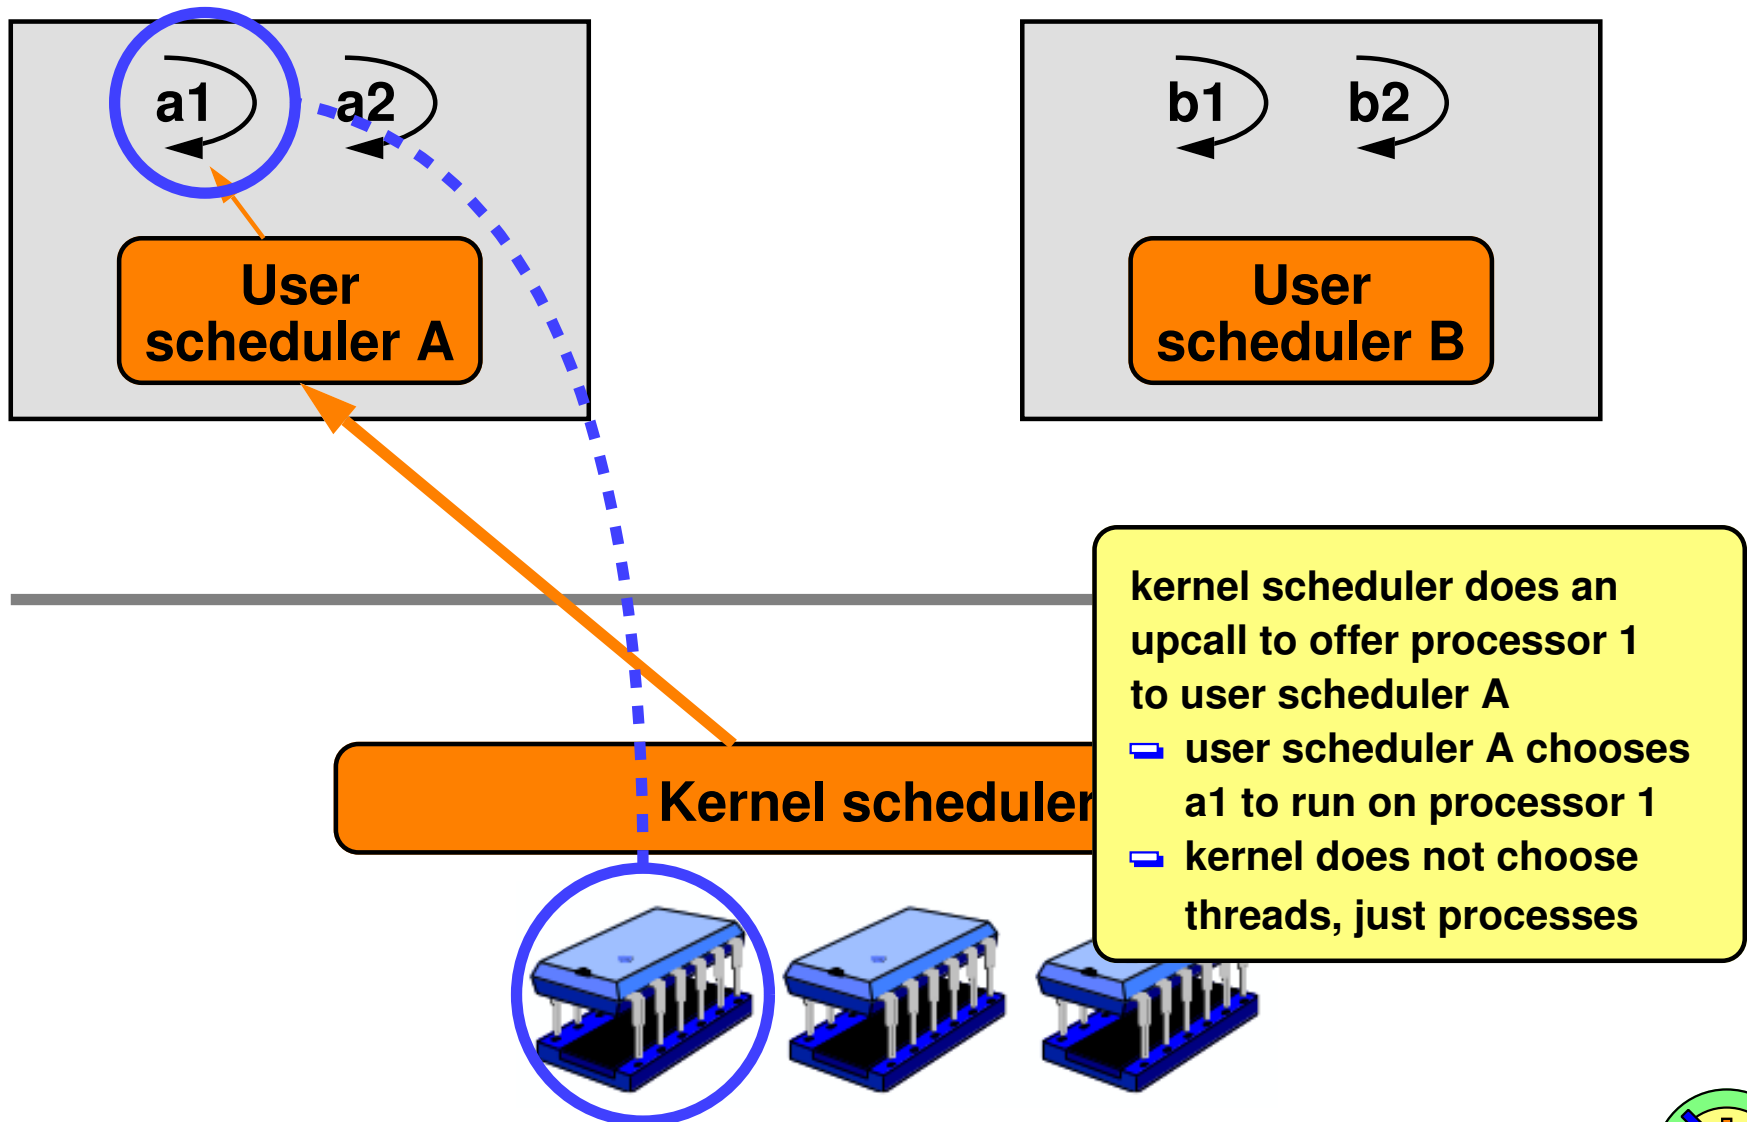
62

Scheduler Activations Model Example



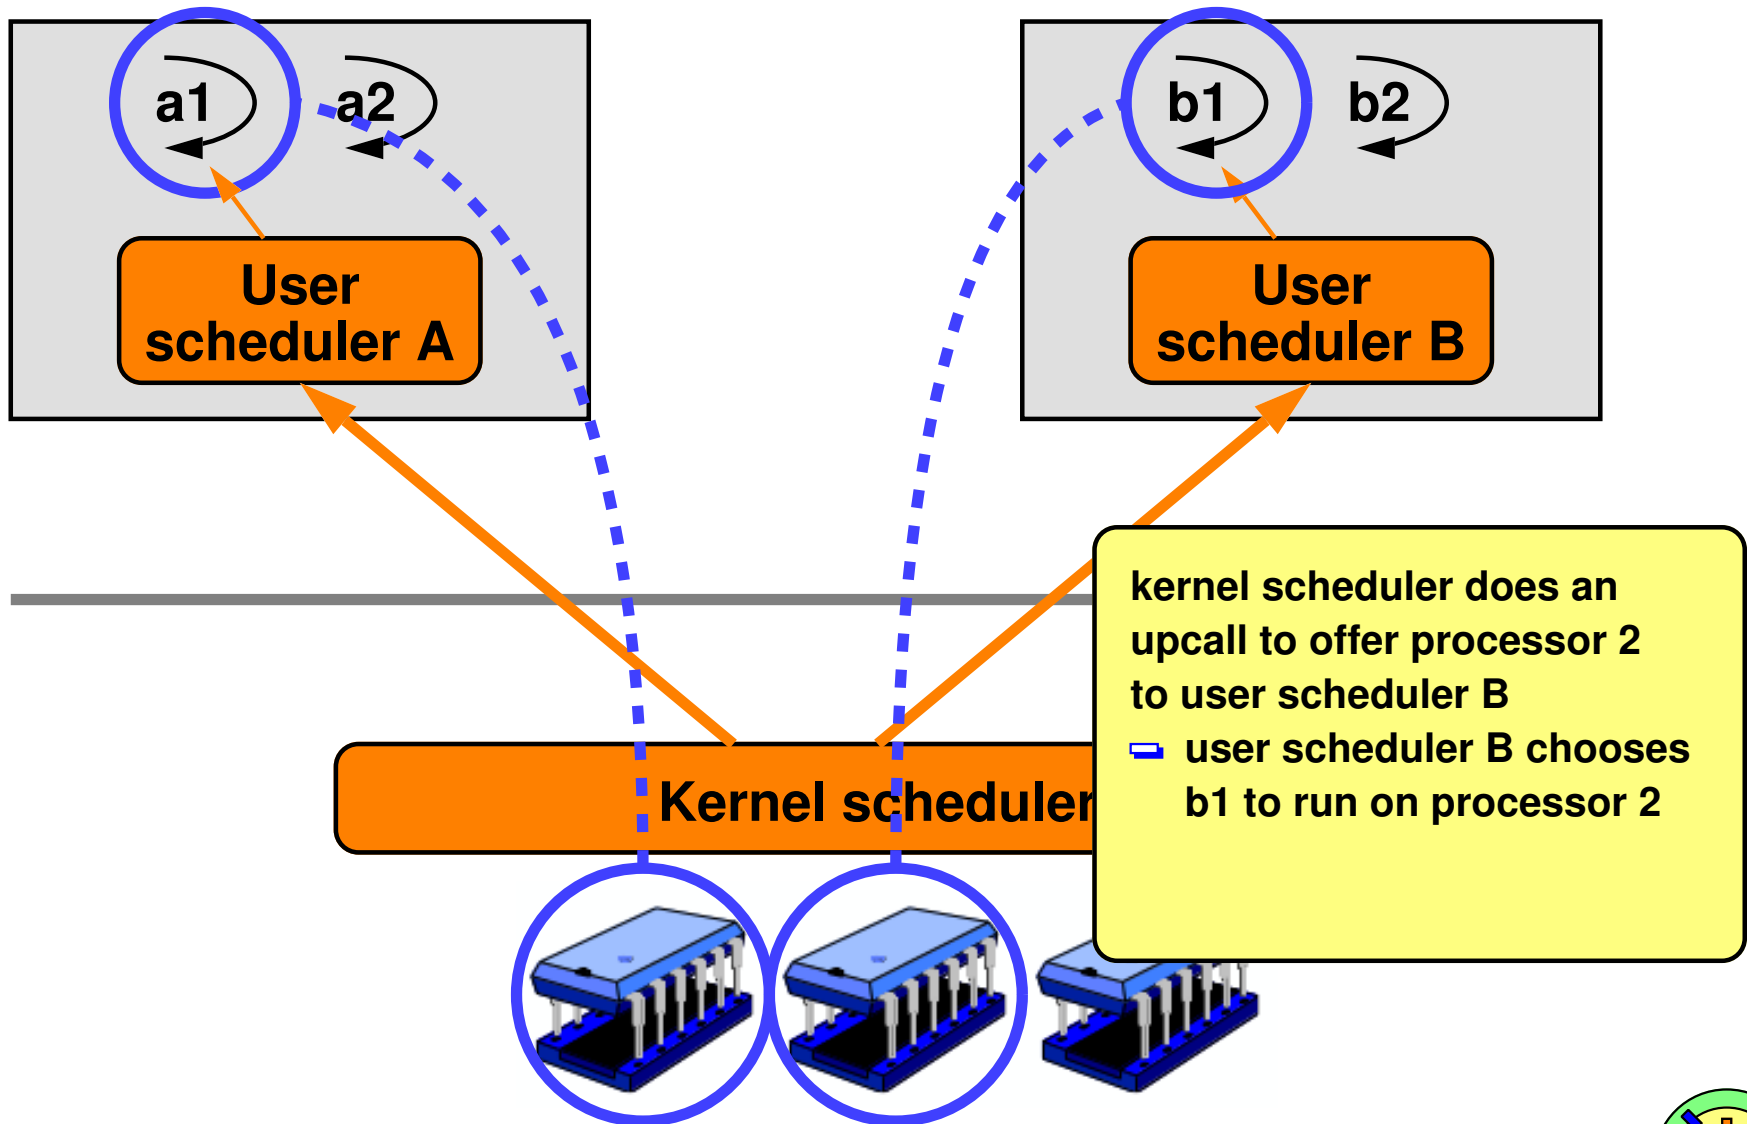
Kernel scheduler does not schedule threads

Scheduler Activations Model Example



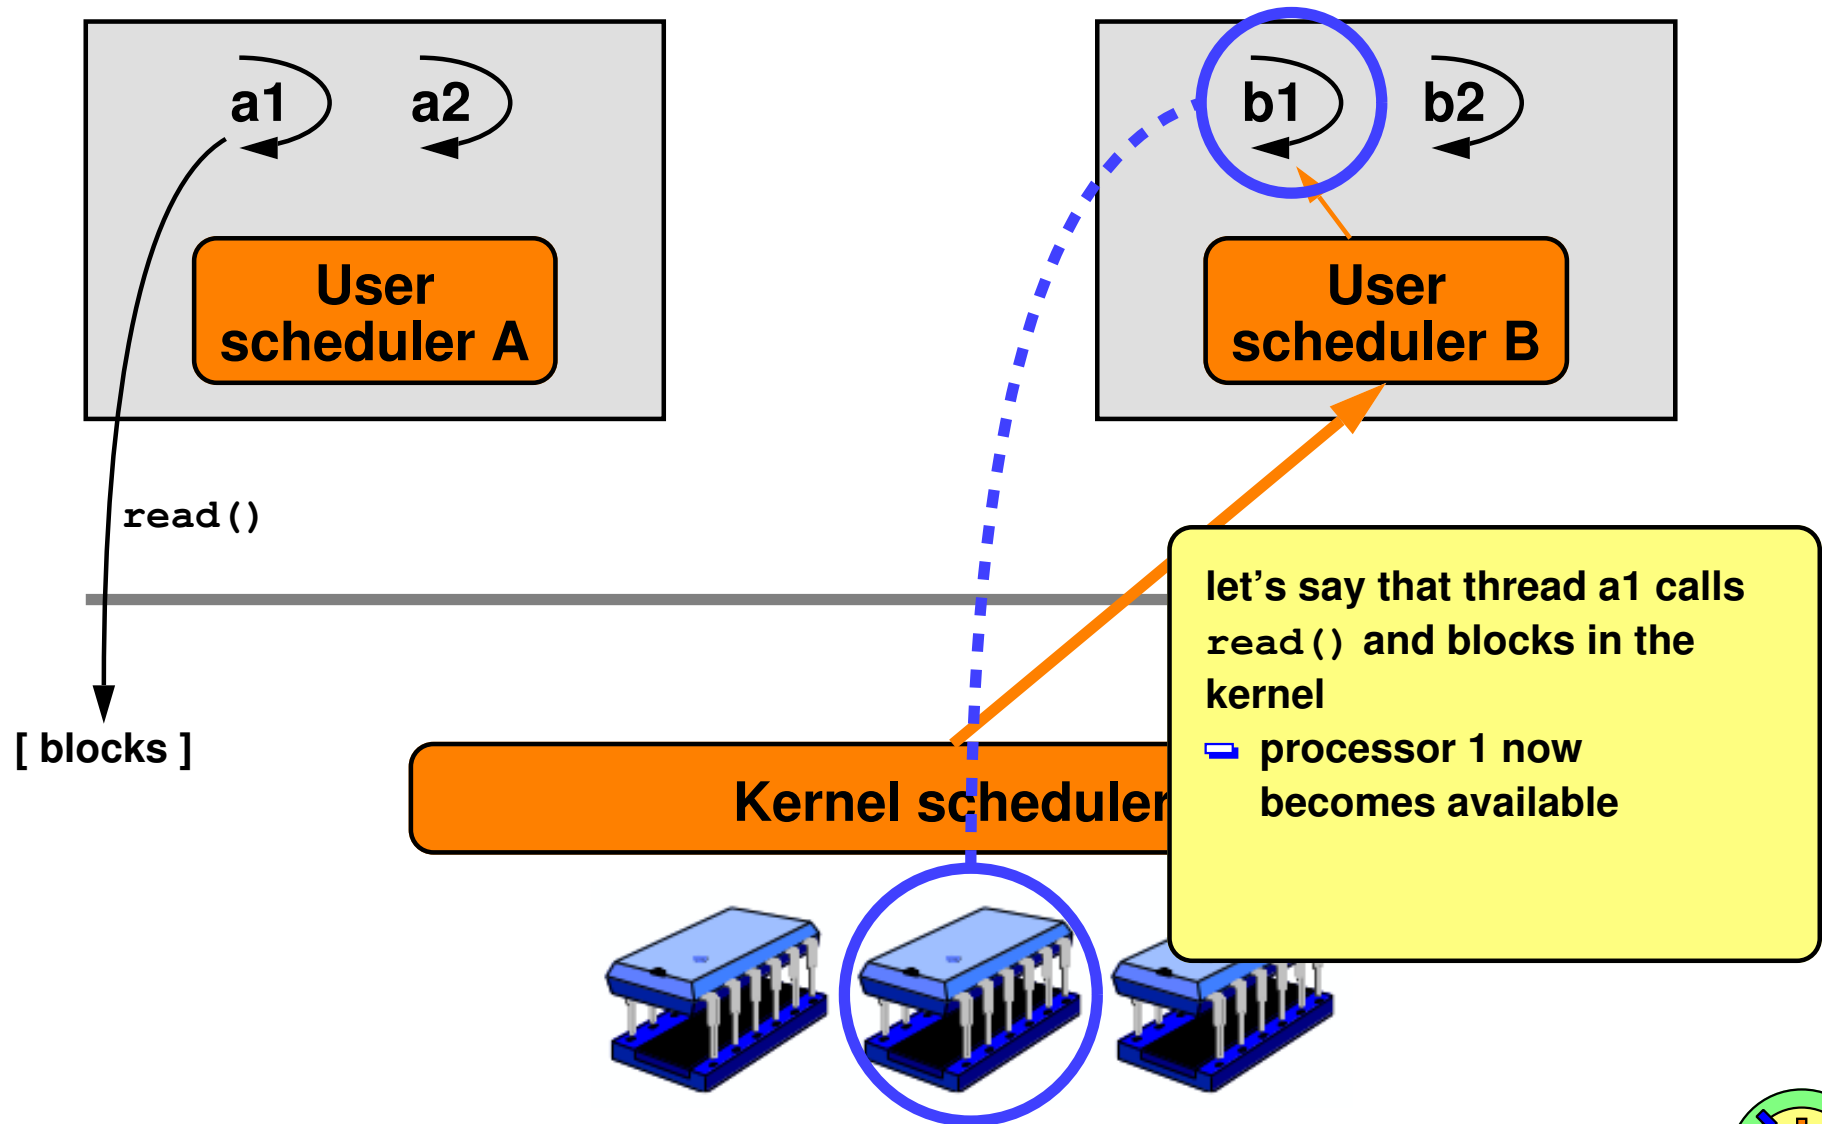
Kernel scheduler does not schedule threads

Scheduler Activations Model Example



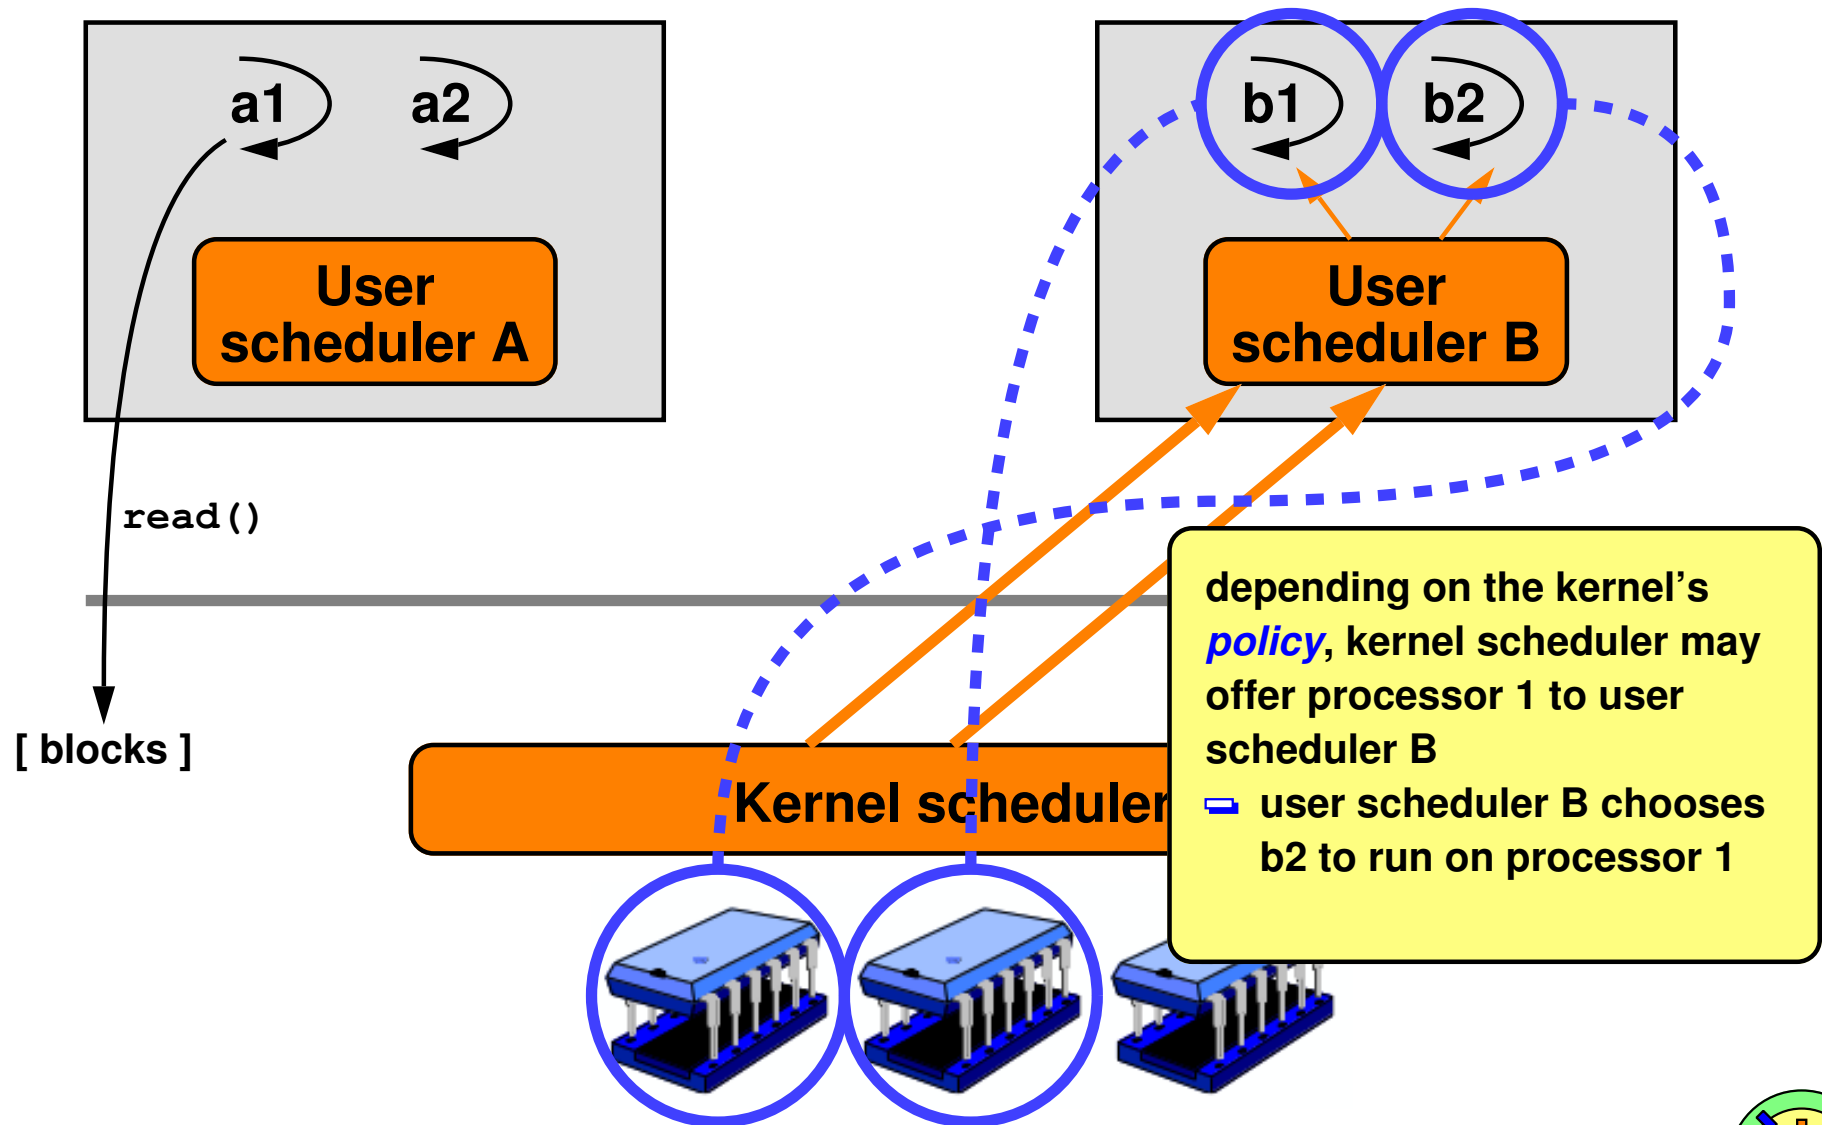
Kernel scheduler does not schedule threads

Scheduler Activations Model Example



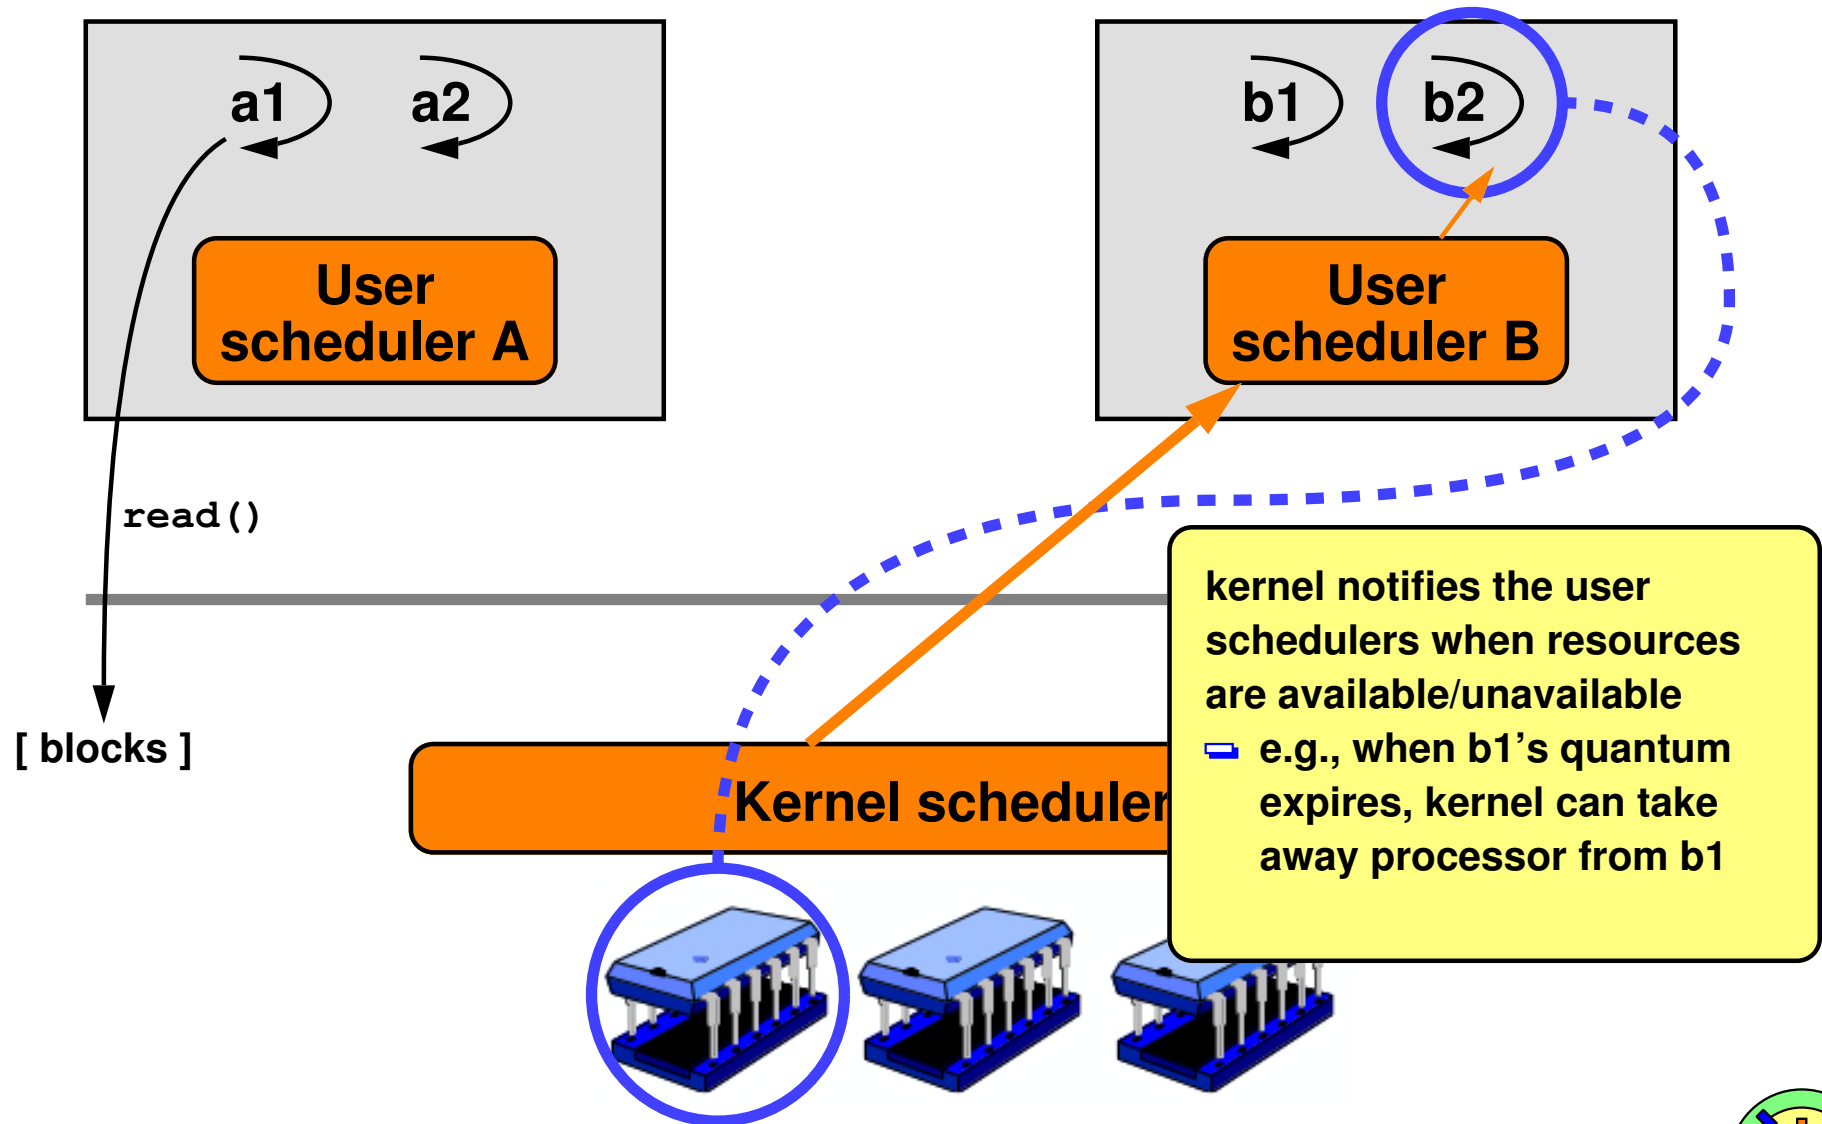
Kernel scheduler can have various scheduling policies

Scheduler Activations Model Example




Kernel scheduler can have various scheduling policies

Scheduler Activations Model Example



Kernel scheduler can have various scheduling policies



68