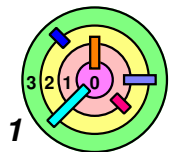



Housekeeping (Lecture 15 - 10/16/2013)

- ➡ Kernel #1 due at 11:45pm on Friday, 10/25/2013
 - if you have code from a previous semester, be very careful and *not copy any code from it*
 - it's best if you just get rid of it
- ➡ Post questions about the kernel #1 to *class Google Group*
 - your classmates may have better answers!
 - and faster turn-around
 - you should state what you have found so other students know that you have tried (and therefore, more willing to help)



Access Protection


- ➡ OS needs to make sure that only authorized processes are allowed access to system resources
 - various ways to provide this
- ➡ Unix (and many other systems, such as Windows) associates with files some indication of which *security principals* are allowed access
 - along with what sort of access is allowed
- ➡ A *security principal* is normally a user or group of users
 - a "user" can be an identity used by processes performing system functions
 - each running process can have several security principals associated with it
 - all processes have a user identification and a set of group identifications
 - for Sixth-Edition Unix, only one user ID and one group ID



Access Protection

- ➡ Each file has associated with it a set of access permissions
- for each of the 3 classes of principals, what sorts of operations on the file are allowed
 - the 3 **classes** are:
 - **user**: owner of the file
 - **group**: group owner of the file
 - **others**: everyone else
 - the operations are grouped into 3 classes:
 - **read**: can read a file or directory
 - **write**: can write a file or directory
 - **execute**: one must have **execute permission for a directory** in order to **follow a path through it**

- ➡ **Rules for checking permissions**
- 1) determines the **smallest class of principals the requester belongs to** (user being smallest and others being largest)
 - 2) then it checks for appropriate permissions with that class



Permissions Example

```
% ls -lR
.:
total 2
drwxr-x--x  2 bill  adm    1024 Dec 17 13:34 A
drwxr----- 2 bill  adm    1024 Dec 17 13:34 B


./A:
total 1
-rw-rw-rw-  1 bill  adm     593 Dec 17 13:34 x

./B:
total 2
-r--rw-rw-  1 bill  adm     446 Dec 17 13:34 x
-rw----rw-  1 trina adm     446 Dec 17 13:45 y
```



Suppose that `bill` and `trina` are members of the `adm` group and `andy` is not

1) Q: May `andy` list the contents of directory `A`?



Permissions Example

```
% ls -lR
.:
total 2
drwxr-x--x  2 bill  adm    1024 Dec 17 13:34 A
drwxr----- 2 bill  adm    1024 Dec 17 13:34 B

./A:
total 1
-rw-rw-rw-  1 bill  adm     593 Dec 17 13:34 x


./B:
total 2
-r--rw-rw-  1 bill  adm     446 Dec 17 13:34 x
-rw---rw-   1 trina  adm     446 Dec 17 13:45 y
```



Suppose that `bill` and `trina` are members of the `adm` group and `andy` is not

1) Q: May `andy` list the contents of directory `A`?

A: No



Permissions Example

```
% ls -lR
.:
total 2
drwxr-x--x  2 bill  adm    1024 Dec 17 13:34 A
drwxr----- 2 bill  adm    1024 Dec 17 13:34 B


./A:
total 1
-rw-rw-rw-  1 bill  adm     593 Dec 17 13:34 x

./B:
total 2
-r--rw-rw-  1 bill  adm     446 Dec 17 13:34 x
-rw----rw-  1 trina adm     446 Dec 17 13:45 y
```



Suppose that `bill` and `trina` are members of the `adm` group and `andy` is not

2) Q: May `andy` read `A/x`?



Permissions Example

```
% ls -lR
```

```
..:
```

```
total 2
```

```
drwxr-x--x  2 bill  adm    1024 Dec 17 13:34 A
```

```
drwxr----- 2 bill  adm    1024 Dec 17 13:34 B
```

```
./A:
```

```
total 1
```

```
-rw-rw-rw-  1 bill  adm      593 Dec 17 13:34 x
```

```
./B:
```

```
total 2
```

```
-r--rw-rw-  1 bill  adm      446 Dec 17 13:34 x
```


```
-rw----rw-  1 trina adm      446 Dec 17 13:45 y
```



Suppose that `bill` and `trina` are members of the `adm` group and `andy` is not

2) Q: May `andy` read `A/x`?

A: Yes



Permissions Example

```
% ls -lR
.:
total 2
drwxr-x--x  2 bill  adm    1024 Dec 17 13:34 A
drwxr----- 2 bill  adm    1024 Dec 17 13:34 B


./A:
total 1
-rw-rw-rw-  1 bill  adm     593 Dec 17 13:34 x

./B:
total 2
-r--rw-rw-  1 bill  adm     446 Dec 17 13:34 x
-rw----rw-  1 trina adm     446 Dec 17 13:45 y
```



Suppose that `bill` and `trina` are members of the `adm` group and `andy` is not

3) Q: May `trina` list the contents of directory `B`?



Permissions Example

```
% ls -lR
.:
total 2
drwxr-x--x  2 bill  adm    1024 Dec 17 13:34 A
drwxr----- 2 bill  adm    1024 Dec 17 13:34 B

./A:
total 1
-rw-rw-rw-  1 bill  adm     593 Dec 17 13:34 x


./B:
total 2
-r--rw-rw-  1 bill  adm     446 Dec 17 13:34 x
-rw----rw-  1 trina adm     446 Dec 17 13:45 y
```



Suppose that `bill` and `trina` are members of the `adm` group and `andy` is not

3) Q: May `trina` list the contents of directory `B`?

A: Yes



Permissions Example

```
% ls -lR
.:
total 2
drwxr-x--x  2 bill  adm    1024 Dec 17 13:34 A
drwxr----- 2 bill  adm    1024 Dec 17 13:34 B


./A:
total 1
-rw-rw-rw-  1 bill  adm     593 Dec 17 13:34 x

./B:
total 2
-r--rw-rw-  1 bill  adm     446 Dec 17 13:34 x
-rw----rw-  1 trina adm     446 Dec 17 13:45 y
```



Suppose that `bill` and `trina` are members of the `adm` group and `andy` is not

4) Q: May `trina` modify `B/y`?



Permissions Example

```
% ls -lR
.:
total 2
drwxr-x--x  2 bill  adm    1024 Dec 17 13:34 A
drwxr----- 2 bill  adm    1024 Dec 17 13:34 B

./A:
total 1
-rw-rw-rw-  1 bill  adm     593 Dec 17 13:34 x


./B:
total 2
-r--rw-rw-  1 bill  adm     446 Dec 17 13:34 x
-rw----rw-  1 trina adm     446 Dec 17 13:45 y
```



Suppose that `bill` and `trina` are members of the `adm` group and `andy` is not

4) Q: May `trina` modify `B/y`?

A: No



Permissions Example

```
% ls -lR
```

```
..:
```

```
total 2
```

```
drwxr-x--x  2 bill  adm    1024 Dec 17 13:34 A
```

```
drwxr----- 2 bill  adm    1024 Dec 17 13:34 B
```

```
./A:
```

```
total 1
```

```
-rw-rw-rw-  1 bill  adm     593 Dec 17 13:34 x
```

```
./B:
```

```
total 2
```


```
-r--rw-rw-  1 bill  adm     446 Dec 17 13:34 x
```

```
-rw----rw-  1 trina adm     446 Dec 17 13:45 y
```



Suppose that `bill` and `trina` are members of the `adm` group and `andy` is not

5) Q: May `bill` modify `B/x`?



Permissions Example

```
% ls -lR
.:
total 2
drwxr-x--x  2 bill  adm    1024 Dec 17 13:34 A
drwxr----- 2 bill  adm    1024 Dec 17 13:34 B

./A:
total 1
-rw-rw-rw-  1 bill  adm     593 Dec 17 13:34 x


./B:
total 2
-r--rw-rw-  1 bill  adm     446 Dec 17 13:34 x
-rw---rw-   1 trina  adm     446 Dec 17 13:45 y
```



Suppose that `bill` and `trina` are members of the `adm` group and `andy` is not

5) Q: May `bill` modify `B/x`?

A: No



Permissions Example

```
% ls -lR
```

```
..:
```

```
total 2
```

```
drwxr-x--x  2 bill  adm    1024 Dec 17 13:34 A
```

```
drwxr----- 2 bill  adm    1024 Dec 17 13:34 B
```

```
./A:
```

```
total 1
```

```
-rw-rw-rw-  1 bill  adm     593 Dec 17 13:34 x
```

```
./B:
```

```
total 2
```


```
-r--rw-rw-  1 bill  adm     446 Dec 17 13:34 x
```

```
-rw----rw-  1 trina adm     446 Dec 17 13:45 y
```



Suppose that `bill` and `trina` are members of the `adm` group and `andy` is not

6) Q: May `bill` read `B/y`?



Permissions Example

```
% ls -lR
```

```
..:
```

```
total 2
```

```
drwxr-x--x  2 bill  adm    1024 Dec 17 13:34 A
```

```
drwxr----- 2 bill  adm    1024 Dec 17 13:34 B
```

```
./A:
```

```
total 1
```

```
-rw-rw-rw-  1 bill  adm     593 Dec 17 13:34 x
```

```
./B:
```

```
total 2
```

```
-r--rw-rw-  1 bill  adm     446 Dec 17 13:34 x
```


```
-rw-- --rw-  1 trina adm     446 Dec 17 13:45 y
```



Suppose that `bill` and `trina` are members of the `adm` group and `andy` is not

6) Q: May `bill` read `B/y`?

A: No




Open

```
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>
int open(const char *path, int options [, mode_t mode])
```



options


- O_RDONLY open for reading only
- O_WRONLY open for writing only
- O_RDWR open for reading and writing
- O_APPEND set the file offset to end of file prior to each write
- O_CREAT if the file does not exist, then create it, setting its mode to mode adjusted by umask
- O_EXCL: if O_EXCL and O_CREAT are set, then open fails if the file exists
- O_TRUNC delete any previous contents of the file
- O_NONBLOCK don't wait if I/O cannot be done immediately



Setting File Permissions


```
#include <sys/types.h>
#include <sys/stat.h>
int chmod(const char *path, mode_t mode)
```

- sets the file permissions of the given file to those specified in mode
- only the owner of a file and the superuser may change its permissions
- nine combinable possibilities for mode (read/write/execute for user, group, and others)
- S_IRUSR (0400), S_IWUSR (0200), S_IXUSR (0100)
- S_IRGRP (040), S_IWGRP (020), S_IXGRP (010)
- S_IROTH (04), S_IWOTH (02), S_IXOTH (01)
- note: numeric prefix of 0 means the number is in octal format




Creating a File

- ➡ Use either `open` or `creat`
 - ▬ `open(const char *pathname, int flags, mode_t mode)`
 - flags must include `O_CREAT`
 - ▬ `creat(const char *pathname, mode_t mode)`
 - ▬ `open` is preferred
- ➡ The mode parameter helps specify the permissions of the newly created file
 - ▬ `permissions = mode & ~umask`



Umask

- ➡ Standard programs create files with "maximum needed permissions" as *mode*
 - compilers: 0777
 - editors: 0666
- ➡ Per-process parameter, *umask*, used to *turn off* undesired permission bits
 - e.g., turn off all permissions for others, write permission for group: set umask to 027
 - compilers: permissions = $0777 \ \& \ \sim(027) = 0750$
 - editors: permissions = $0666 \ \& \ \sim(027) = 0640$
 - set with `umask ()` system call or (usually) `umask` shell command






Ch 6: File Systems

Bill Cheng

<http://merlot.usc.edu/cs402-f13>




Files




Requirements

- ➡ **Permanent storage**
 - ▬ resides on disk (or alternatives)
 - ▬ survives software and hardware crashes
 - (including loss of disk?)
- ➡ **Quick, easy, and efficient**
 - ▬ satisfies needs of most applications
 - how do applications use permanent storage?



Applications

- ➡ **Software development**
 - text editors
 - linkers and loaders
 - source-code control
- ➡ **Document processing**
 - editing
 - browsing
- ➡ **Web stuff**
 - serving
 - browsing
- ➡ **Program execution**
 - paging



Needs




Directories

- convenient naming
- fast lookup



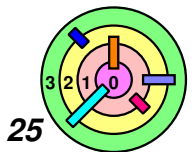
File access

- sequential is very common!
- "random access" is relatively rare




6.1 The Basics of File Systems

- ➡ *UNIX's S5FS*
- ➡ Disk Architecture
- ➡ Problems with S5FS
- ➡ Improving Performance
- ➡ Dynamic Inodes

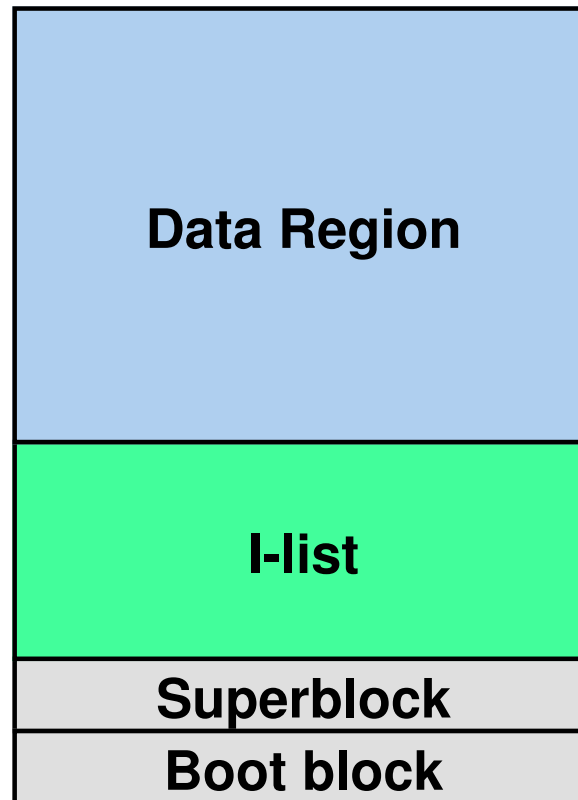


S5FS

- ➡ **A simple file system**
 - ▬ **slow**
 - ▬ **not terribly tolerant to crashes**
 - ▬ **reasonably efficient in space**
 - **no compression**
- ➡ **Concerns**
 - ▬ **on-disk data structures**
 - **file representation**
 - **free space**




S5FS Layout

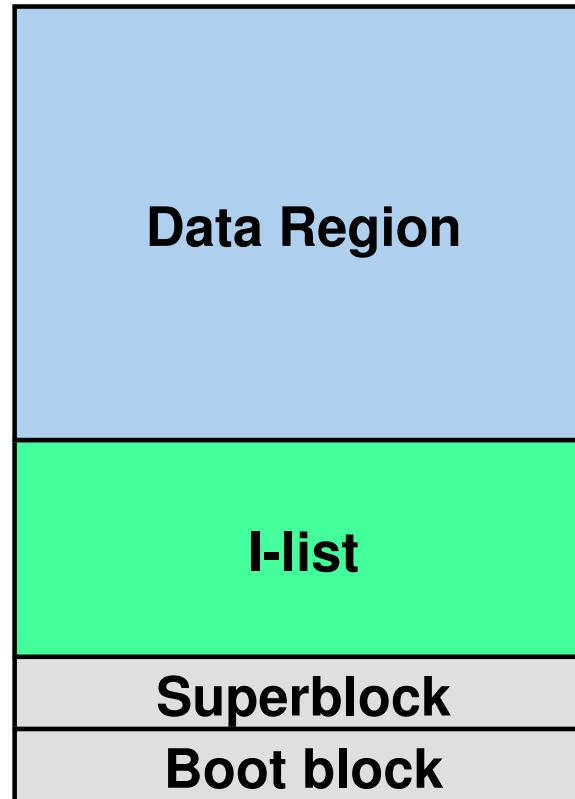


➡ A disk is simply an array of blocks of 1KB each (old Unix: 512B)

➡ A "linear view" (1-D array of blocks) of the disk



S5FS Layout




The *superblock*

- describes the layout of the rest of the file system
- contains the *heads* of the *free lists*


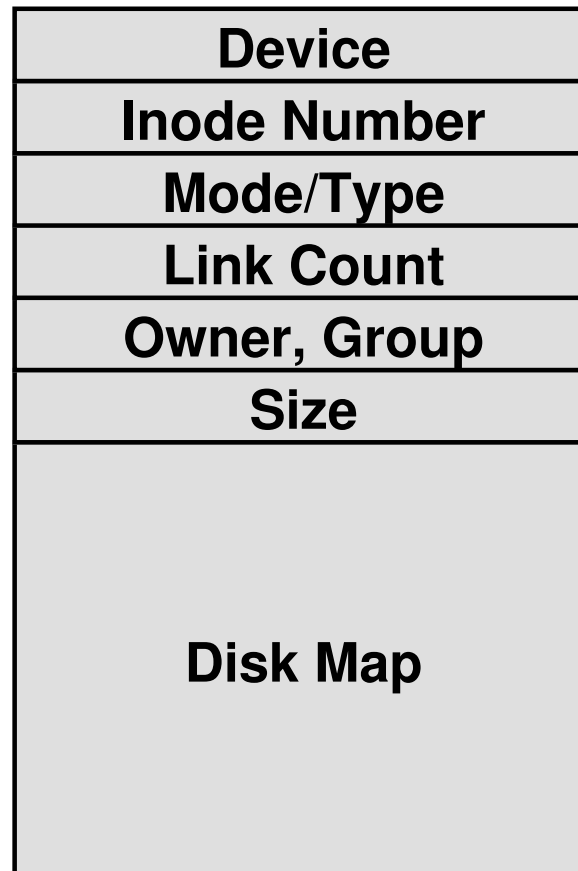


The *i-list* is an *array* of *index nodes (inodes)*

- each representing a *file*




S5FS: Inode




Disk Map

→ assuming blocksize = 1KB


0
1
2
3
4
5
6
7
8
9
10
11
12





Disk Map





- assuming blocksize = 1KB
- up to 10KB




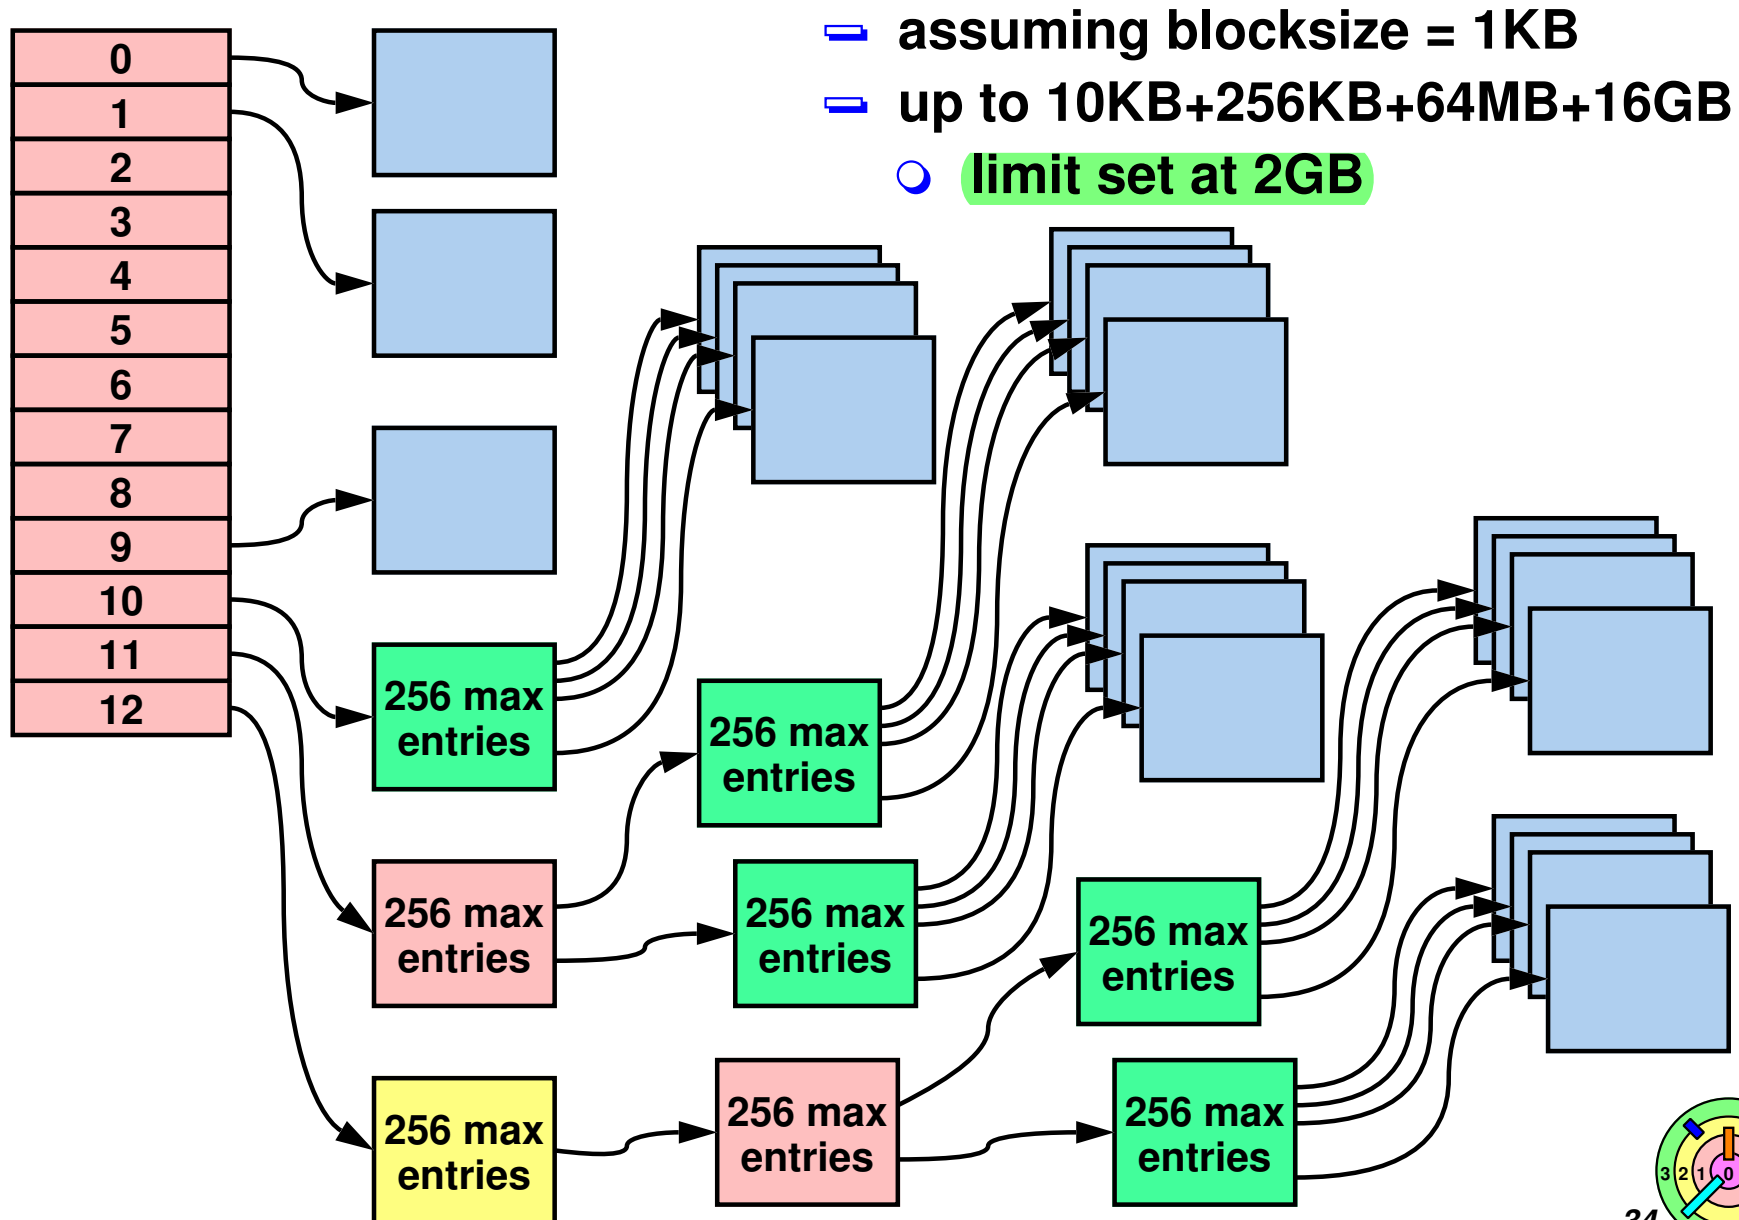
Disk Map



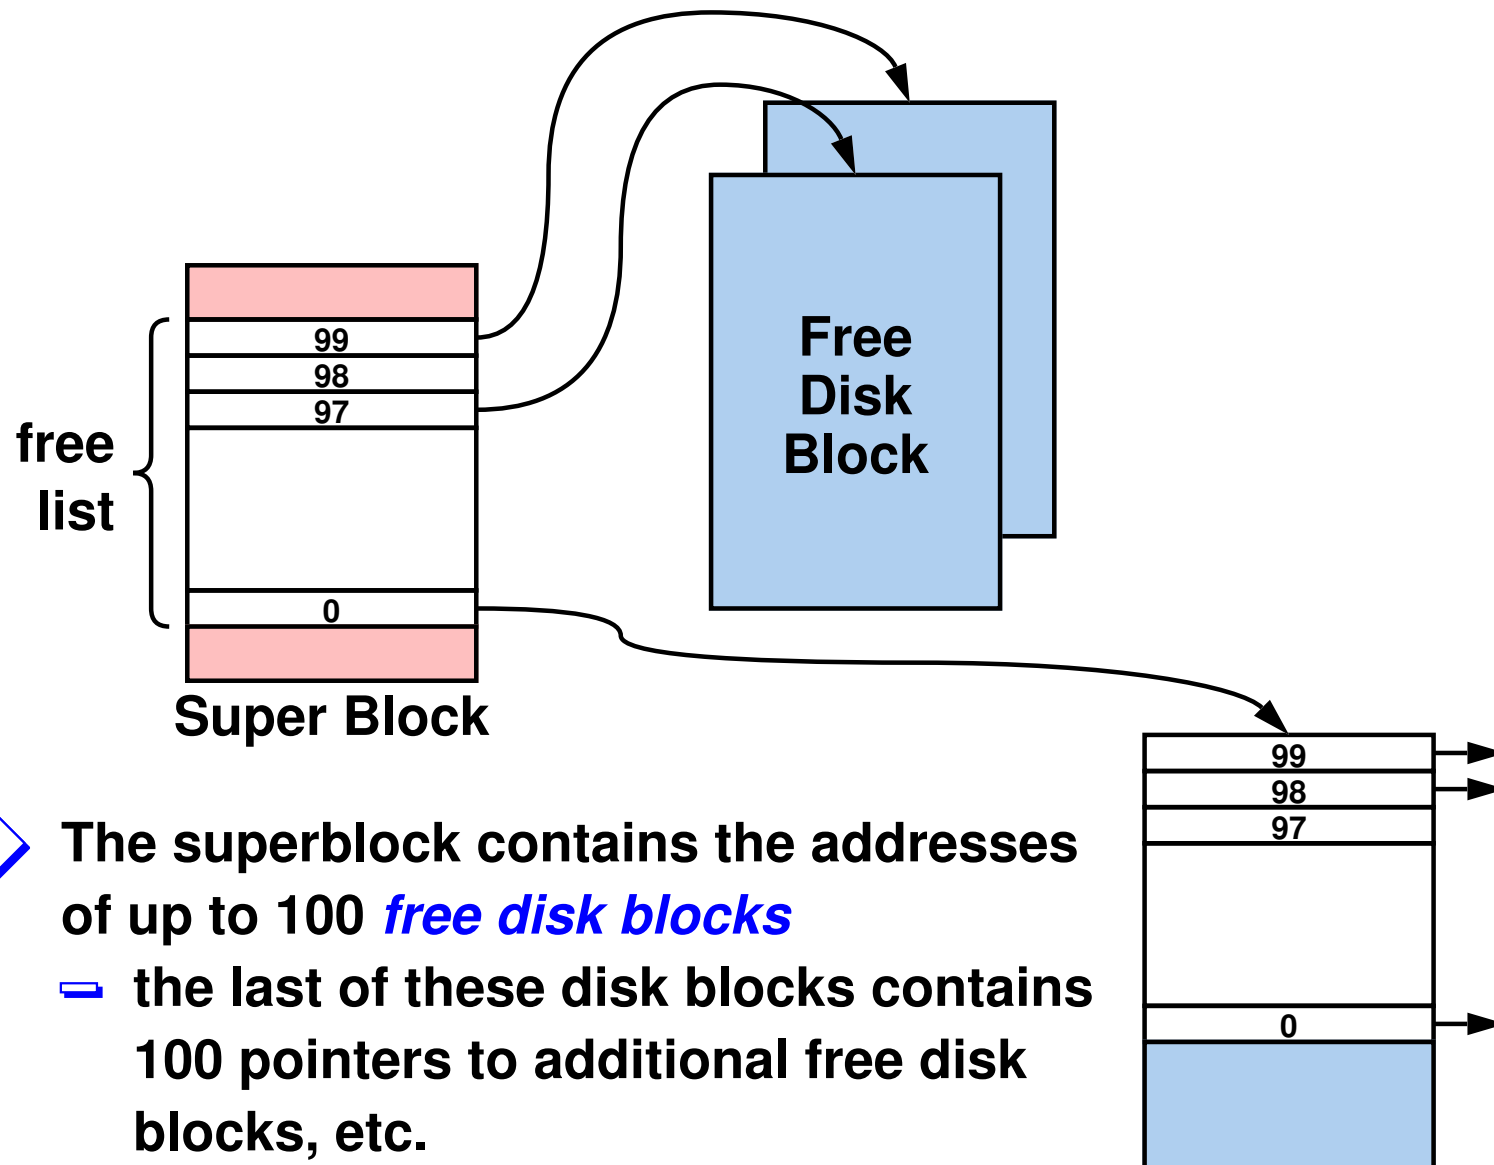
Disk Map



Disk Map

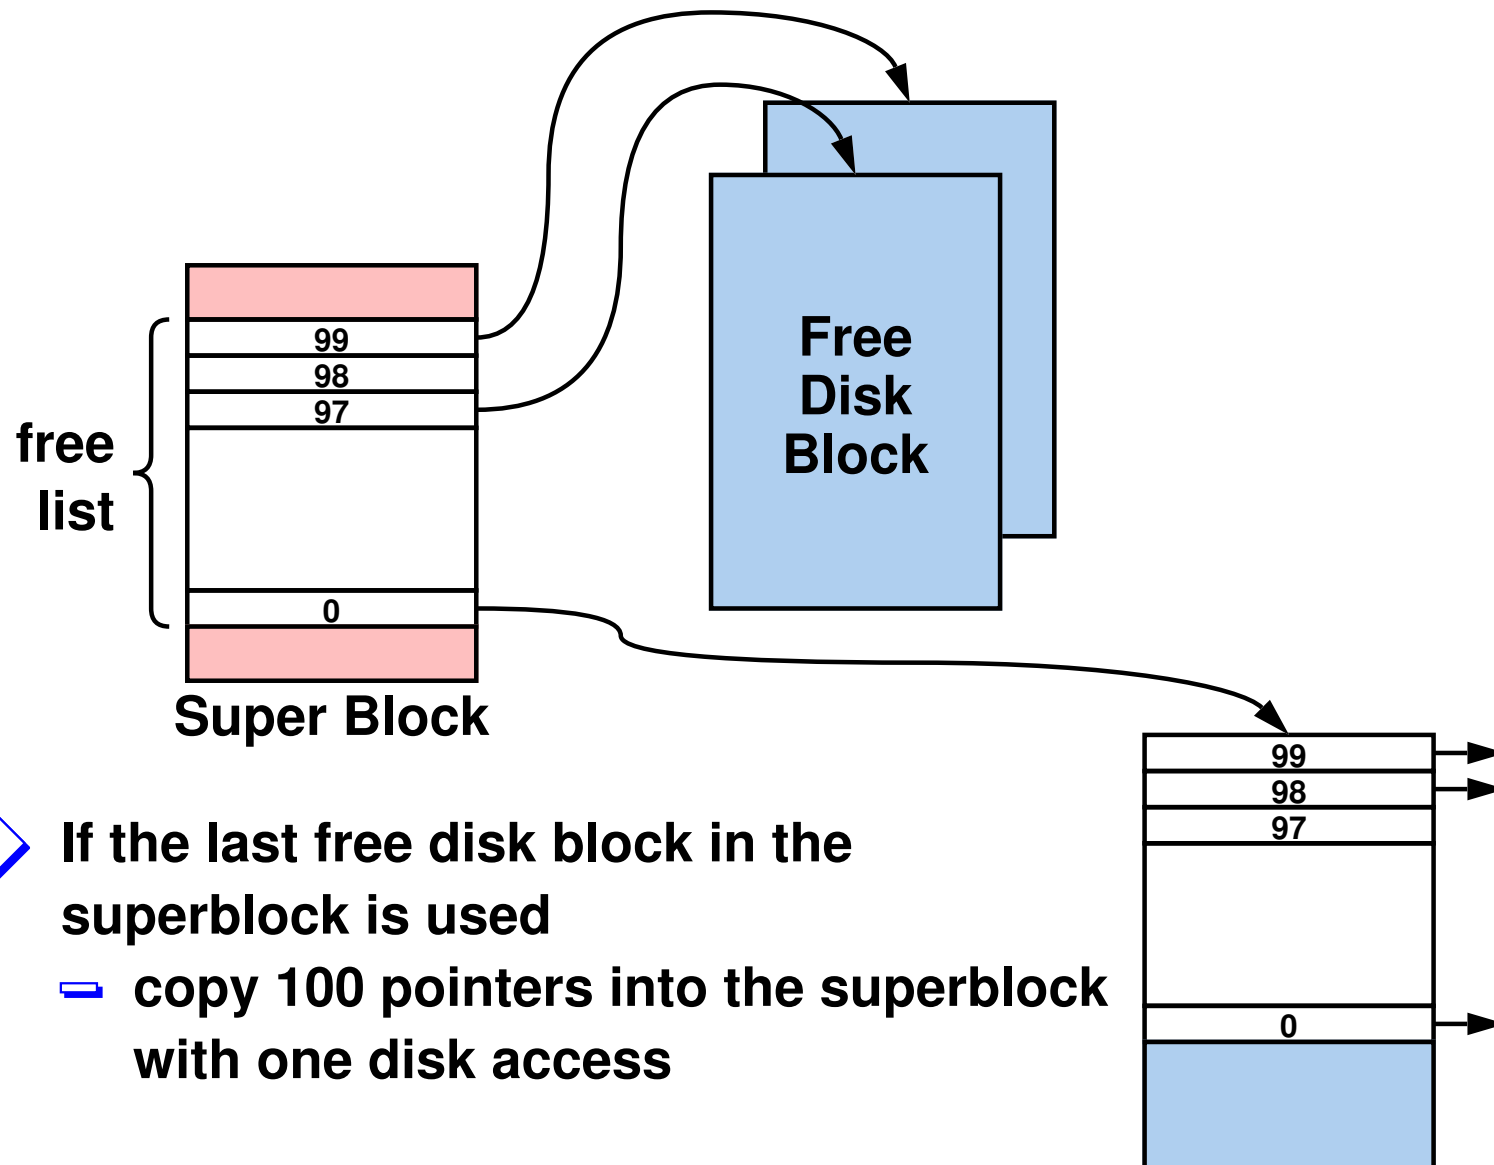


S5FS Free List



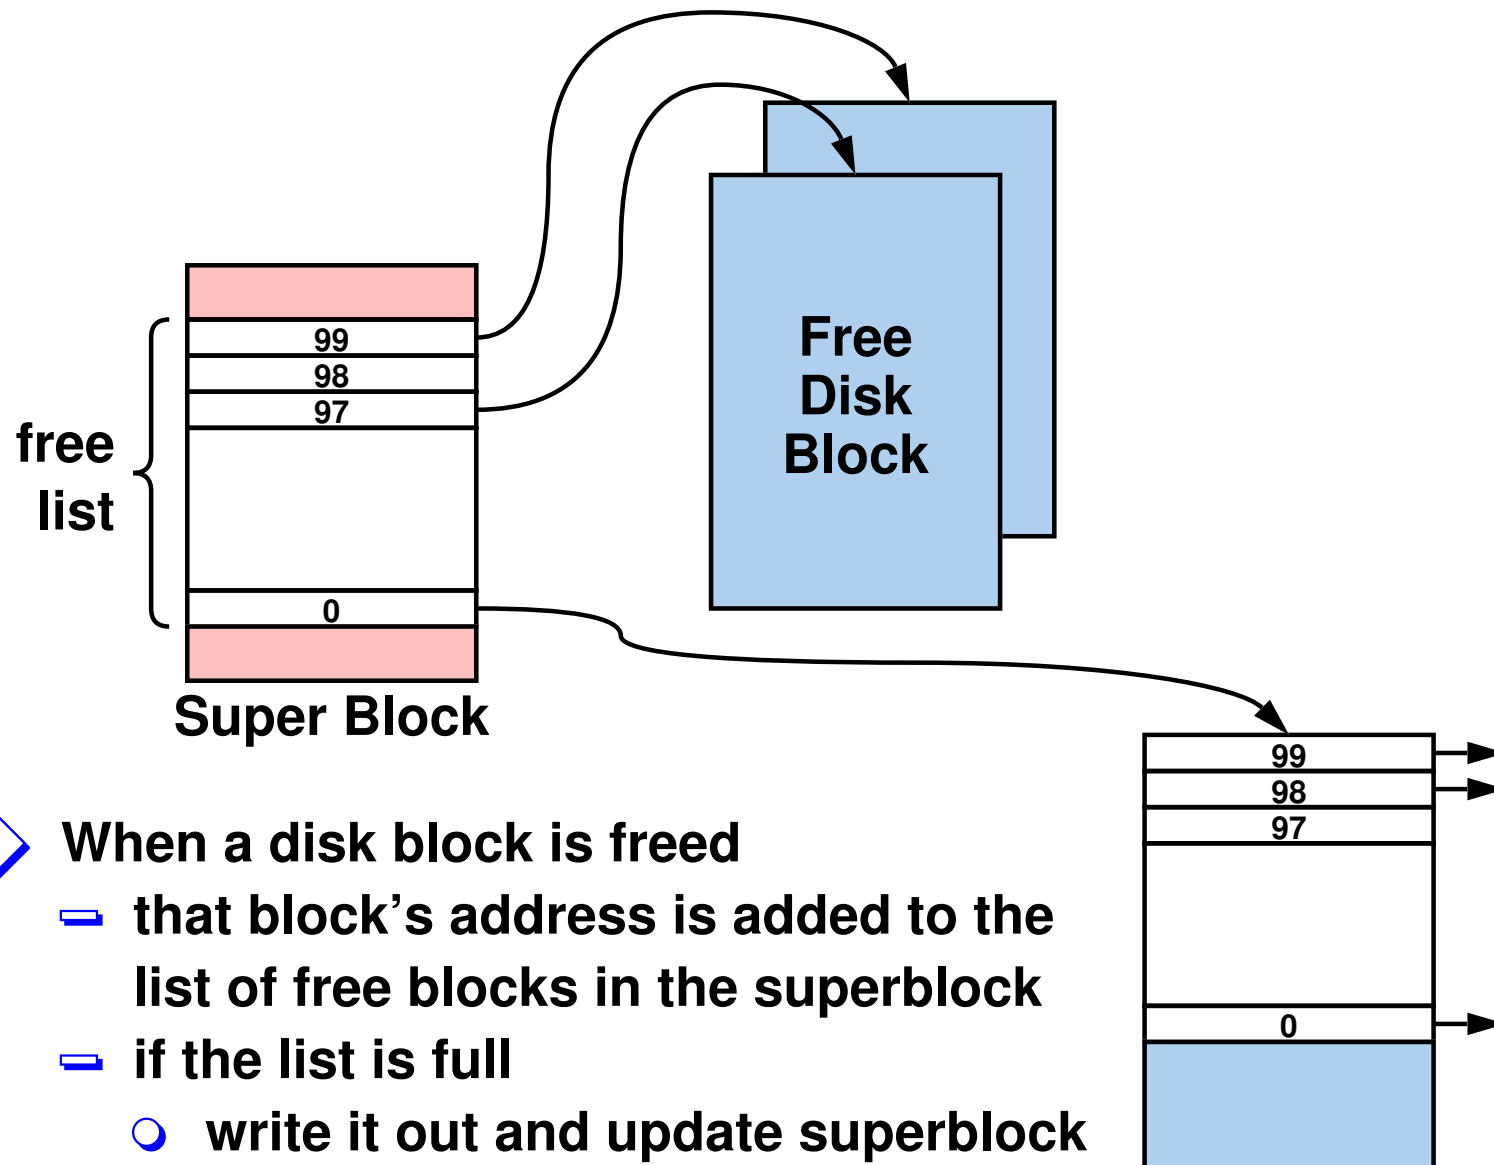
- ➡ The superblock contains the addresses of up to 100 *free disk blocks*
- the last of these disk blocks contains 100 pointers to additional free disk blocks, etc.
 - can find *all* the free disk blocks this way

S5FS Free List




- ➡ If the last free disk block in the superblock is used
- copy 100 pointers into the superblock with one disk access

S5FS Free List




When a disk block is freed

- that block's address is added to the list of free blocks in the superblock
- if the list is full
 - write it out and update superblock



S5FS Free Inode List




16	0	
15		
14		
13	0	
12	0	
11	0	
10		
9	0	
8		
7		
6	0	
5		
4	0	
3		
2		
1		

I-list




Inodes (in the i-list) are marked free or not free

- no additional organization in the i-list
- the superblock *cache*s free inodes (i.e., in the *inode cache*)



S5FS Free Inode List




16	0	
15		
14		
13	0	
12	0	
11	0	
10		
9	0	
8		
7		
6	0	
5		
4	0	
3		
2		
1		

I-list




The *inode cache*

- to allocate an inode, simply mark it not free and remove it from the inode cache
- to free an inode, simply mark it free and add to the inode cache if there is room




S5FS Free Inode List




16	0	
15		
14		
13	0	
12	0	
11	0	
10		
9	0	
8		
7		
6	0	
5		
4	0	
3		
2		
1		

I-list

- ➡ If the inode cache is empty
- ➡ scan the i-list to refill it
 - to help out with the scan, the inode cache contains the index of the first free inode in the i-list
 - ◆ need to maintain this entry when necessary




S5FS Free Inode List




16	0	
15		
14		
13	0	
12	0	
11	0	
10		
9	0	
8		
7		
6	0	
5		
4	0	
3		
2		
1		

I-list

- ➡ To *create a file*
- ➡ get a *free block*
 - update *free list*
 - ➡ get a *free inode*
 - update *i-list* and *inode cache*



S5FS Free Inode List




16	0	
15		
14		
13	0	
12	0	
11	0	
10		
9	0	
8		
7		
6	0	
5		
4	0	
3		
2		
1		

I-list



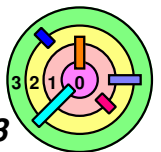
To *delete a file*

- ➡ add disk block(s) to *free list*
- ➡ mark inode free in *i-list* and may be update *inode cache*



6.1 The Basics of File Systems


- ➡ UNIX's S5FS
- ➡ *Disk Architecture*
- ➡ Problems with S5FS
- ➡ Improving Performance
- ➡ Dynamic Inodes



Disk Architecture


➡ How do you get to a particular byte of data on the disk?

— what does a disk look like?



➡ Clearly, this looks nothing like the S5FS layout

— or any logical file system layout



Rhinopias Disk Drive

Rotation speed	10,000 RPM
Number of surfaces	8
Sector size	512 bytes
Sectors/track	500-1000; 750 average
Tracks/surface	100,000
Storage capacity	307.2 billion bytes
Average seek time	4 milliseconds
One-track seek time	.2 milliseconds
Maximum seek time	10 milliseconds

