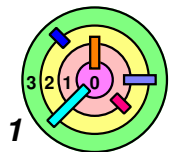



# Housekeeping (Lecture 16 - 10/21/2013)

- ➡ Kernel #1 due at 11:45pm this Friday, 10/25/2013
  - if you have code from a previous semester, be very careful and **not copy any code from it**
    - it's best if you just get rid of it
- ➡ If you modify additional files, make sure you include them in your submission
  - you would need to change your top-level Makefile in this case
  - you should be able to get the assignment to work without having to do this, but it's okay that you have to change other files
- ➡ **Grading guidelines** is the only way we will grade
  - make sure you have tried everything there
- ➡ After submission, make sure you **Verify Your Kernel Submission**
- ➡ **Midterm exam** coverage posted
  - see the News section of the class web page




# 6.1 The Basics of File Systems

- ➡ UNIX's S5FS
- ➡ Disk Architecture
- ➡ *Problems with S5FS*
- ➡ Improving Performance
- ➡ Dynamic Inodes




## S5FS on Rhinopias (A Marketing Disaster ...)

- ➡ Rhinopias's maximum transfer speed?
  - 63.9 MB/sec
- ➡ S5FS's average transfer speed on Rhinopias?
  - average seek time:
    - < 4 milliseconds (say 2)
  - average rotational latency:
    - ~3 milliseconds
  - per-sector transfer time:
    - negligible
  - time/sector: 5 milliseconds
  - transfer time: 102.4 KB/sec (.16% of maximum)



# 6.1 The Basics of File Systems

- ➡ UNIX's S5FS
- ➡ Disk Architecture
- ➡ Problems with S5FS
- ➡ *Improving Performance*
- ➡ Dynamic Inodes



# What to Do About It?




## Hardware

- employ pre-fetch buffer
  - filled by hardware with what's underneath head
  - helps reads a bit; doesn't help writes




## Software

- better on-disk data structures
  - increase block size
  - minimize seek time
  - reduce rotational latency





# FFS



- ➡ **Better on-disk organization**
- ➡ **Longer component names in directories**
- ➡ **Retains disk map of S5FS**



# Larger Block Size




# Larger Block Size







# The Down Side ...

**Smaller  
Block Size**





# The Down Side ...

**Smaller  
Block Size**




# The Down Side ...


**Smaller  
Block Size**





# The Down Side ...



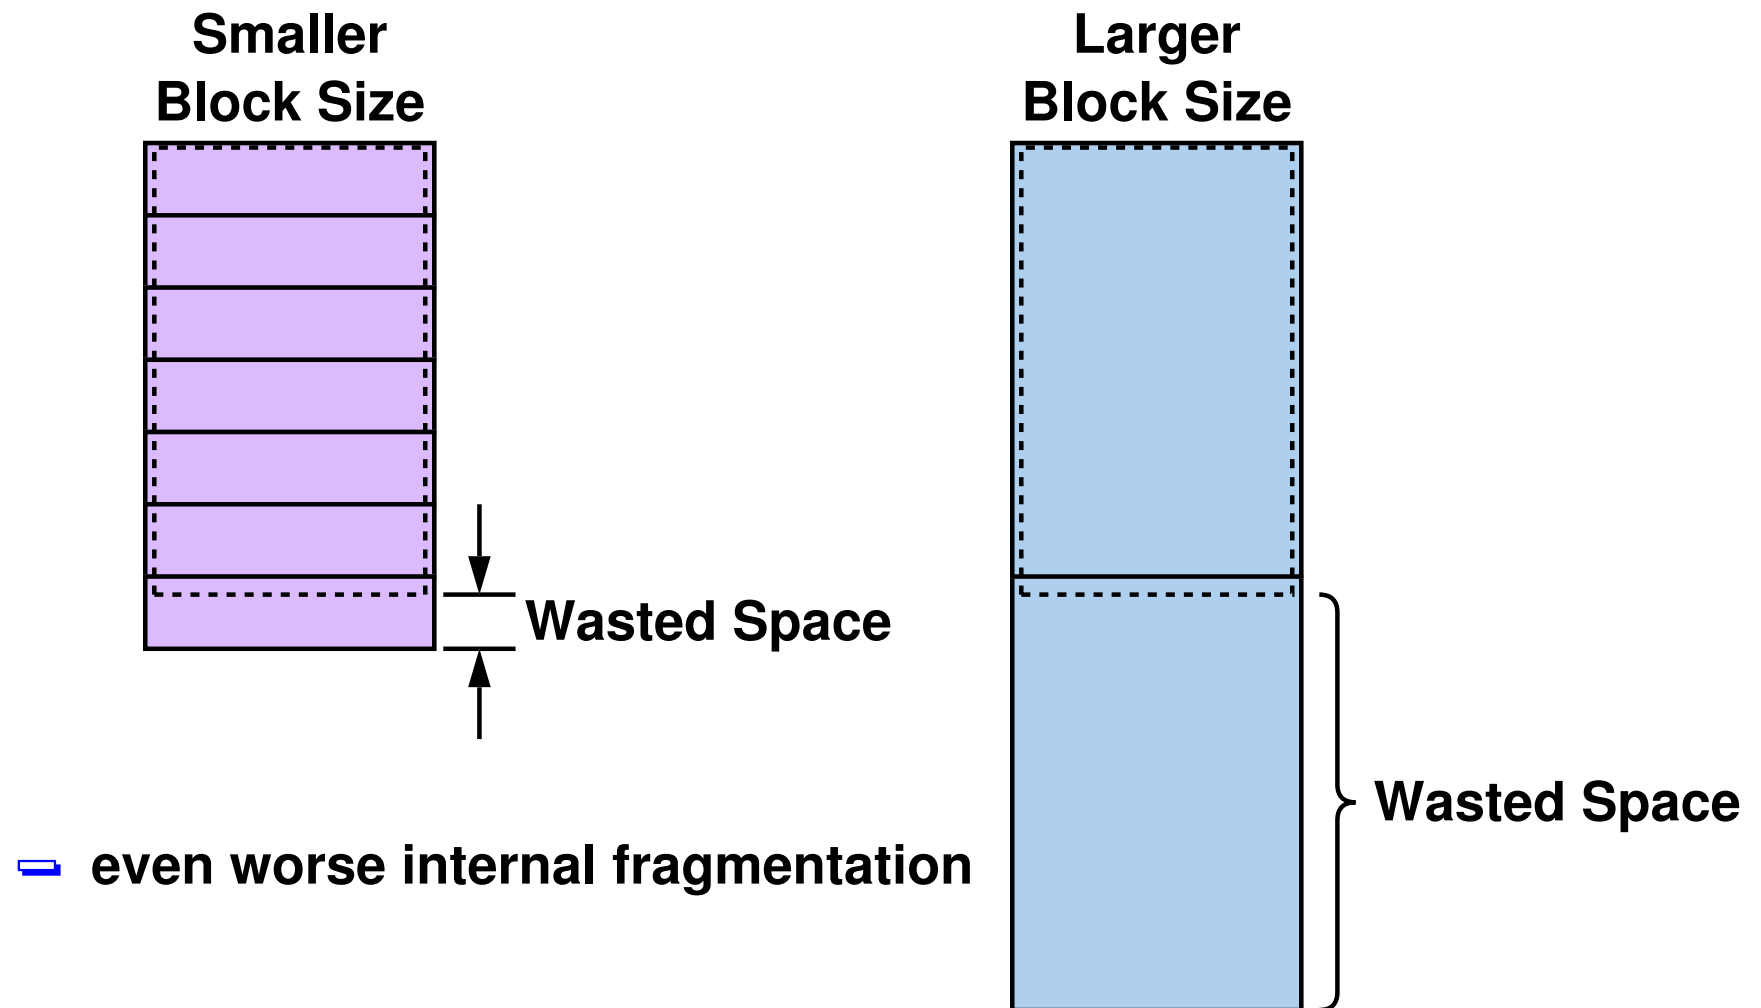
— internal fragmentation





# The Down Side ...




## The Down Side ...




# Two Block Sizes ...



## Two Block Sizes ...




- e.g., 16KB blocks and 1KB fragments
- best of both worlds









# Rules

- ➡ File-system blocks may be split into fragments that can be independently assigned to files
  - fragments assigned to a file must be contiguous and in order
- ➡ The number of fragments per block (1, 2, 4, or 8) is fixed for each file system
- ➡ Allocation in fragments may only be done on what would be the last block of a file, and only for small files




# Use of Fragments (1)







 File A

 File B




## Use of Fragments (2)






 File A

 File B

— A can grow by 2 segments





## Use of Fragments (3)




 File A

 File B




# Minimizing Seek Time


- ➡ Keep related things close to one another
- ➡ Separate unrelated things



# Cylinder Groups



⇒ recall that *seeking* to the *next sector* is much *faster*




# Minimizing Seek Time




**The practice:**


- attempt to put new inodes in the same cylinder group as their directories**
- put inodes for new directories in cylinder groups with "lots" of free space**
- put the beginning of a file (first 10KB, i.e., direct blocks) in the inode's cylinder group**
- put additional portions of the file (each 2MB) in cylinder groups with "lots" of free space**



# Locality Of File Access




— if access "d.c", likely to access "e.c"






# Locality Of File Access




— if access "d.c", likely to access "e.c"





## How Are We Doing?

- ➡ Configure Rhinopias with 20 cylinders per group
  - 2-MB file fits entirely within one cylinder group
  - average seek time within cylinder group is  $\sim .3$  milliseconds
  - average rotational delay still 3 milliseconds
  - .12 milliseconds required for disk head to pass over 8KB block
  - 3.42 milliseconds for each block
  - 2.4 million bytes/second average transfer time
  - 20-fold improvement
  - 3.7% of maximum possible




# Minimizing Latency (1)




# Numbers


- ➡ **Rhinopias spins at 10,000 RPM**
  - **6 milliseconds/revolution**
- ➡ **100 microseconds required to service disk-completion interrupt and start next operation**
  - **typical of early 1980s**
- ➡ **Each block takes 120 microseconds to traverse disk head**
- ➡ **Reading successive blocks is expensive!**



## Minimizing Latency (2)




**Block interleaving**



## How're We Doing Now? (part 1)

- ➡ Time to read successive blocks (two-way interleaving):
- after request for second block is issued, must wait 20 microseconds for the beginning of the block to rotate under disk head
  - factor of 300 improvement!

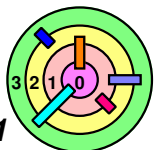


## How're We Doing Now? (part 2)




**Same setup as before**

- 2-MB file within one cylinder group**
- actually fits in one cylinder**
- block interleaving employed: every other block is skipped**
- .3-millisecond seek to that cylinder**
- 3-millisecond rotational delay for first block**
- 50 blocks/track, but 25 read in each revolution**
- 10.24 revolutions required to read all of file**
- 32.4 MB/second (50% of maximum possible)**



## Further Improvements?


- ➡ S5FS: 0.16% of capacity
- ➡ FFS without block interleaving: 3.8% of capacity
- ➡ FFS with block interleaving: 50% of capacity
- ➡ What next?






# Larger Transfer Units

- ➡ **Allocate in whole tracks or cylinders**
  - too much wasted space
- ➡ **Allocate in blocks, but group them together**
  - transfer many at once



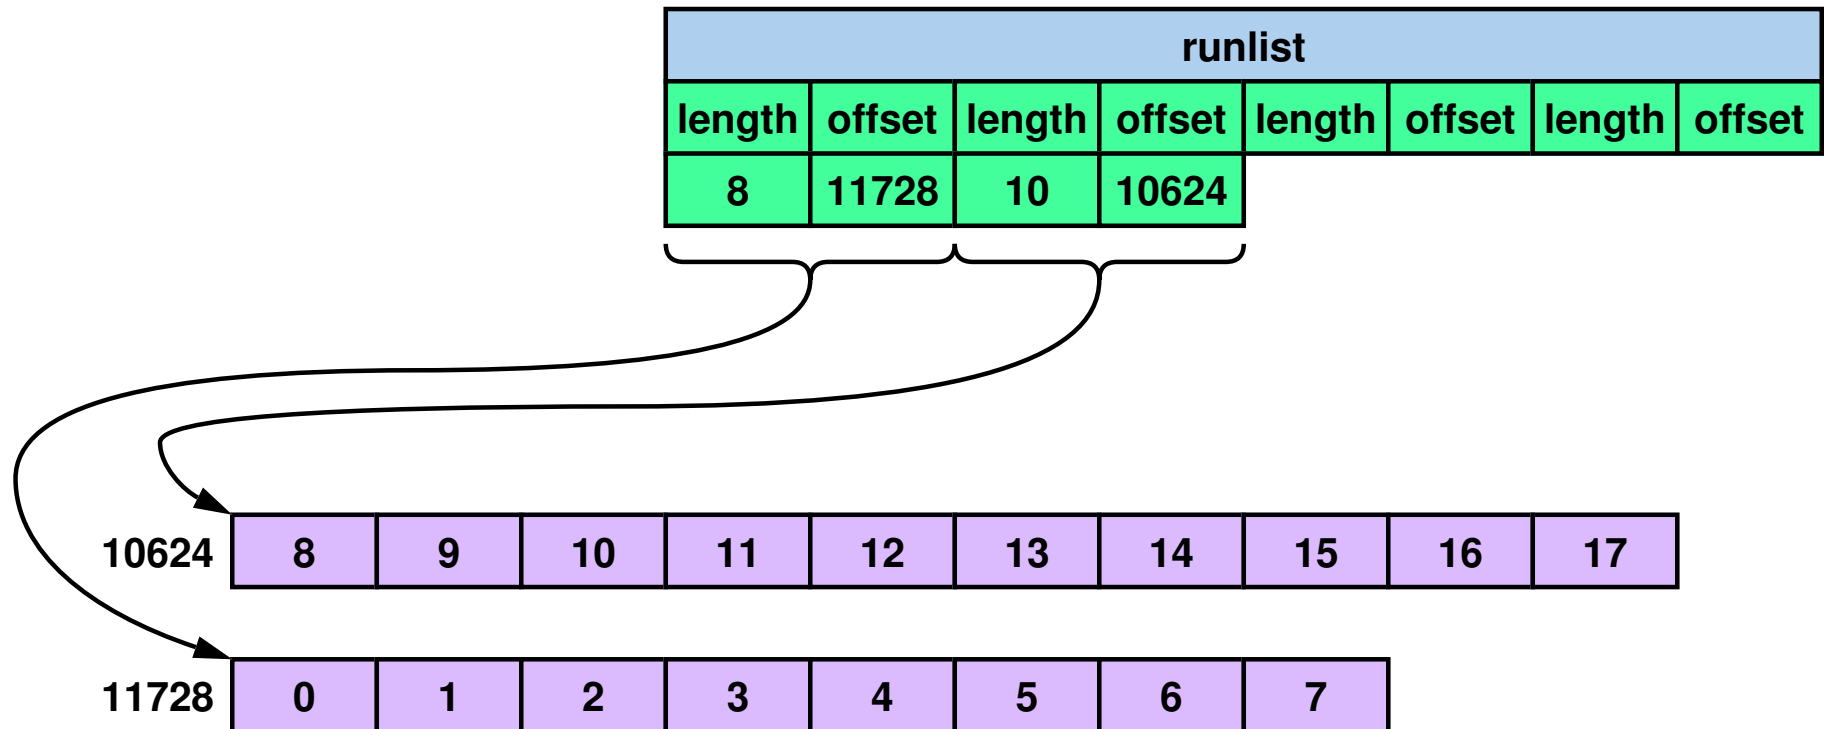
# Block Clustering

- ➡ Allocate space in blocks, eight at a time
- ➡ Linux's Ext2 (an FFS clone):
  - allocate eight blocks at a time
  - extra space is available to other files if there is a shortage of space
- ➡ FFS on Solaris (~1990)
  - delay disk-space allocation until:
    - 8 blocks are ready to be written
    - or the file is closed




# Extents

## ➡ Windows



# Problems with Extents

- ➡ Could result in highly fragmented disk space
  - lots of small areas of free space
    - external fragmentation
  - solution: use a *defragmenter* to *coalesce* free space
- ➡ Random access
  - linear search through a long list of extents
  - solution: multiple levels



# Extents in NTFS




Top-level run list							
length	offset	length	offset	length	offset	length	offset
50000	1076	10000	9738	36000	5192	2200	14024

Runlist							
length	offset	length	offset	length	offset	length	offset
8	11728	10	10624				

10624	50008	50009	50010	50011	50012	50013	50014	50015	50016	50017
11728	50000	50001	50002	50003	50004	50005	50006	50007		




# Are We There Yet?




# S5FS & FFS Data Placement Example

File Data:  
(e.g., PDF)





File On-disk  
Representation:



S5FS:





FFS:





# CPU, Memory, Disk Speeds Over Time

Capacity/Speed






# CPU, Memory, Disk Speeds Over Time

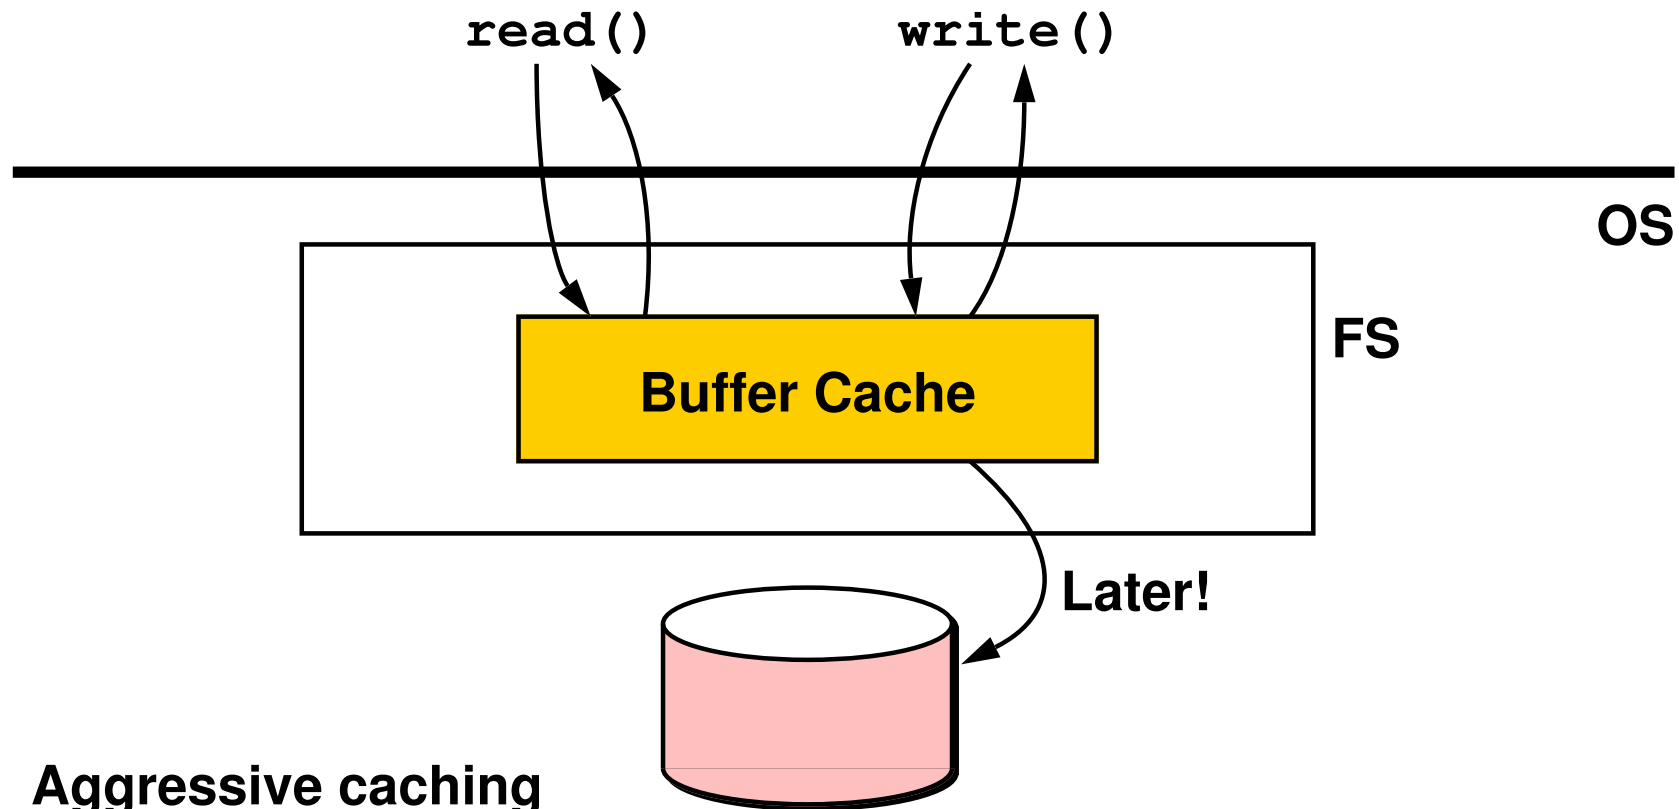


# A Different Approach

- ➡ We have lots of primary memory
  - enough to cache all commonly used files
- ➡ Read time from disk doesn't matter
- ➡ Time for writes does matter




# The Buffer Cache



## Aggressive caching

- most read and write will have a **cache hit**
- writes to the disk can wait, may be for quite a while
  - longer the wait, higher the **risk**
- file system optimized for **writing!**
  - how? you organize the disk as a very long **log**






# Log-Structured File Systems



Main principles

= *never delete*

= *append only*




# Log-Structured File Systems




Main principles

= *never delete*

= *append only*




0 1 2 ...




# Log-Structured File Systems

- ➡ How does "never delete" and "always append" help with performance?
  - ▬ minimize seek latency
  - ▬ minimize rotational latency
    - write a cylinder at a time
- ➡ ***Sprite FS*** (a log-structured file system)
  - ▬ through batching, a single, long write can write out everything





# LFS Data Placement Example

File On-disk  
Representation:




LFS:




# LFS Data Placement Example

➡ What happens if you want to modify the file?  
 — how does "append-only" really work?

➡ Ex: you create file A and then file B



- you modify file A, e.g., append to the last block of file A
- the new file will be referred as A'







# LFS Data Placement Example

➡ **What happens if you want to modify the file?**  
 ➡ **how does "append-only" really work?**

 **Ex: you create file A and then file B**



- ➡ you modify file A, e.g., append to the last block of file A
- ➡ the new file will be referred as A'




# LFS Data Placement Example

➡ What happens if you want to modify the file?  
 — how does "append-only" really work?

➡ Ex: you create file A and then file B

LFS:




*Inode Map:*

A'

B

- you modify file A, e.g., append to the last block of file A
- the new file will be referred as A'
  - the inode has changed as well




# LFS Data Placement Example

➡ What happens if you want to modify the file?  
 — how does "append-only" really work?


➡ Ex: you create file A and then file B

LFS:




*Inode Map:*

- you modify file A, e.g., append to the last block of file A
- the new file will be referred as A'
  - the inode has changed as well
- a *piece* of the *inode map* is appended to the log
  - fixed regions (previous version and current version) on the disk keeps track of *all* the *inode map pieces*
    - ◆ known as *checkpoint file*




## More On Inode Map

- ➡ ***Inode Map*** cached in primary memory
  - ▬ indexed by inode number
  - ▬ points to inode on disk
  - ▬ written out to disk in pieces as updated
  - ▬ ***checkpoint file*** contains locations of pieces
    - written to disk occasionally
    - two copies: current and previous
- ➡ **Commonly/Recently used inodes and other disk blocks cached in primary memory**



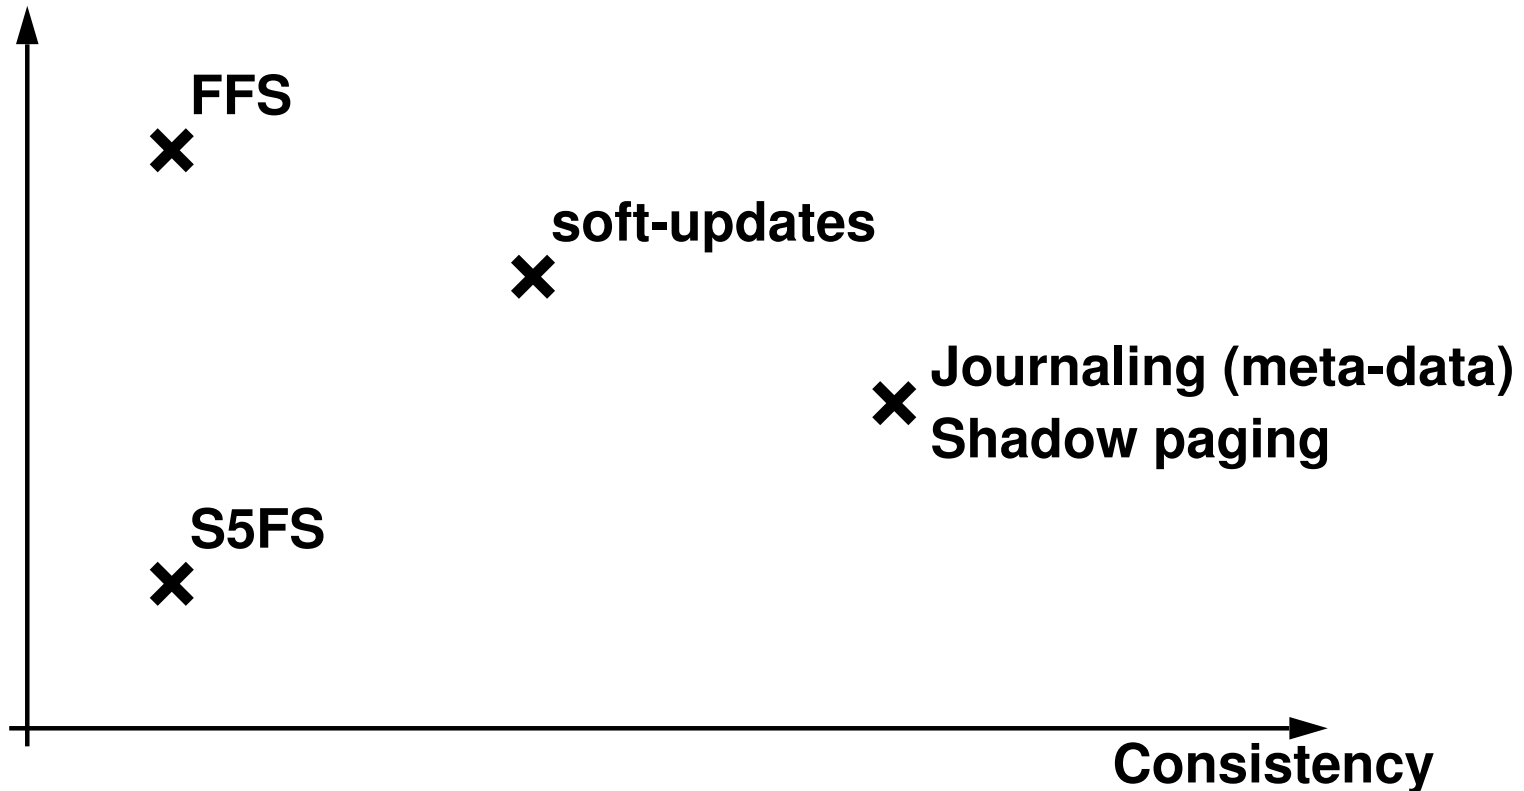
# 6.2 Crash Resiliency

- ➡ What Goes Wrong
- ➡ Dealing with Crashes




# Overveiw

Performance




- soft-update provides *recoverable consistency*
- journaling and shadow paging provide *transactional consistency*





## In the Event of a Crash ...

- ➡ **Most recent updates did not make it to disk**
  - **is this a big problem?**
  - **equivalent to crash happening slightly earlier**
    - **but you may have received (and believed) a message:**
      - ◆ **"file successfully updated"**
      - ◆ **"homework successfully handed in"**
      - ◆ **"stock successfully purchased"**
  - **there's worse ...**





# File-System Consistency (1)










# File-System Consistency (1)





# File-System Consistency (1)




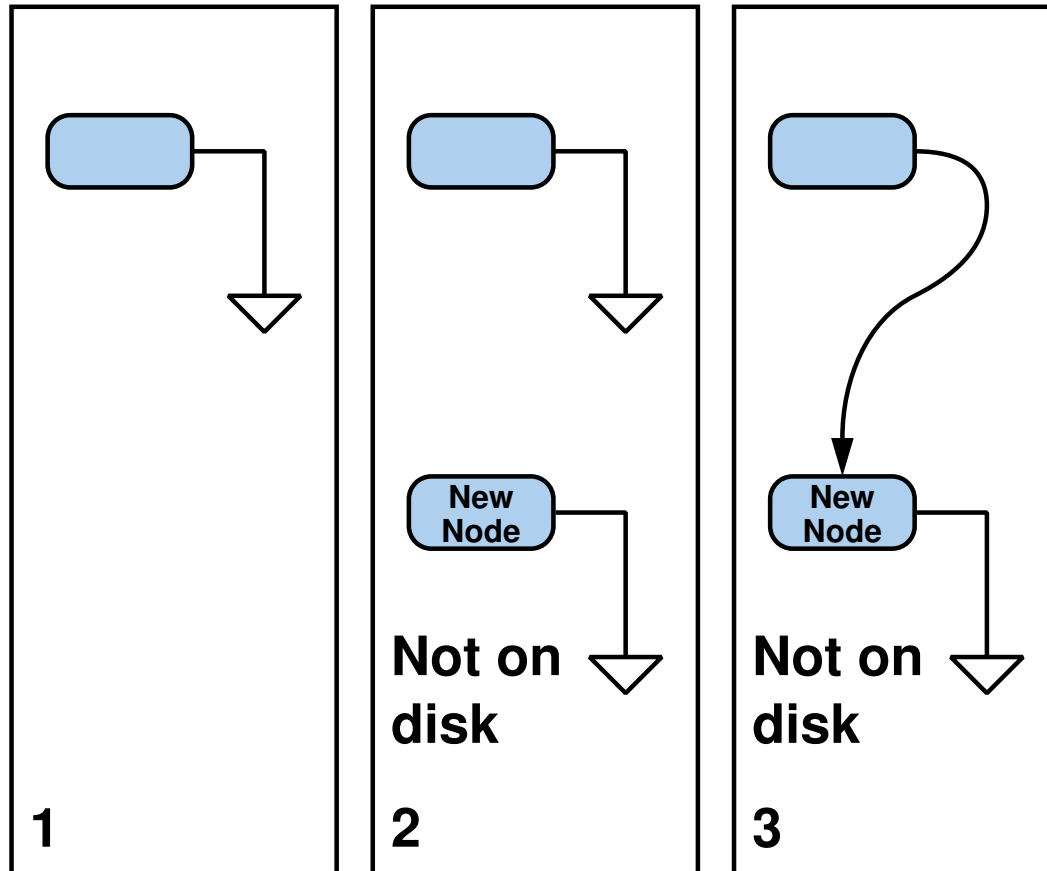
# File-System Consistency (2)



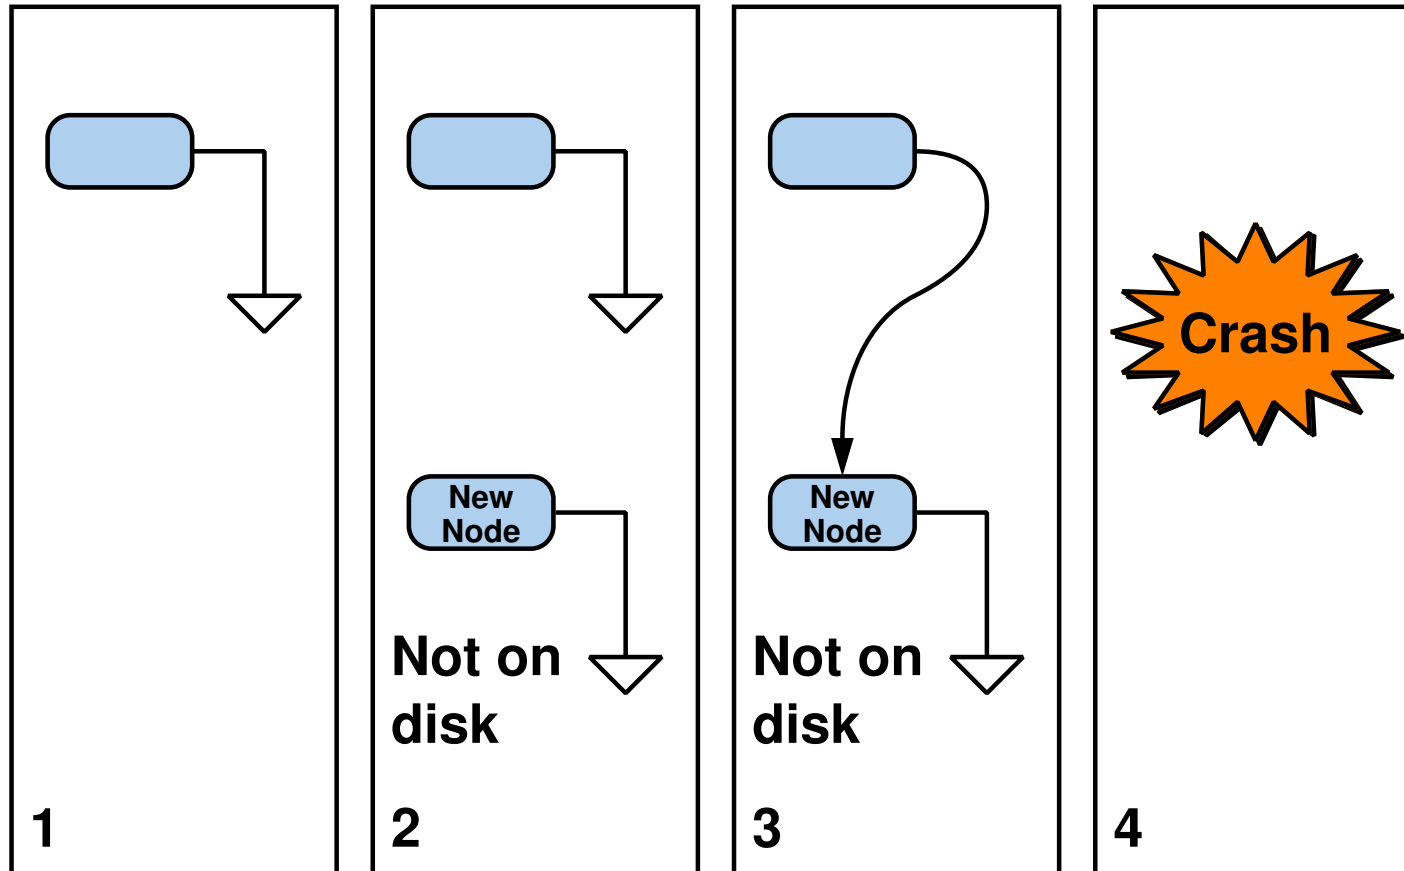
# File-System Consistency (2)



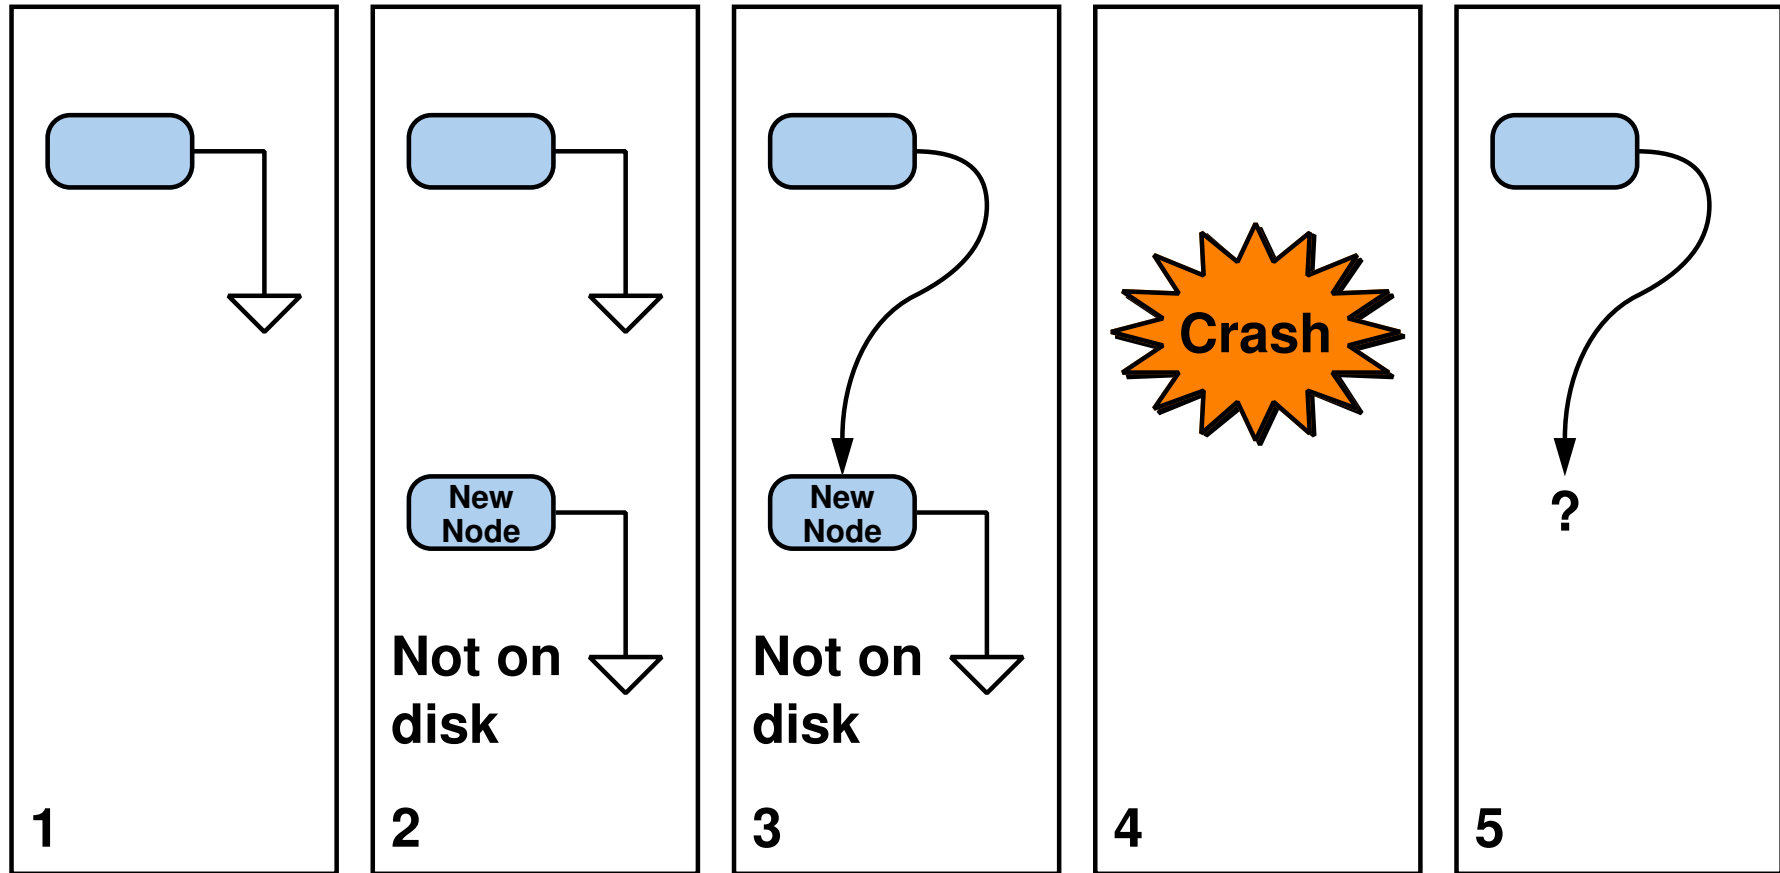
# File-System Consistency (2)



# File-System Consistency (2)




# File-System Consistency (2)




## A More Realistic Example

➡ Let's say that you are appending to file A




- the buffer cache does not know about the relationship between blocks x, y, and z
- techniques like locking (i.e., lock the disk so that it cannot crash when it's locked) won't work






# A More Realistic Example

➡ Let's say that you are appending to file A




- what about this order and crash timing?
- what about other combinations?





# How to Cope ...

- ➡ Don't crash
- ➡ Perform multi-step disk updates in an order such that disk is always consistent, i.e., the *consistency-preserving approach*
- ➡ Perform multi-step disk updates as *transactions*, i.e., implemented so that either all steps take effect or none do




# Maintaining Consistency

1) Write this  
synchronously  
to disk




# Maintaining Consistency

1) Write this  
synchronously  
to disk





2) Then write this  
asynchronously  
via the cache




# Innocuous Inconsistency

After crash:



# Innocuous Inconsistency

After crash:

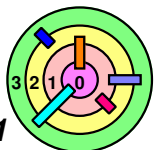


# Soft Updates




## Main idea

- order disk operations to preserve meta-data consistency
  - innocuous inconsistency is considered ok



# Back To The Example




➡ Let's say that you are appending to file A






# Back To The Example

➡ Let's say that you are appending to file A



# Back To The Example


➡ Let's say that you are appending to file A



- is this bad?
- how bad is it?


# Back To The Example

➡ Let's say that you are appending to file A




# Back To The Example

➡ Let's say that you are appending to file A




## Soft Update Example 2



➡ Create a new file with one data block



Choice 1




Choice 2




# Soft Update Example 3


➡ Move a file



Choice 1




Choice 2




# Soft Update

- ➡ **An implementation of the consistency-preserving approach**
  - ▬ **the idea is simple:**
    - **update cache in an order that maintains consistency**
    - **write cache contents to disk in same order in which cache was updated**
  - ▬ **isn't, because reality is more complicated**
    - **(assuming speed is important)**



# Which Order?




— this is easy






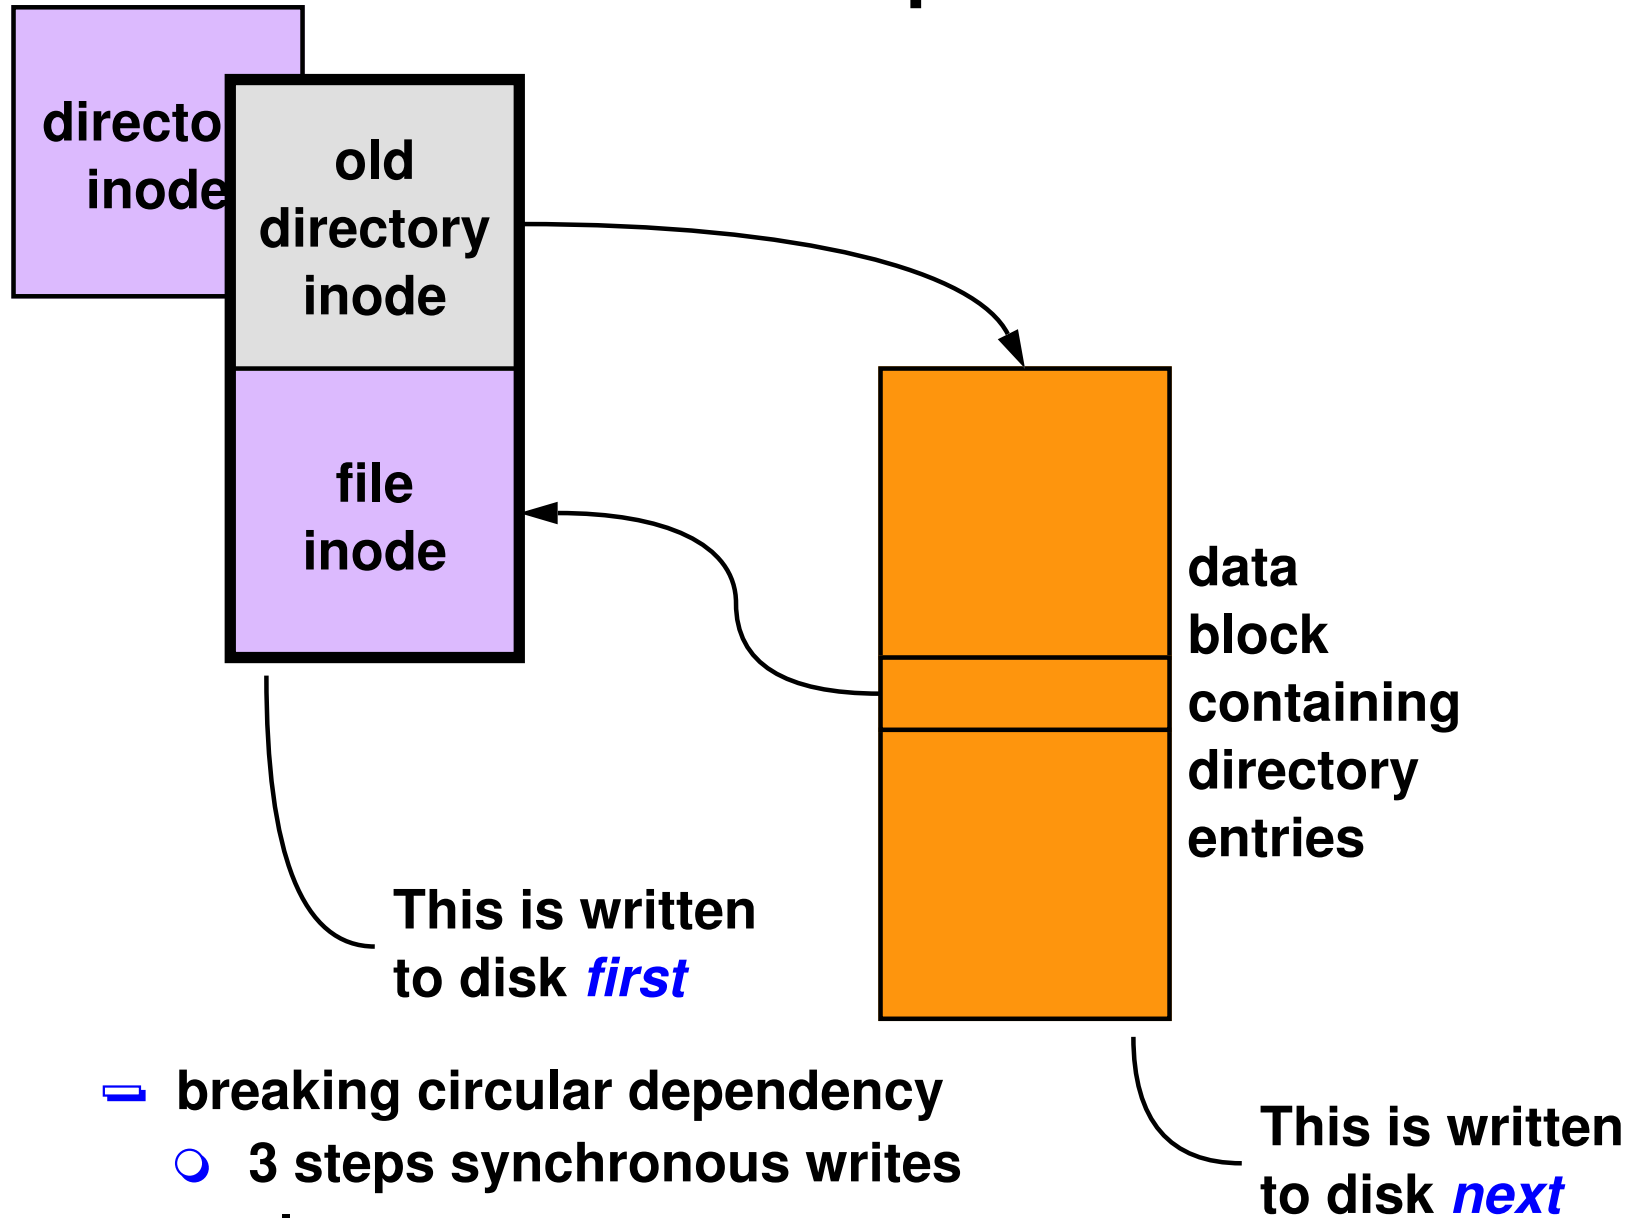
## However ...




- circular dependency
  - in reality, in order to save the number of disk writes, multiple objects can be packed into a disk block



# Soft Updates



- breaking circular dependency
- 3 steps synchronous writes
- slow



# Soft Updates in Practice

- ➡ Implemented for FFS in 1994
- ➡ Used in FreeBSD's FFS
  - improves performance (over FFS with synchronous writes)
  - disk updates may be *many seconds* behind cache updates
  - need to *reclaim lost disk blocks* as background activity after the system restarts

